

# KubeSecOps

Karthik Gaekwad  
Principal Engineer  
Oracle Inc

# Hello!

- I'm Karthik Gaekwad
- NOT a DBA



- <https://cloudnative.oracle.com/>
- Cloud Native evangelist at Oracle Cloud Infrastructure
- Used to be a developer on the OKE Team.



**@iteration1**

# Hello!



ORACLE®

StackEngine



- Been in Industry 15 years.
- In general, I like building stuff with friends.
  - A maintainer for Gauntlt- Open source security scanner.
- Love Teaching and building community.
  - Run Devopsdays Austin, Container Days, Cloud Austin.
  - Chair All Day Devops Cloud Native track.
  - LinkedIn Learning Author for Learning Kubernetes (and more).



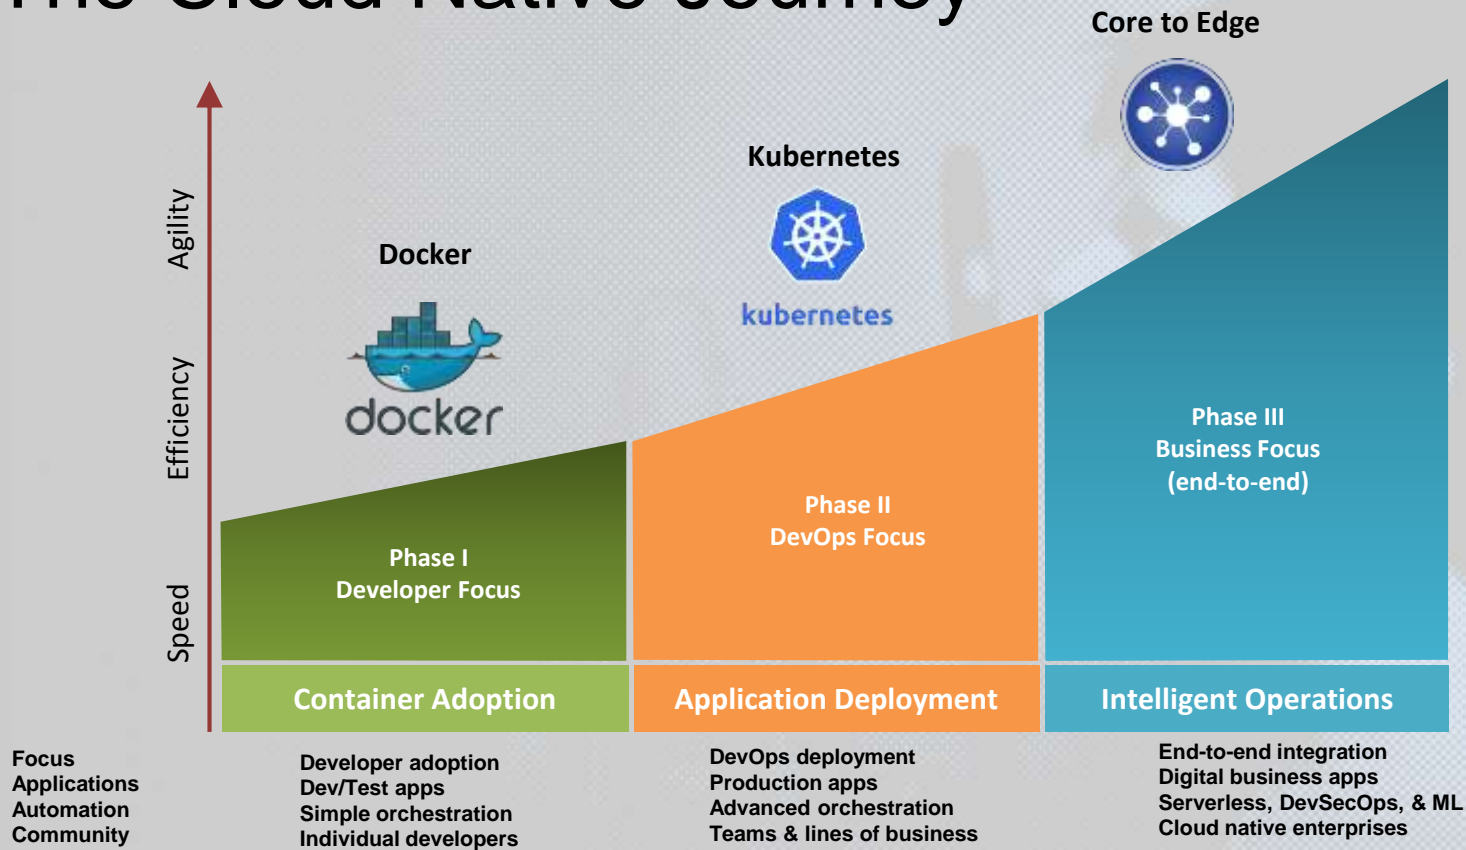
Need an OCI Trial Account?

**ORACLE®**

**Cloud Infrastructure**

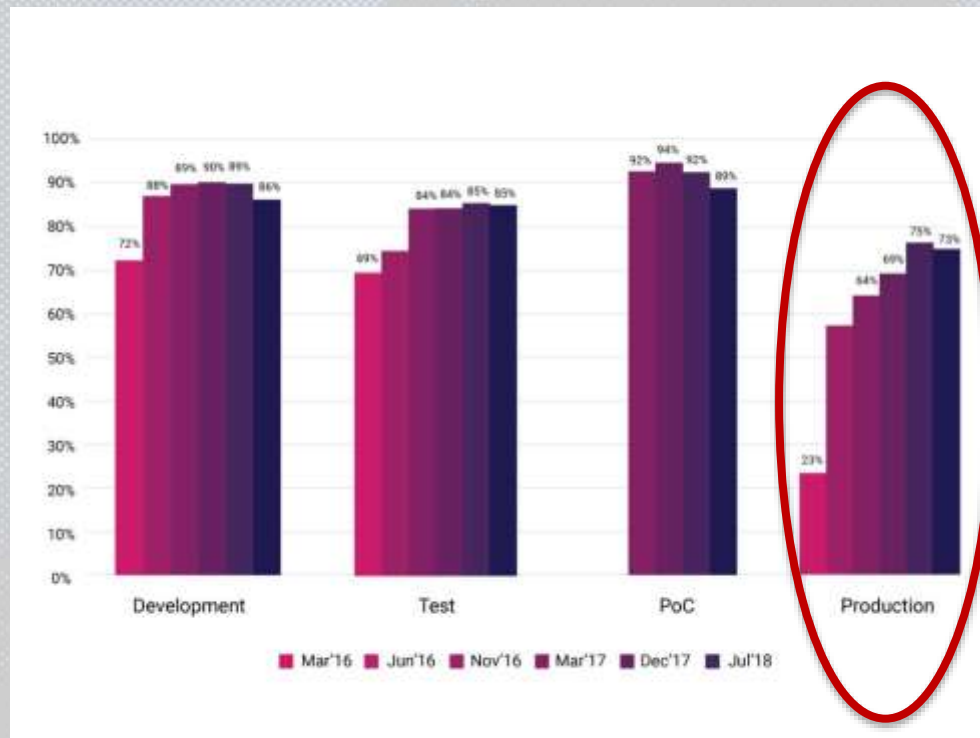
**<http://bitly.com/ocicloud>**

# The Cloud Native Journey



# Latest CNCF Survey: August 2018

Where Does Your  
Company Use Containers?




# Latest CNCF Survey: August 2018

Where Does Your  
Company Use Containers?




# Container Management IS Kubernetes






# Good News, Bad News...

**Good:** On average, CNCF project usage is up over 200% since the Dec 2017!



But...

# Complexity, Culture, Training, & Security Issues Remain



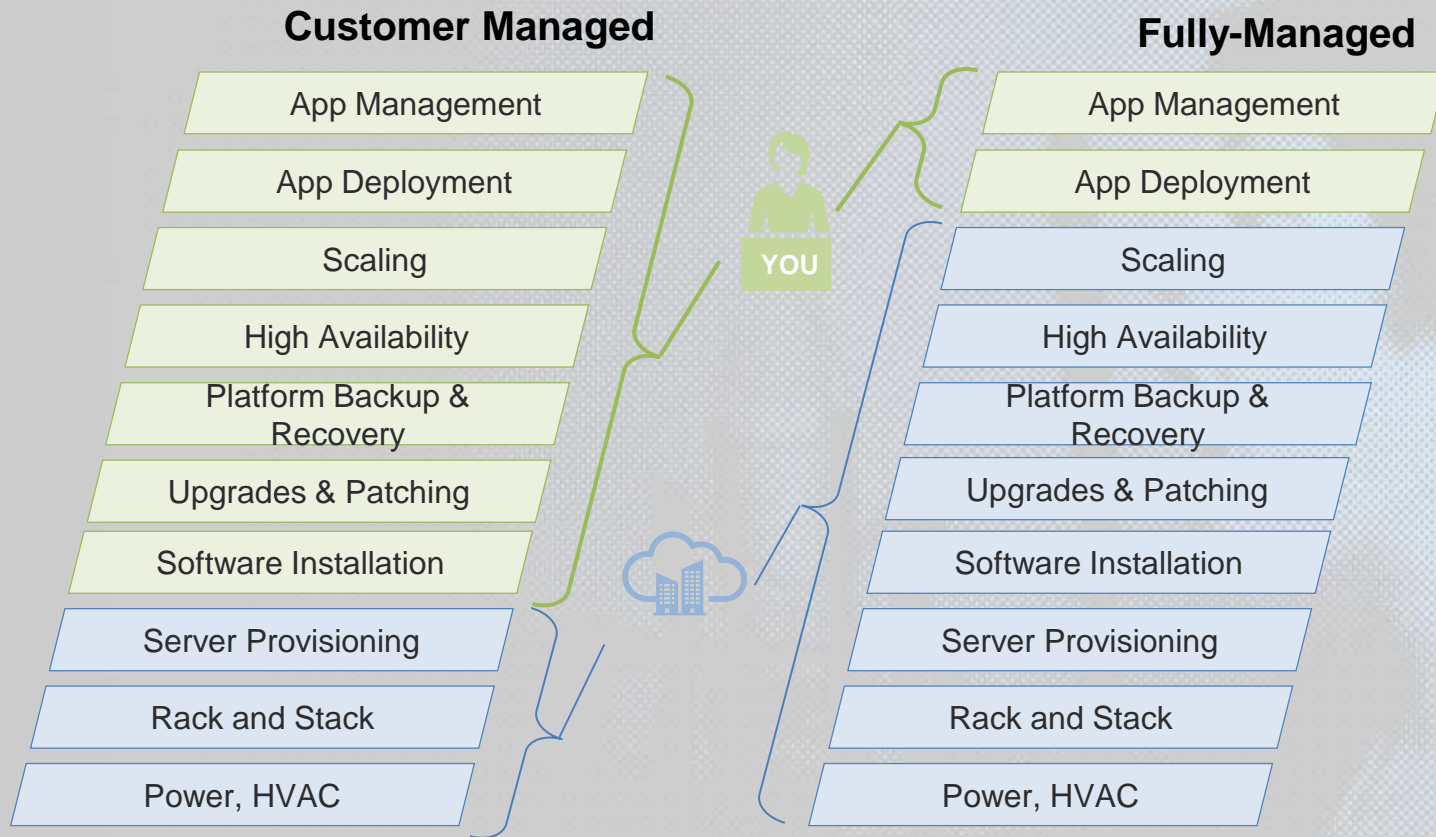
# Kubernetes & Cloud Native Challenges

- Managing, maintaining, upgrading Kubernetes Control Plane
  - API Server, etcd, scheduler etc....
- Managing, maintaining, upgrading Kubernetes Data Plane
  - In place upgrades, deploy parallel cluster etc....
- Figuring out container networking & storage
  - Overlays, persistent storage etc... - it should just work
- Managing Teams
  - How do I manage & control team access to my clusters?
- Security, security, security



Source: Oracle Customer Survey 2018

# How Teams Address Issues?



## Benefits

- ✓ Faster Time to Deploy
- ✓ Lower Risk
- ✓ Accelerate Innovation



A close-up shot of Gene Wilder as Gene Wilder, wearing a purple velvet suit, a white shirt, a tan bow tie, and a brown top hat. He has curly blonde hair and is smiling, resting his head on his right hand. The background is dark and out of focus.

Which brings us to security...

Security....



Where no news, is good news!

# Unsecured K8s dashboards



< Back

Research

## Lessons from the Cryptojacking Attack at Tesla

by RedLock CSI Team | 02.20.18, 6:00 AM

- Unsecured Kubernetes Dashboard with account creds.
- Used this to mine cryptocurrency.
- 2017: Aviva
- 2018: Tesla, Weight Watchers
- <https://redlock.io/blog/cryptojacking-tesla>



# Kubelet credentials hack

The screenshot shows a HackerOne report page. At the top, the HackerOne logo and navigation links are visible. The report title is "SSRF in Exchange leads to ROOT access in all instances". Below the title, there are details about the report: it was reported on May 23, 2018, at 4:59pm, and it was assigned to Shopify. The report is categorized as "Server-Side Request Forgery (SSRF)" and has a bounty of \$25,000. The report is marked as "Resolved (Closed)".

**SUMMARY BY SHOPIFY**

Shopify infrastructure is isolated into subsets of infrastructure. @0xach reported it was possible to gain root access to any container in one particular subset by exploiting a server side request forgery bug in the screenshotting functionality of Shopify Exchange. Within an hour of receiving the report, we disabled the vulnerable service, began auditing applications in all subsets and remediating across all our infrastructure. The vulnerable subset did not include Shopify.com.

After auditing all services, we fixed the bug by deploying a metadata containment proxy to disable access to metadata information. We also disabled access to internal IPs on all infrastructure subsets. We awarded this \$25,000 as a Shopify Core RCE since some applications in this subset do have access to some Shopify core data and systems.

**THE EXPLOIT CHAIN - How to get root access on all Shopify instances**

1 - Access Google Cloud Metadata

- 1: Create a store (sothers.shopify.com)
- 2: Edit the template: password.liquid and add the following content:

```
<script>
window.location="https://metadata.google.internal/computeMetadata/v1beta1/instance/permissions/default"
// This one don't work here because Google Cloud sets the "X-Frame-Options: DENYFRAME" header.
</script>
```

- Shopify: Server Side request Forgery
- Get kubelet certs/private key
- Root access to any container in part of infrastructure.
- <https://hackerone.com/reports/341876>



← → ↻ https://www.shodan.io/search?query=port%3A2379+product%3A%22etcd%22

**SHODAN** port:"2379" product:"etcd" 🔍 🌐 Explore Downloads Reports Developer Pricing Enterprise Access


🔪 Exploits 🗺 Maps 📄 Share Search 📄 Download Results 📄 Create Report

---

**TOTAL RESULTS**

2,367

**TOP COUNTRIES**



Country	Count
China	933
United States	602
Germany	163
France	110
Singapore	67

**TOP ORGANIZATIONS**

Organization	Count
Amazon.com	323
Hangzhou Alibaba Advertising Co.,Ltd.	276
Tencent cloud computing	207
Hetzner Online GmbH	70
Digital Ocean	70

**TOP OPERATING SYSTEMS**

OS	Count
Linux 3.x	1

**TOP VERSIONS**

Version	Count
3.3.2	299
3.2.18	136
3.3.9	126
3.2.22	123
3.2.17	82

---

**10.24.237.101**

**Tencent cloud computing**

Added on 2018-09-26 19:18:25 GMT

🇨🇳 China

[Details](#)

etcd

Name: etcd\_10.0.128.13

Version: 3.3.2

Uptime: 79h30m33.232183055s

Peers: http://10.0.128.13:2380

---

**10.236.207.100**

[2 compute.amazonaws.com](#)

**Amazon Corporate Services Pty**

Added on 2018-09-26 19:58:42 GMT

🇦🇺 Australia, Sydney

[Details](#)

[cloud](#)

etcd

Name: NFR-50nodeap-southeast-2002

Version: 3.3.2

Uptime: 54h22m1.357953836s

Peers: http://13.236.207.66:2380

---

**10.11.36.2**

**Spectrum Business**

Added on 2018-09-26 18:52:04 GMT

🇺🇸 United States, Flower Mound

[Details](#)

etcd

Name: core2

Version: 2.2.5

Uptime: 5h38m59.402175596s

Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

---

**192.168.4.6**

**Red Hat**

Added on 2018-09-26 18:48:50 GMT

🇺🇸 United States

[Details](#)

etcd

Name: rdcloud-devstack4.rdcloud

Version: 3.2.17

Uptime: 633h2m36.382059797s

Peers: http://192.168.4.6:2380

---

**04.100.120.122.1**

[amazonaws.com](#)

**Amazon.com**

Added on 2018-09-26 18:45:20 GMT

🇺🇸 United States, Ashburn

[Details](#)

etcd

Name: master-0

Version: 3.2.18

Uptime: 20m10.490053853s


← → ↻ https://www.shodan.io/search?query=port%3A2379+product%3A%22etcd%22

SHODAN port:"2379" product:"etcd" 🔍 🌐 Explore Downloads Reports Developer Pricing Enterprise Access

🗺️ Maps 📄 Share Search 📄 Download Results 📄 Create Report

TOTAL RESULTS

2,367



China	933
United States	802
Germany	163
France	110
Singapore	67

TOP ORGANIZATIONS

Amazon.com	323
Hangzhou Alibaba Advertising Co.,Ltd.	276
Tencent cloud computing	267
Hetzner Online GmbH	70
Digital Ocean	70

TOP OPERATING SYSTEMS

Linux 3.x	1
-----------	---

TOP VERSIONS

3.3.2	299
3.2.18	136
3.3.9	126
3.2.22	123
3.2.17	82

**10.24.230.181**

**Tencent cloud computing**

Added on 2018-09-26 19:18:25 GMT

🇨🇳 China

[Details](#)

etcd

Name: etcd\_10.0.128.13

Version: 3.3.2

Uptime: 79h30m33.232183055s

Peers: http://10.0.128.13:2380

**10.236.207.100**

**Amazon Corporate Services Pty**

Added on 2018-09-26 19:58:42 GMT

🇦🇺 Australia, Sydney

[Details](#)

cloud

etcd

Name: NFR-50nodeap-southeast-2002

Version: 3.3.2

Uptime: 54h22m1.357953836s

Peers: http://13.236.207.66:2380

**10.11.36.2**

**Spectrum Business**

Added on 2018-09-26 18:53:04 GMT

🇺🇸 United States, Flower Mound

[Details](#)

etcd

Name: core2

Version: 2.2.5

Uptime: 5h38m59.402175596s

Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

**192.168.4.6**

**Red Hat**

Added on 2018-09-26 18:48:50 GMT

🇺🇸 United States

[Details](#)

etcd

Name: rdcloud-devstack4.rdcloud

Version: 3.2.17

Uptime: 633h2m36.382059797s

Peers: http://192.168.4.6:2380

**04.100.120.122.1**

**Amazon.com**

Added on 2018-09-26 18:45:20 GMT

🇺🇸 United States, Ashburn

[Details](#)

etcd

Name: master-0


Version: 3.2.18

Uptime: 20m10.490053853s

SHODAN port:"2379" product:"etcd"

Maps Share Search Download Results Create Report

TOTAL RESULTS: 2,367



Country	Count
China	933
United States	802
Germany	163
France	110
Singapore	57

TOP OPERATING SYSTEMS

OS	Count
Linux 3.x	1

TOP VERSIONS

Version	Count
3.3.2	299
3.2.18	136
3.3.9	126
3.2.22	123
3.2.17	82

10.24.200.101  
Tencent cloud computing  
Added on 2018-09-26 19:18:25 GMT  
China  
Details

etcd  
Name: etcd\_10.0.128.13  
Version: 3.3.2  
Uptime: 79h30m33.232183055s  
Peers: http://10.0.128.13:2380

10.236.207.66  
Amazon Web Services Pty  
Added on 2018-09-26 19:08:42 GMT  
Sydney  
Details

etcd  
Name: NFR-50nodeap-southeast-2002  
Version: 3.3.2  
Uptime: 54h22m1.357953836s  
Peers: http://13.236.207.66:2380

10.11.36.2  
Spectrum Business  
Added on 2018-09-26 18:52:04 GMT  
United States, Flower Mound  
Details

etcd  
Name: core2  
Version: 3.2.5  
Uptime: 5h38m59.402175596s  
Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

192.168.4.6  
Red Hat  
Added on 2018-09-26 18:48:50 GMT  
United States  
Details

etcd  
Name: rdocloud-devstack4.rdocloud  
Version: 3.2.17  
Uptime: 633h2m36.382059797s  
Peers: http://192.168.4.6:2380

0.100.120.122.1  
Amazon.com  
Added on 2018-09-26 18:45:20 GMT  
United States, Ashburn  
Details

etcd  
Name: master-0  
Version: 3.2.18  
Uptime: 20h10.490053853s

SHODAN port:"2379" product:"etcd"

Maps Share Search Download Results Create Report

TOTAL RESULTS: 2,367

World map showing search results by country.

Country	Count
China	933
United States	802
Germany	163
France	110
Singapore	57

TOP VERSIONS

Version	Count
3.3.2	299
3.2.18	136
3.3.9	126
3.2.22	123
3.2.17	82

etcd

Name: etcd\_10.0.128.13  
Version: 3.3.2  
Uptime: 79h30m33.232183055s  
Peers: http://10.0.128.13:2380

etcd

Name: NFR-50nodeap-southeast-2002  
Version: 3.3.2  
Uptime: 54h22m1.357953836s  
Peers: http://13.236.207.66:2380

etcd

Name: core2  
Version: 2.2.5  
Uptime: 5h38m59.402175596s  
Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

etcd

Name: rdcloud-devstack4.rdcloud  
Version: 3.2.17  
Uptime: 633h2m36.382059797s  
Peers: http://192.168.4.6:2380

etcd

Name: master-0  
Version: 3.2.18  
Uptime: 20m10.40053853s



SHODAN

port:2379\* product:etcd

Explore Downloads Reports Developer Pricing Enterprise Access

Maps Share Search Download Results Create Report

TOTAL RESULTS: 2,367

China 933  
United States 802  
Germany 163  
France 110  
Singapore 1

Amazon.com  
Hangzhou Alibaba Cloud Services  
Tencent cloud computing  
Hetzner Online GmbH  
Digital Ocean

etcd  
Name: etcd\_10.0.128.13  
Version: 3.3.2  
Uptime: 79h30m33.232103055s  
Peers: http://10.0.128.13:2380

etcd  
Name: NFR-jendrap-0001  
Version: 3.3.2  
Uptime: 54h22m1.3579538s  
Peers: http://13.236.201.10:2380

etcd  
Name: core2  
Version: 2.2.5  
Uptime: 5h38m59.4021755s  
Peers: http://10.11.36.2:2380

etcd  
Name: rdocloud-devstack4.rdocloud  
Version: 3.2.17  
Uptime: 633h2m36.382059797s  
Peers: http://192.168.4.6:2380

etcd  
Name: master-0  
Version: 3.2.18  
Uptime: 20m10.400053053s  
Peers: http://10.0.0.1:2380

Linux 3.x 1

TOP VERSIONS

3.3.2	299
3.2.18	136
3.3.9	126
3.2.22	123
3.2.17	82

A close-up portrait of Gene Wilder as Charlie Bucket from the 1971 film "Charlie and the Chocolate Factory". He is wearing a brown top hat, a purple velvet jacket, and a light-colored bow tie. He has curly blonde hair and a surprised expression. The background is blurred, showing festive lights.

How did we get here?

A person in a flight suit is shown from the side, operating a complex cockpit. The cockpit is filled with numerous analog gauges, dials, and control panels. The person's right hand is visible, reaching towards a control knob. The overall scene is dimly lit, with the primary light source coming from the cockpit's instruments. The text "Kubernetes is too complicated" is overlaid in a bright yellow font across the upper portion of the image.

“Kubernetes is too complicated”

A person in a flight suit is shown from the side, operating a complex control panel. The panel is filled with numerous circular gauges, dials, and buttons, some of which are illuminated with green and red lights. The person's hand is visible, interacting with one of the controls. The overall scene suggests a high-tech, possibly military or aviation, environment.

**“Kubernetes is too complicated”**

**“We hope it’ll get easier”**





**I want to get better!  
But where to start?**

# Let's look at:

## Attack Surface

- More importantly, how to limit damage

## Security related features in K8s

- The more you know, the better you build

## Opensource Tooling to help

- Because we all need help

# Attack Surface

# Attack Surface

**Goal: Reduce the attack surface**

Analysis for:

- Host
- Container (Images and running)
- Kubernetes Cluster



# Attack Surface: Host

- These are the machines you're running Kubernetes on.
- Age old principles of Linux still apply:
  - Enable SELinux
  - AppArmor
  - Seccomp
  - Hardened Images
- Goal: Minimize privilege to applications running on the host
- Good news: Already a wealth of information on this subject!
  - <http://lmgty.com/?q=how+to+reduce+attack+surface+linux>

# Attack Surface: Container Images

**GOAL: Know your base image  
when building containers**

# Attack Surface: Container Images

**GOAL: Know your base image when building containers**

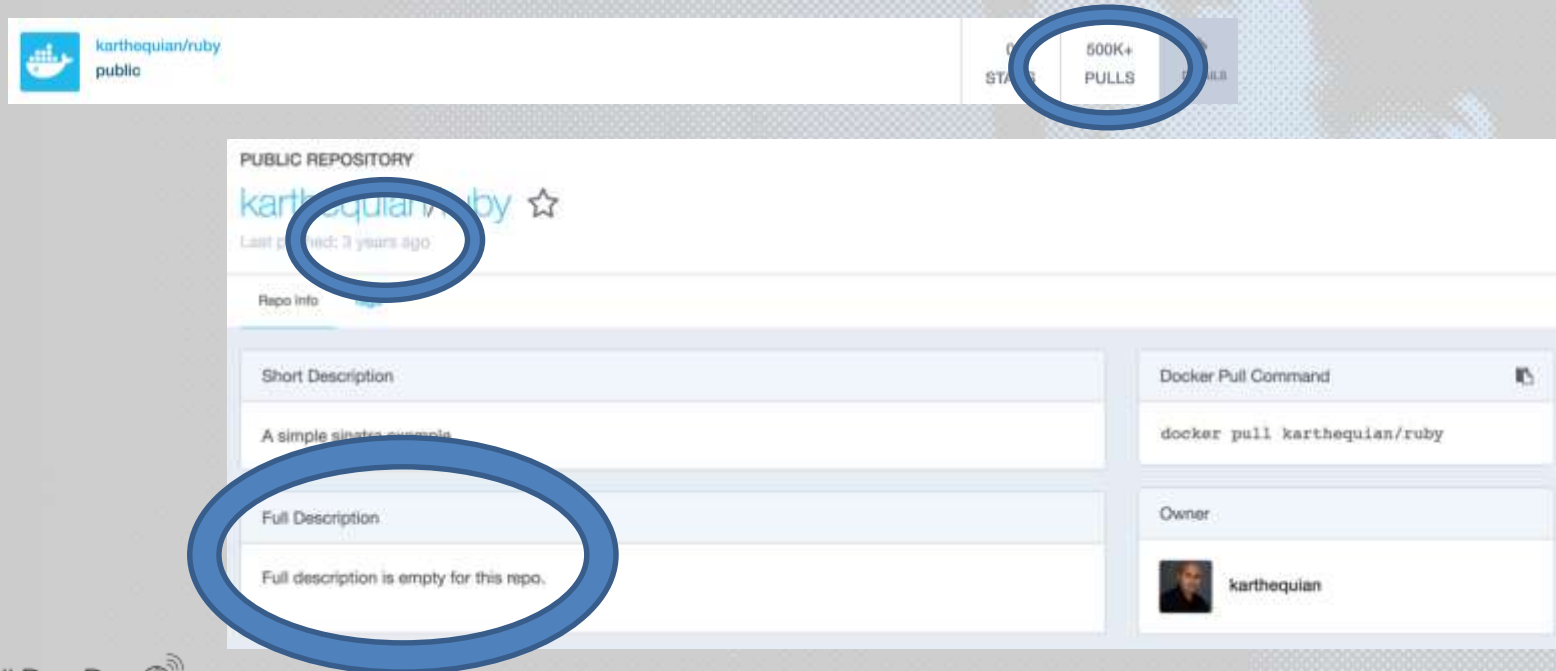
The screenshot shows the Docker Hub page for the repository `karthequian/ruby`. The repository name is circled in blue. The '500K+ PULLS' badge is also circled in blue. The repository description is 'A simple sinatra example'. The full description is empty.

**Repository Information:**

- Repository Name: `karthequian/ruby`
- Status: Public
- Pulls: 500K+
- Last pushed: 3 years ago
- Short Description: A simple sinatra example
- Full Description: Full description is empty for this repo.
- Docker Pull Command: `docker pull karthequian/ruby`
- Owner: karthequian

# Attack Surface: Container Images

**GOAL: Know your base image when building containers**





# Attack Surface: Container Images

**GOAL: Know your base image when building containers**

When in doubt, stick to an official images!



Or start from a sane base image (example: alpine linux)

# Attack Surface: Container Images

**GOAL: Smaller the image, the better**

- Less things for an attacker to exploit.
- Quicker to push, quicker to pull.

# Attack Surface: Container Images

## **GOAL: Don't rely on :latest tag**

- :latest image yesterday might not be :latest image tomorrow
- Instead, you'd want to know what specific version you're operating with.
- Side benefit: If there is a new vulnerability announced for OS version x.y.z, you know immediately whether you're running that version!

# Attack Surface: Container Images

**GOAL: Check for vulnerabilities  
periodically**

- Plenty of ways to do this in registries. We'll cover more in the tooling section



# Attack Surface: Running Containers

## **GOAL: Don't run as root**

- Containers running as root might be completely unnecessary for the actual application.
- If compromised, attacker can do a lot more things..
- Pod security policies can help (we'll see how later).

# Attack Surface: Running Containers

## **GOAL: Limit host mounts**

- Be wary of images that require broad access to paths on the host
- Limit your host mount to a smaller subset of directories
- Reduces blast radius on compromise

# **Attack Surface: Kubernetes Cluster**

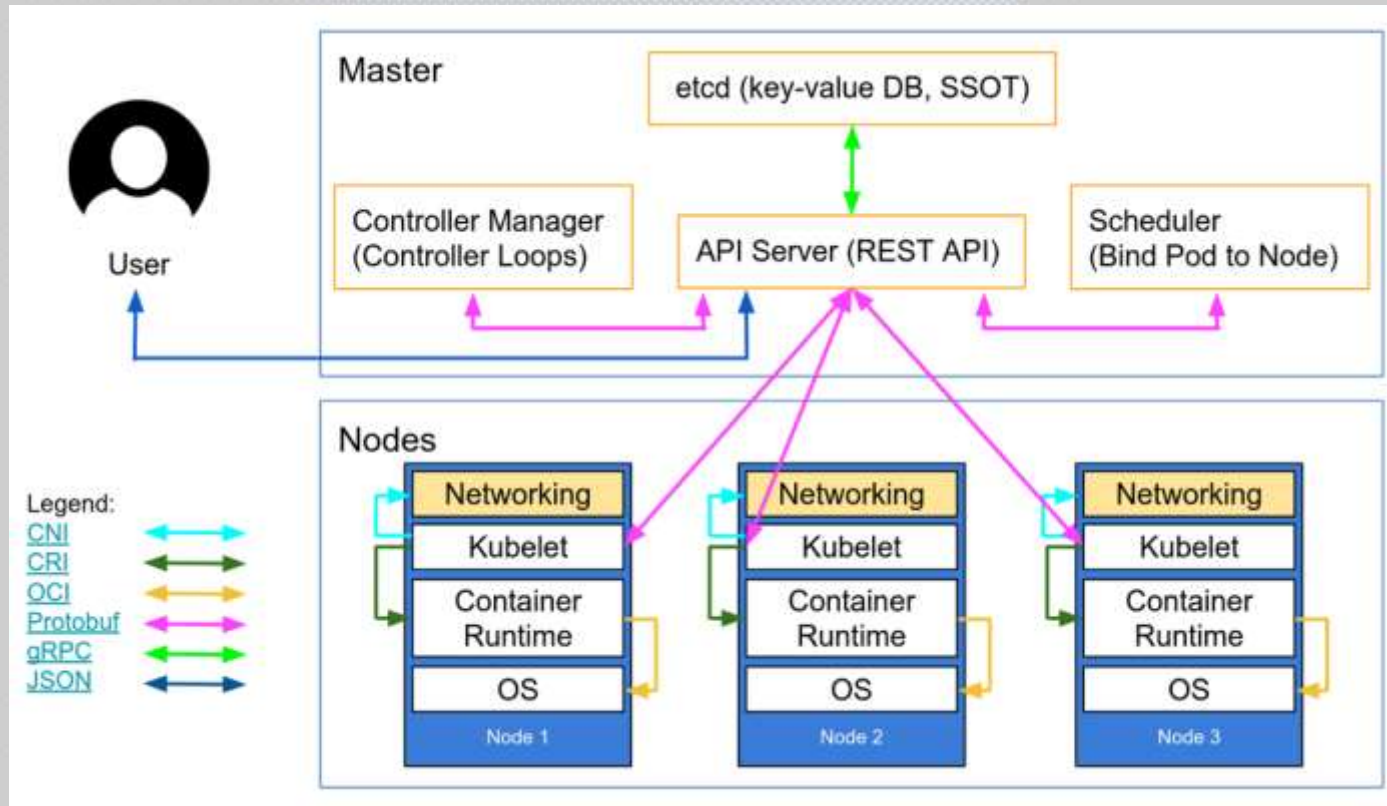
# Kubernetes Cluster- TLS

## TLS ALL THE THINGS





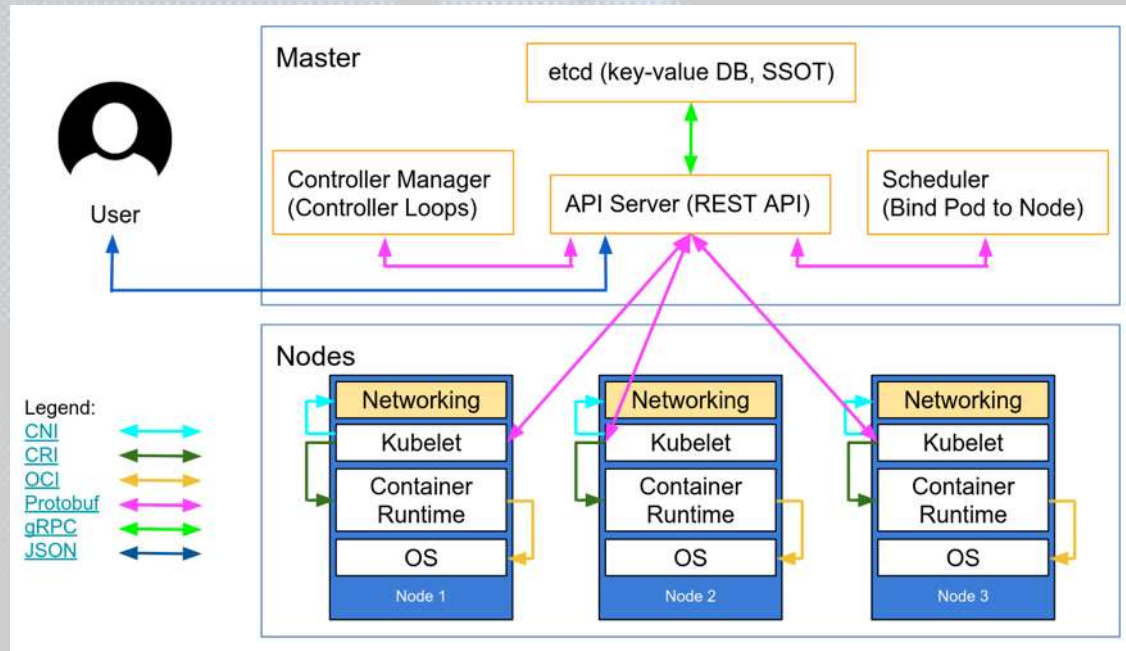
# Kubernetes Cluster- TLS



# Kubernetes Cluster- TLS

## TLS Checklist:

1. User and Master
2. Nodes and Master
3. Everything etcd
4. Kubelet to API Server



We're a little  
better off now.

But what else to do?



# K8s Features

How can the platform help me make secure choices?





# K8s Features

Authentication

Authorization

Audit Logging

Network Policies

Pod security policies

Kubernetes Secrets



# Authentication and Authorization

Do you know how you are authenticating with Kubernetes?

Many ways to Authenticate

- Client Certs
- Static token file
- Service Account tokens
- OpenID
- Webhook Mode
- And more (<https://kubernetes.io/docs/reference/access-authn-authz/authentication/>)

A close-up photograph of a small, grey and white striped kitten. The kitten is lying under a thick, brown, fuzzy blanket, with only its head and front paws visible. It has large, round, orange-brown eyes and is looking directly at the camera with a curious expression. The background is softly blurred, showing a patterned pillow and a light-colored surface.

**Goal: Pick a strategy that  
fits your use case**

**Whatever you do,  
DO NOT YOLO!**



**If you DO NOT YOLO...**



**You can pick an authz strategy..**



# Authentication and Authorization

<https://kubernetes.io/docs/reference/access-authn-authz/authorization/>

## Authorization Modules

- **Node** - A special-purpose authorizer that grants permissions to kubelets based on the pods they are scheduled to run. To learn more about using the Node authorization mode, see [Node Authorization](#).
- **ABAC** - Attribute-based access control (ABAC) defines an access control paradigm whereby access rights are granted to users through the use of policies which combine attributes together. The policies can use any type of attributes (user attributes, resource attributes, object, environment attributes, etc). To learn more about using the ABAC mode, see [ABAC Mode](#).
- **RBAC** - Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within an enterprise. In this context, access is the ability of an individual user to perform a specific task, such as view, create, or modify a file. To learn more about using the RBAC mode, see [RBAC Mode](#)
  - When specified RBAC (Role-Based Access Control) uses the `rbac.authorization.k8s.io` API group to drive authorization decisions, allowing admins to dynamically configure permission policies through the Kubernetes API.
  - To enable RBAC, start the apiserver with `--authorization-mode=RBAC`.
- **Webhook** - A WebHook is an HTTP callback: an HTTP POST that occurs when something happens; a simple event-notification via HTTP POST. A web application implementing WebHooks will POST a message to a URL when certain things happen. To learn more about using the Webhook mode, see [Webhook Mode](#).

# Authentication and Authorization

- Pro tip: Nobody uses ABAC anymore. Don't be that guy....
- RBAC is the defacto standard
  - Based on roles and role bindings
  - Good set of defaults: <https://github.com/uruddarraju/kubernetes-rbac-policies>
- Can use multiple authorizers together, but can get confusing.
  - 1<sup>st</sup> authorizer to authorize passes authz

# Kubernetes Cluster- Audit Logs

- Wat?
- “Kubernetes auditing provides a security-relevant chronological set of records documenting the sequence of activities that have affected system by individual users, administrators or other components of the system.”
- Answers: What/when/who/where information on security events.
- **Your job:** Periodically watch Kubernetes Audit logs
- <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>





See, you know how to take the reservation, you just don't know how to hold the reservation and that's really the most important part of the reservation, the holding. Anybody can just take them.

— *Jerry Seinfeld* —

**AZ QUOTES**



# Kubernetes Cluster- Network Policies

- Consider adding a network policy to the cluster...
- Default Policy: All pods can talk to all other pods.
- Consider limiting this with a Network Policy
- <https://kubernetes.io/docs/concepts/services-networking/network-policies/>

# Kubernetes Cluster- Pod Security Policies

- Consider adding Pod Security policies
- PodSecurityPolicy: A Defined set of conditions a pod must run with.
- Think of this as authorization for pods.

# Kubernetes Cluster: Pod Security Policies

Control Aspect	Field Names
Running of privileged containers	<code>privileged</code>
Usage of host namespaces	<code>hostPID</code> , <code>hostIPC</code>
Usage of host networking and ports	<code>hostNetwork</code> , <code>hostPorts</code>
Usage of volume types	<code>volumes</code>
Usage of the host filesystem	<code>allowedHostPaths</code>
White list of Flexvolume drivers	<code>allowedFlexVolumes</code>
Allocating an FSGroup that owns the pod's volumes	<code>fsGroup</code>
Requiring the use of a read only root file system	<code>readOnlyRootFilesystem</code>
The user and group IDs of the container	<code>runAsUser</code> , <code>supplementalGroups</code>
Restricting escalation to root privileges	<code>allowPrivilegeEscalation</code> , <code>defaultAllowPrivilegeEscalation</code>
Linux capabilities	<code>defaultAddCapabilities</code> , <code>requiredDropCapabilities</code> , <code>allowedCapabilities</code>
The SELinux context of the container	<code>seLinux</code>
The AppArmor profile used by containers	<code>annotations</code>
The seccomp profile used by containers	<code>annotations</code>
The sysctl profile used by containers	<code>annotations</code>

Capability for an admin to control specific actions

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/#what-is-a-pod-security-policy>

# Kubernetes Secrets

- **GOAL: Use Kubernetes secrets to store sensitive data instead of config maps.**
- Also look at: secrets encryption provider.
  - Controls how etcd encrypts API data
  - **---experimental-encryption-provider-config**
- <https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/>



# Opensource Tooling



# Keep tabs on the CNCF Security landscape

**CNCF Cloud Native Interactive Landscape**

CNCF Cloud Native Interactive Landscape (CNCF Cloud Native Interactive Landscape) provides a good introduction. The associated landscape (gig) and associated landscape (gig) are specifically provided below. Please open a pull request to correct any errors. (Copyrights are not owned by CNCF, last updated: 2019-08-08 16:01:00)

You are viewing 20 cards with a total of 18,196 stars, market cap of \$1100 and funding of \$410M.

1,000 10,000 100,000

Including CNCF Projects (1)

**Security** (1)

- Security** (Cloud Native Computing Foundation) (CNCF)
- The Open Policy Agent (OPA)** (Cloud Native Computing Foundation) (CNCF)

Excluding CNCF Projects (1)

- Open Policy Agent (OPA)** (Cloud Native Computing Foundation) (CNCF)

1,000 10,000 100,000

CNCF Member Projects (10)

Project	Funding	Stars	Market Cap
aqua	Funding: \$10M	1,000	\$10M
clair	Funding: \$10M	1,000	\$10M
Datica	Funding: \$10M	1,000	\$10M
kube-bench	Funding: \$10M	1,000	\$10M
kube-hunter	Funding: \$10M	1,000	\$10M
NeuVector	Funding: \$10M	1,000	\$10M
OpenSCAP	Funding: \$10M	1,000	\$10M
ORACLE POLICY AUTOMATION	Funding: \$10M	1,000	\$10M
Sysdig Falco	Funding: \$10M	1,000	\$10M
Twistlock	Funding: \$10M	1,000	\$10M

Excluding CNCF Member Projects (10)

Project	Funding	Stars	Market Cap
anchore	Funding: \$10M	1,000	\$10M
BLACKBUCK	Funding: \$10M	1,000	\$10M
Grafeas	Funding: \$10M	1,000	\$10M
ORF / Hydra	Funding: \$10M	1,000	\$10M
StackRox	Funding: \$10M	1,000	\$10M
WhiteSource	Funding: \$10M	1,000	\$10M

Copyrights data is used under license from Copyrights to CNCF. For more information, please see the [license](#) file.

<https://landscape.cncf.io/landscape=security-compliance>

# CNCF Projects



- “The Update Framework”
- Is a framework or a methodology.
- Used for secure software updates.
- Based on ideas surrounding trust and integrity.



- Is a project.
- Based on TUF.
- A solution to secure software updates and distribution.
- Used in Docker Trusted Registry.



# Clair



- Open source project for the static analysis of vulnerabilities in containers.
- Find vulnerable images in your repo.
- Built into quay.io, but you can add to your own repo.
- <https://github.com/coreos/clair>





Quay Security Scanner has detected **13** vulnerabilities.  
Patches are available for **4** vulnerabilities.

- 1 High-level vulnerabilities.
- 1 Medium-level vulnerabilities.
- 2 Low-level vulnerabilities.
- 5 Negligible-level vulnerabilities.
- 4 Unknown-level vulnerabilities.

### Image Vulnerabilities

☐ Only show fixable

CVE	SEVERITY	PACKAGE	CURRENT VERSION	FIXED VERSION	INTRODUCTION RANGE
CVE-2013-7443	High	lnux	3.16.7-ckt20-1+deb8u3	Fixed	apt-get up.
CVE-2015-5276	Medium	gcc-4.9	4.9.3-10	Fixed	file:b5391.
CVE-2016-2858	Low	glibc	2.19-18+deb8u3	Fixed	file:b5391.
CVE-2016-0623	Low	lnux	3.16.7-ckt20-1+deb8u3	Fixed	apt-get up.
CVE-2005-3660	Negligible *	lnux	3.16.7-ckt20-1+deb8u3	Fixed	apt-get up.
CVE-2015-4093	Negligible *	lnux	3.16.7-ckt20-1+deb8u3	Fixed	apt-get up.
CVE-2008-4108	Negligible *	python-defaults	2.7.9-1	Fixed	apt-get up.
CVE-2013-8830	Negligible	lnux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4	apt-get up.
CVE-2013-4392	Negligible *	systemd	215-17+deb8u3	Fixed	file:b5391.
CVE-2015-7515	Unknown	lnux	3.16.7-ckt20-1+deb8u3	Fixed	apt-get up.
CVE-2015-8816	Unknown	lnux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4	apt-get up.
CVE-2016-2547	Unknown	lnux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4	apt-get up.
CVE-2016-2545	Unknown	lnux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4	apt-get up.

# Kube-bench



- Checks whether a Kubernetes cluster is deployed according to security best practices.
- Run this after creating your K8s cluster.
- <https://github.com/aquasecurity/kube-bench>
- Defined by the CIS Benchmarks Docs:  
<https://www.cisecurity.org/cis-benchmarks/>
- Run it against your Kubernetes Master, or Kubernetes node.

# Kube-bench example

```
~$ kubectl logs kube-bench-node
[INFO] 2 Worker Node Security Configuration
[INFO] 2.1 Kubelet
[FAIL] 2.1.1 Ensure that the --allow-privileged argument is set to false (Scored)
[PASS] 2.1.2 Ensure that the --anonymous-auth argument is set to false (Scored)
[PASS] 2.1.3 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Scored)
[PASS] 2.1.4 Ensure that the --client-ca-file argument is set as appropriate (Scored)
[PASS] 2.1.5 Ensure that the --read-only-port argument is set to 0 (Scored)
[FAIL] 2.1.6 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Scored)
[FAIL] 2.1.7 Ensure that the --protect-kernel-defaults argument is set to true (Scored)
[FAIL] 2.1.8 Ensure that the --make-iptables-util-chains argument is set to true (Scored)
[FAIL] 2.1.9 Ensure that the --keep-terminated-pod-volumes argument is set to false (Scored)
[FAIL] 2.1.10 Ensure that the --hostname-override argument is not set (Scored)
[FAIL] 2.1.11 Ensure that the --event-qps argument is set to 0 (Scored)
[PASS] 2.1.12 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Scored)
[PASS] 2.1.13 Ensure that the --cadvisor-port argument is set to 0 (Scored)
[FAIL] 2.1.14 Ensure that the RotateKubeletClientCertificate argument is set to true
[FAIL] 2.1.15 Ensure that the RotateKubeletServerCertificate argument is set to true
[INFO] 2.2 Configuration Files
[FAIL] 2.2.1 Ensure that the kubelet.conf file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.2 Ensure that the kubelet.conf file ownership is set to root:root (Scored)
[FAIL] 2.2.3 Ensure that the kubelet service file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.4 2.2.4 Ensure that the kubelet service file ownership is set to root:root (Scored)
[FAIL] 2.2.5 Ensure that the proxy kubeconfig file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.6 Ensure that the proxy kubeconfig file ownership is set to root:root (Scored)
[WARN] 2.2.7 Ensure that the certificate authorities file permissions are set to 644 or more restrictive (Scored)
[WARN] 2.2.8 Ensure that the client certificate authorities file ownership is set to root:root
```



# Kubesecc

- Helps you quantify risk for Kubernetes resources.
- Run against your K8s applications (deployments/pods/daemonsets etc)
- <https://kubesecc.io/> from controlplane
- Can be used standalone, or as a kubectl plugin (<https://github.com/stefanprodan/kubectl-kubesecc>)



# Kubesecc example

```
~$ kubectl -n kube-system plugin scan deployment/kubernetes-dashboard
scanning deployment kubernetes-dashboard
deployment/kubernetes-dashboard kubesecc.io score 3
```

-----

Advise

1. containers[].securityContext.runAsNonRoot == true

Force the running image to run as a non-root user to ensure least privilege

2. containers[].securityContext.capabilities.drop

Reducing kernel capabilities available to a container limits its attack surface

3. containers[].securityContext.readOnlyRootFilesystem == true

An immutable root filesystem can prevent malicious binaries being added to PATH and increase attack cost

4. containers[].securityContext.runAsUser > 10000

Run as a high-UID user to avoid conflicts with the host's user table

5. containers[].securityContext.capabilities.drop | index("ALL")

Drop all capabilities and add only those required to reduce syscall attack surface

```
~$ █
```

# Kubeaudit



- Opensourced from Shopify.
- Auditing your applications in your K8s cluster.
- <https://github.com/Shopify/kubeaudit>
- Little more targeted than Kubesec.

kubeaudit is a program that will help you audit your Kubernetes clusters. Specify -l to run kubeaudit using ~/.kube/config otherwise it will attempt to create an in-cluster client.

#patcheswelcome

#### Usage:

kubeaudit [command]

#### Available Commands:

allowpe	Audit containers that allow privilege escalation
caps	Audit container for capabilities
help	Help about any <a href="#">command</a>
image	Audit container images
nonroot	Audit containers running as root
np	Audit namespace network policies
priv	Audit containers running as root
rootfs	Audit containers with read only root filesystems
sat	Audit automountServiceAccountToken = true pods against an empty (default) service account
version	Print the version number of kubeaudit

#### Flags:

-a, --allPods	Audit againsts pods in all the phases (default Running Phase)
-h, --help	help for kubeaudit
-j, --json	Enable json logging
-c, --kubeconfig string	config file (default is \$HOME/.kube/config)
-l, --local	Local mode, uses ~/.kube/config as configuration
-f, --manifest string	yaml configuration to audit
-v, --verbose string	Set the debug level (default "INFO")

Use "kubeaudit [command] --help" for more information about a command.



# Kubeaudit example

```
~$ /Users/karthik/Downloads/kubeaudit_0.2.0_darwin_amd64/kubeaudit allowpe -c /Users/karthik/.kube/config
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=dotnetworld Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=javademo Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-alertmanager Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-kube-state-metrics Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-pushgateway Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-server Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=wishlist-deployment Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=contour Namespace=heptio-contour
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=kube-dns Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=kube-dns-autoscaler Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=kubernetes-dashboard Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=oci-volume-provisioner Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=tiller-deploy Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=daemonSet Name=prom-demo-prometheus-node-exporter Namespace=default
ERRO[0003] AllowPrivilegeEscalation not set which allows privilege escalation, please set to false KubeType=daemonSet Name=kube-flannel-ds Namespace=kube-system
ERRO[0003] AllowPrivilegeEscalation not set which allows privilege escalation, please set to false KubeType=daemonSet Name=kube-proxy Namespace=kube-system
```



A close-up shot of Gene Wilder as Charlie Bucket. He is wearing a brown top hat, a purple suit jacket, a white shirt, and a large tan bow tie. He has a wide-eyed, excited expression and is resting his chin on his right hand. The background is slightly out of focus, showing a yellow vertical pole and a blue wall with a white circular object. The text "So much time and so little to do." is overlaid at the bottom in yellow.

“So much time and so little to do.”

# Couple more resources to look at:

- 11 ways not to get hacked:  
<https://kubernetes.io/blog/2018/07/18/11-ways-not-to-get-hacked>
- K8s security (from Image Hygiene to Network Policy):  
<https://speakerdeck.com/mhausenblas/kubernetes-security-from-image-hygiene-to-network-policies>

KEEP CALM  
AND  
KUBE ON  
*@iteration1*



# Thank You All Day DevOps Sponsors

## Platinum Sponsors



## Gold Sponsors



## Media Sponsors







Meet me in the Slack channel for Q&A

[bit.ly/addo-slack](https://bit.ly/addo-slack)