



redhat.<sup>®</sup>

## Kubernetes Security

### From Image Hygiene to Network Policies

Michael Hausenblas **@mhausenblas**

Developer Advocate, Red Hat

2018-05-29, DevOpsCon, Berlin

# \$ whois mhausenblas

- Developer Advocate @ Red Hat (Go, Kubernetes, OpenShift)
- Developer Advocate @ Mesosphere (Mesos, DC/OS, Kubernetes)
- Chief Data Engineer @ MapR (HDFS, HBase, Drill, etc.)
- Applied research (4y in Ireland, 7y in Austria)
- Nowadays mainly developing tools in Go (Python, Node, Java, C++)
- Kinda developer turned ops (aka appops)



admin

developer

architect



SRE

infosec

PM

# Overview and terminology

container images

running container

authn & authz

communication

apps

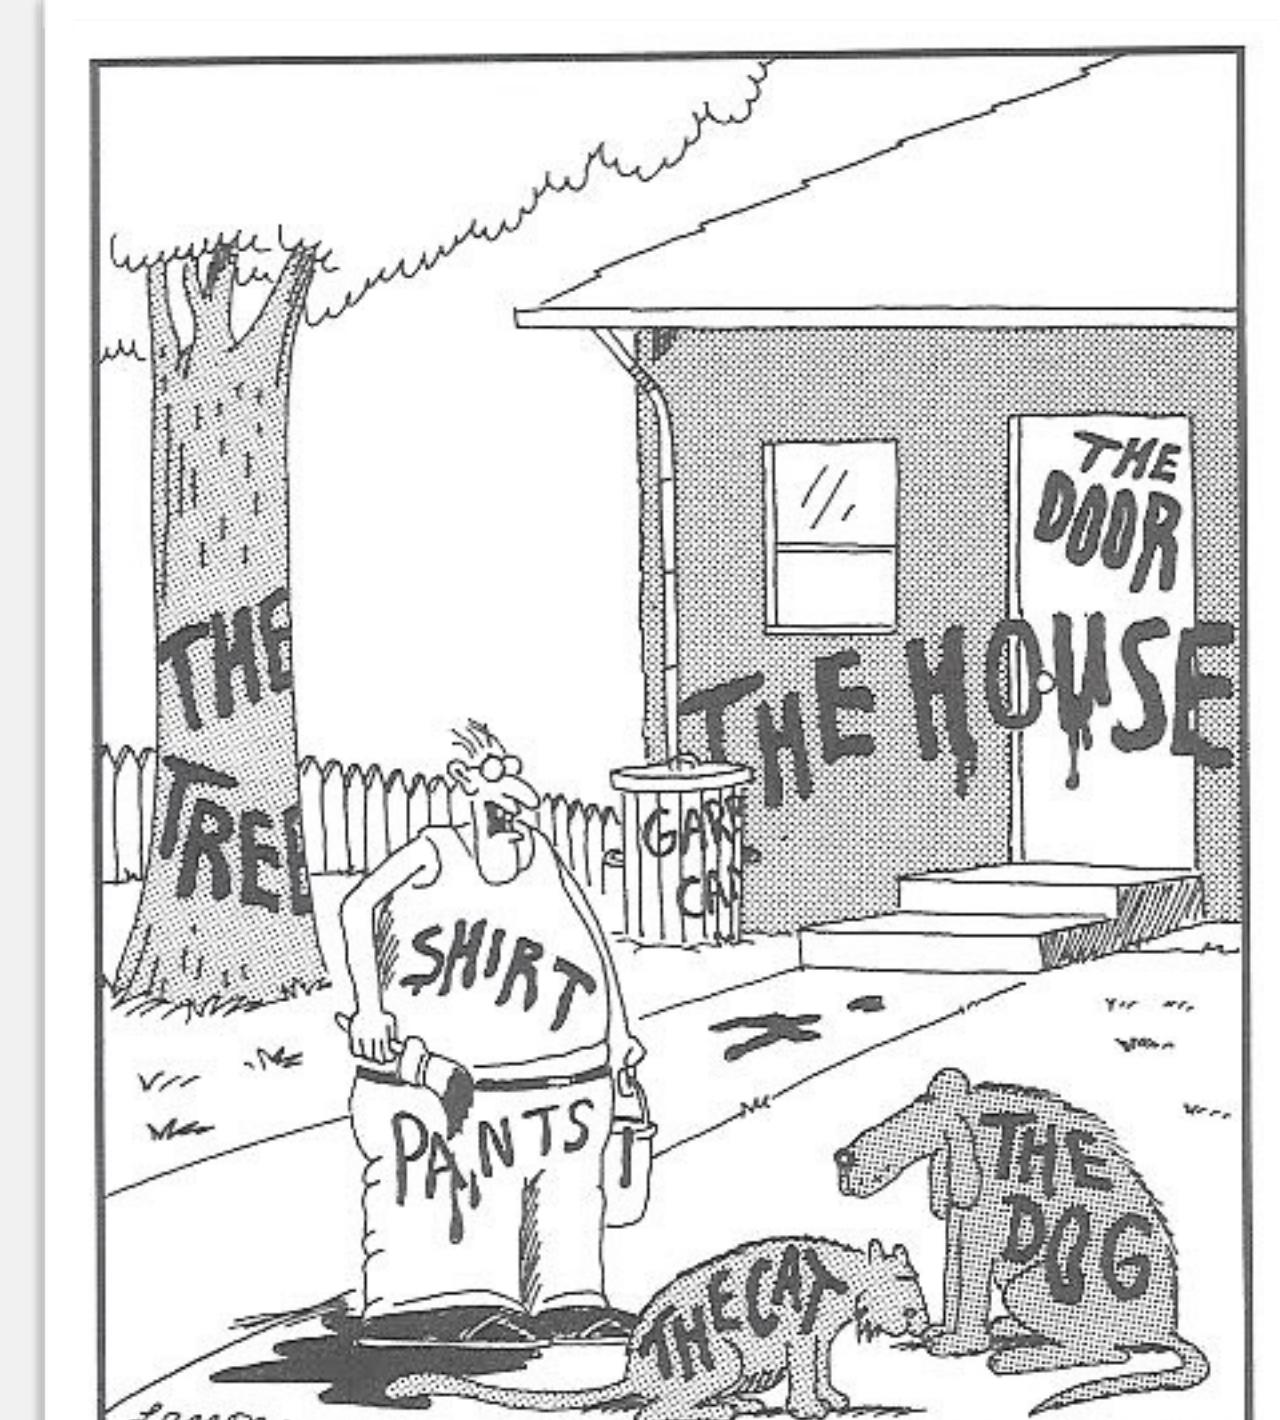
control plane

An expert is a person who  
has made all the mistakes  
which can be made, in a  
narrow field.

Updated from Niels Bohr, Danish physicist, Nobel Prize for Physics in 1922

# Terminology

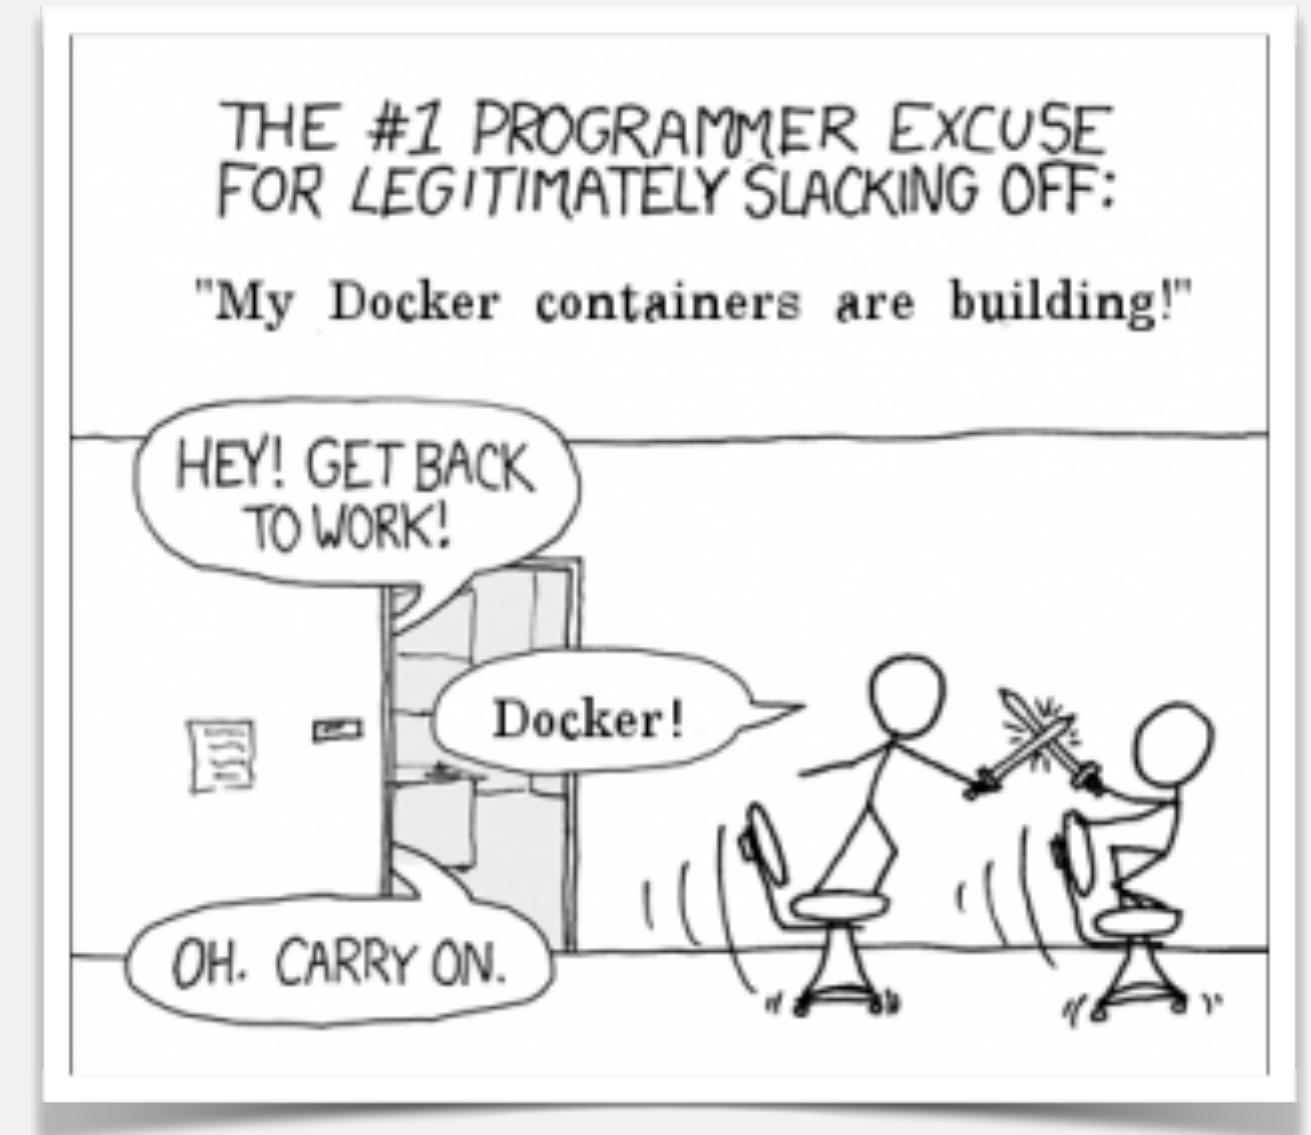
- build-time vs run-time
- immutability & automation
- responsibilities
- moving parts



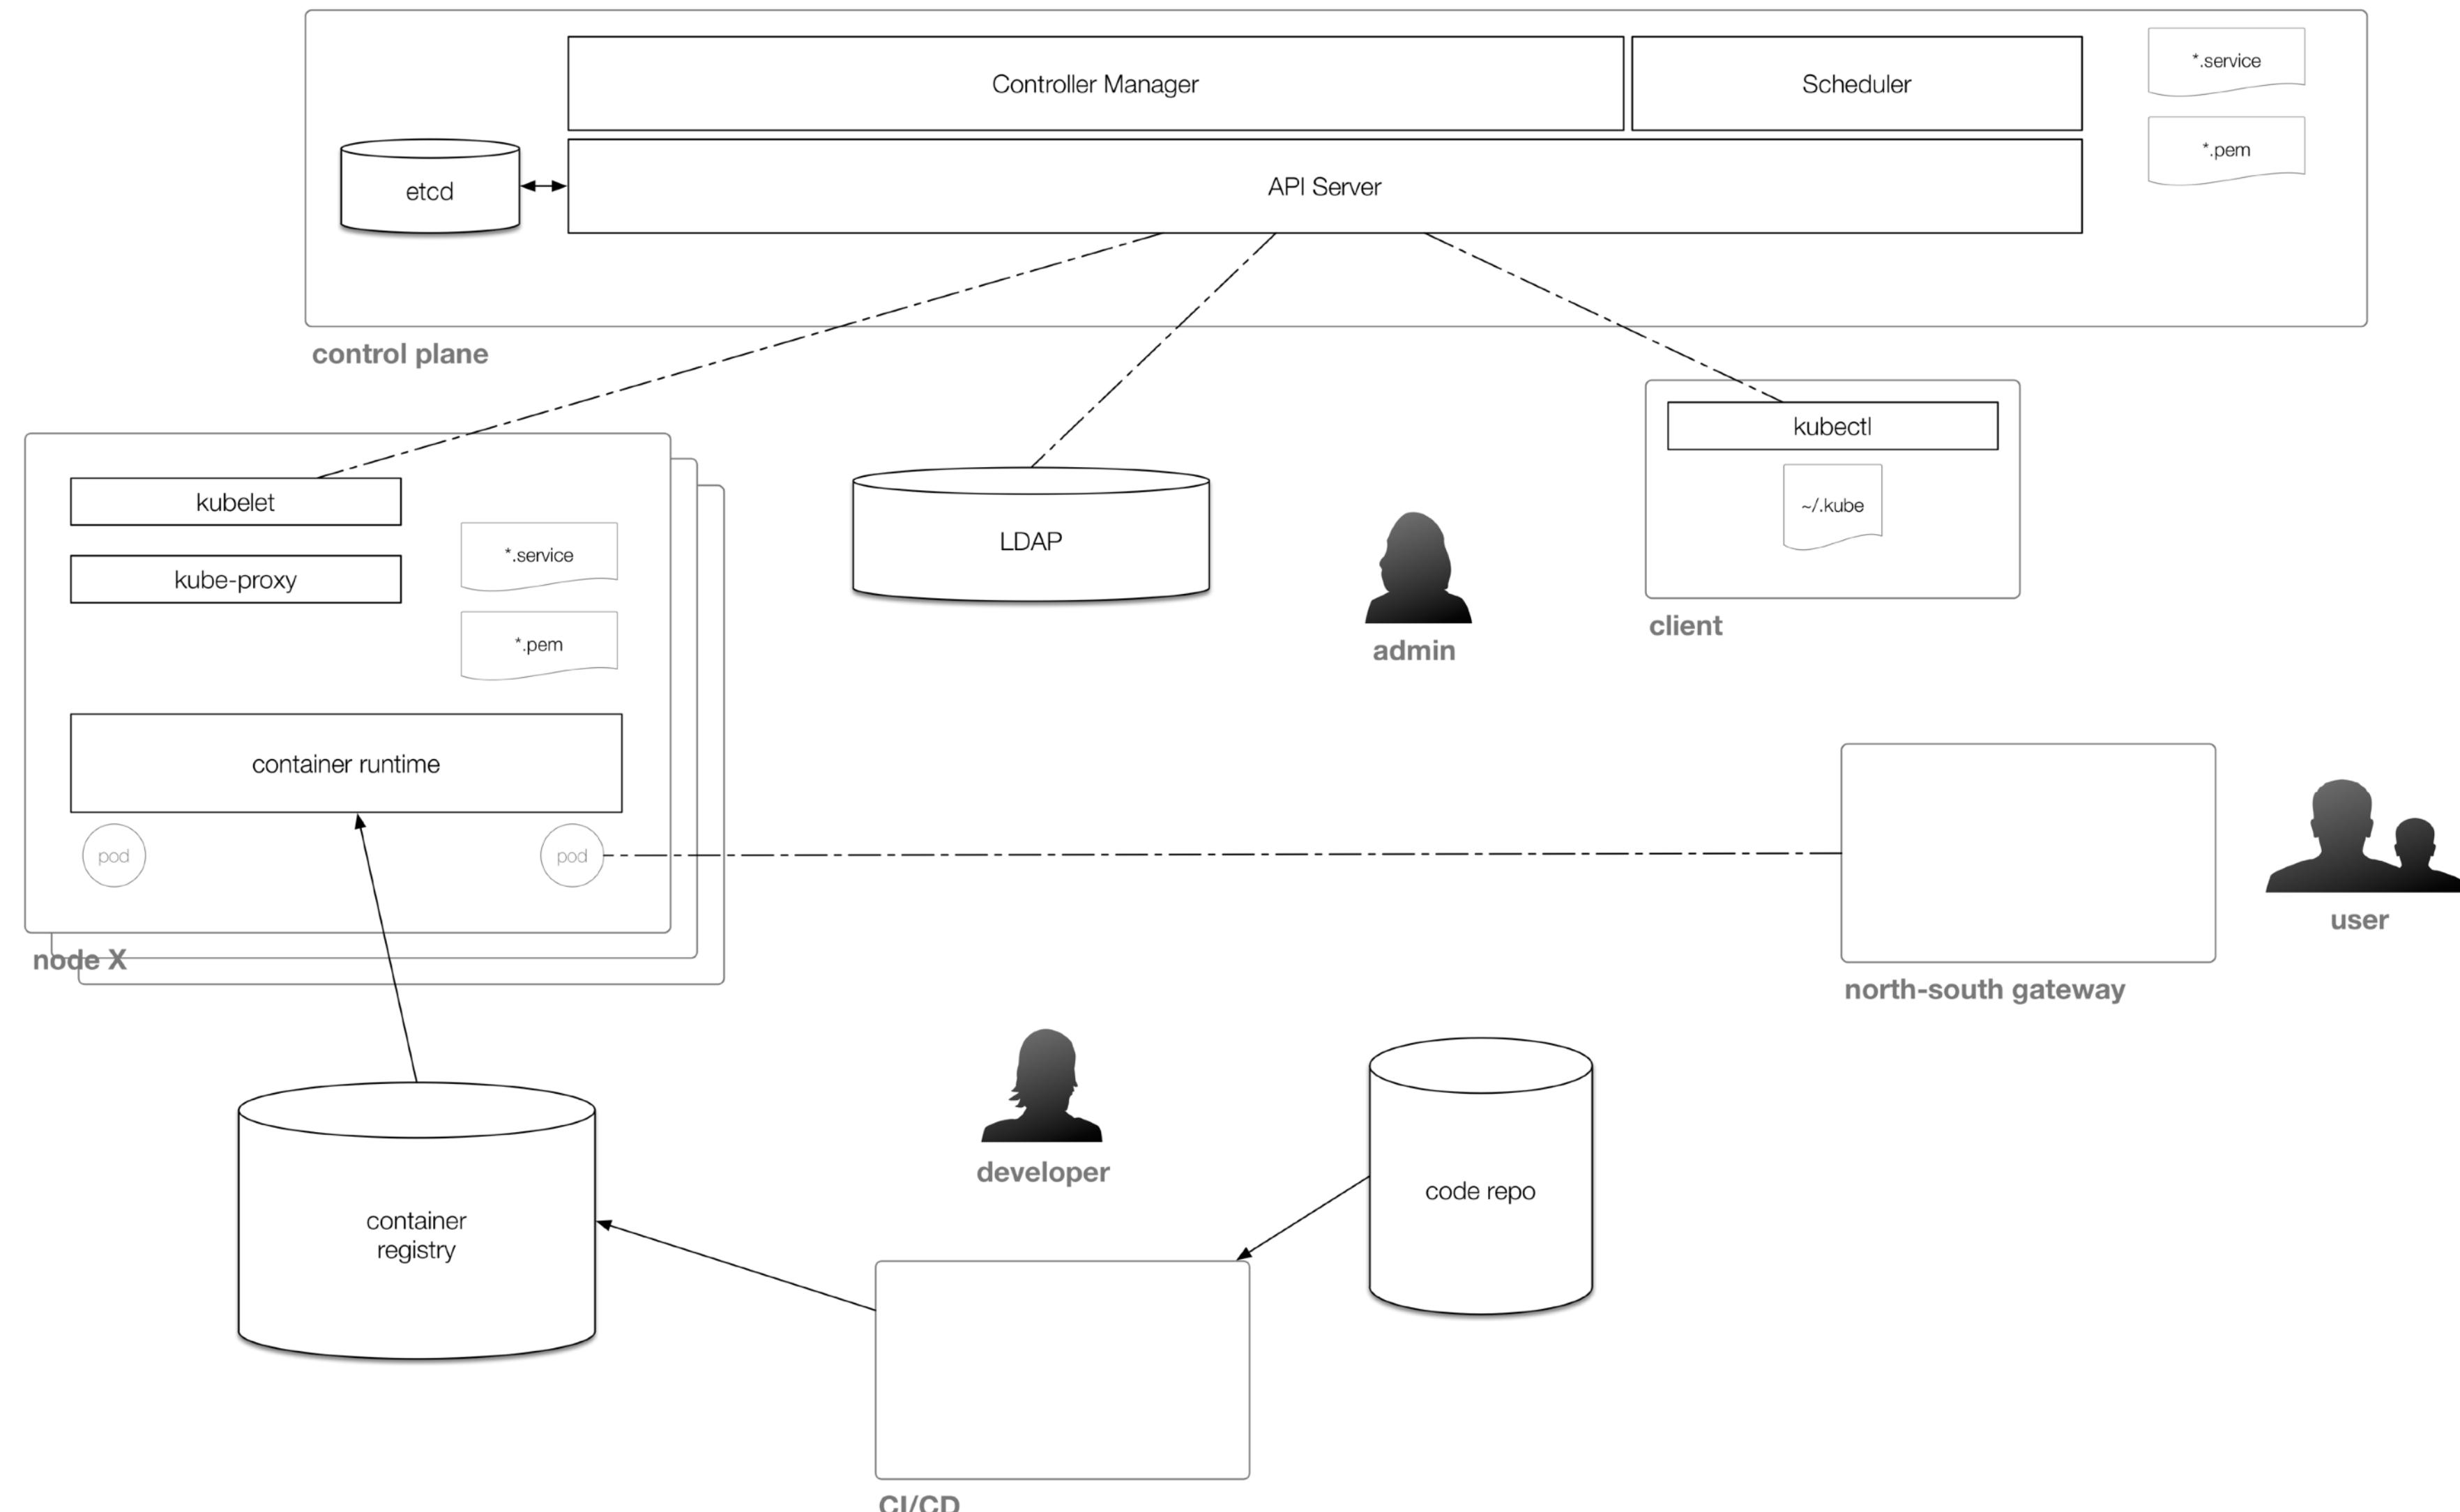
"Now! ... That should clear up  
a few things around here!"

# Responsibilities

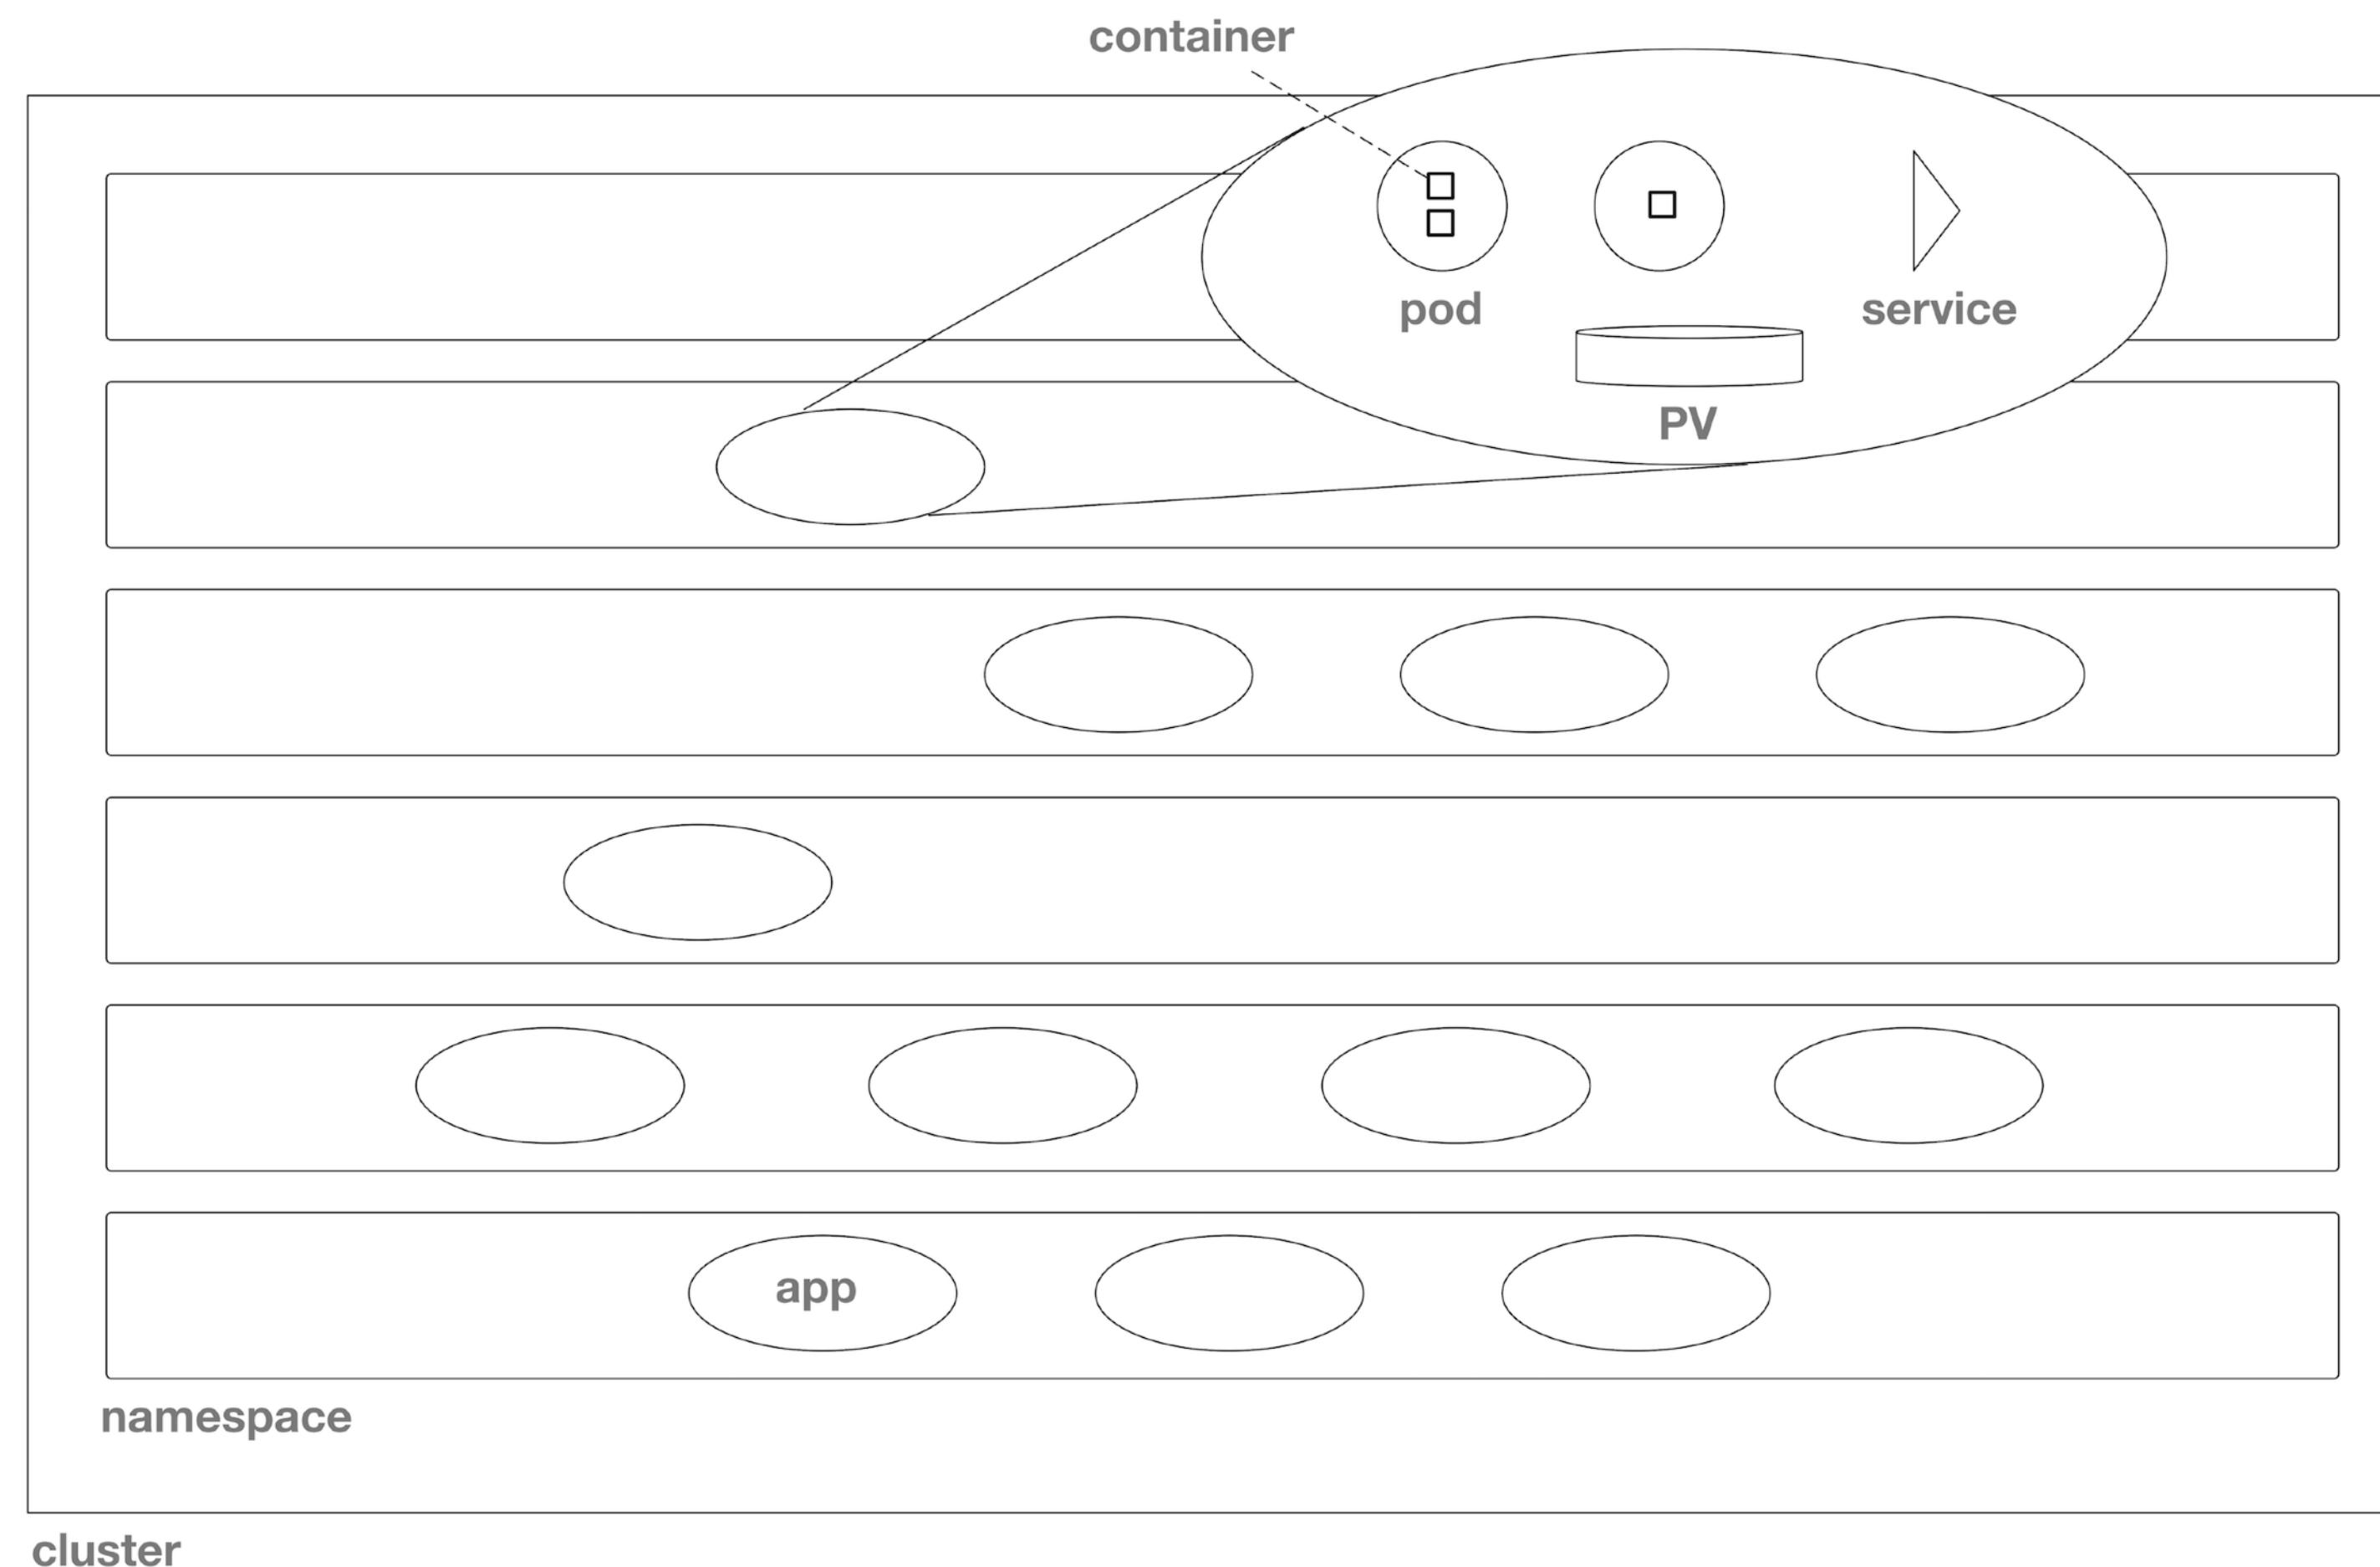
- infrastructure admin
- namespace admin
- developer



# Moving parts—physical view



# Moving parts—logical view



# Building container images

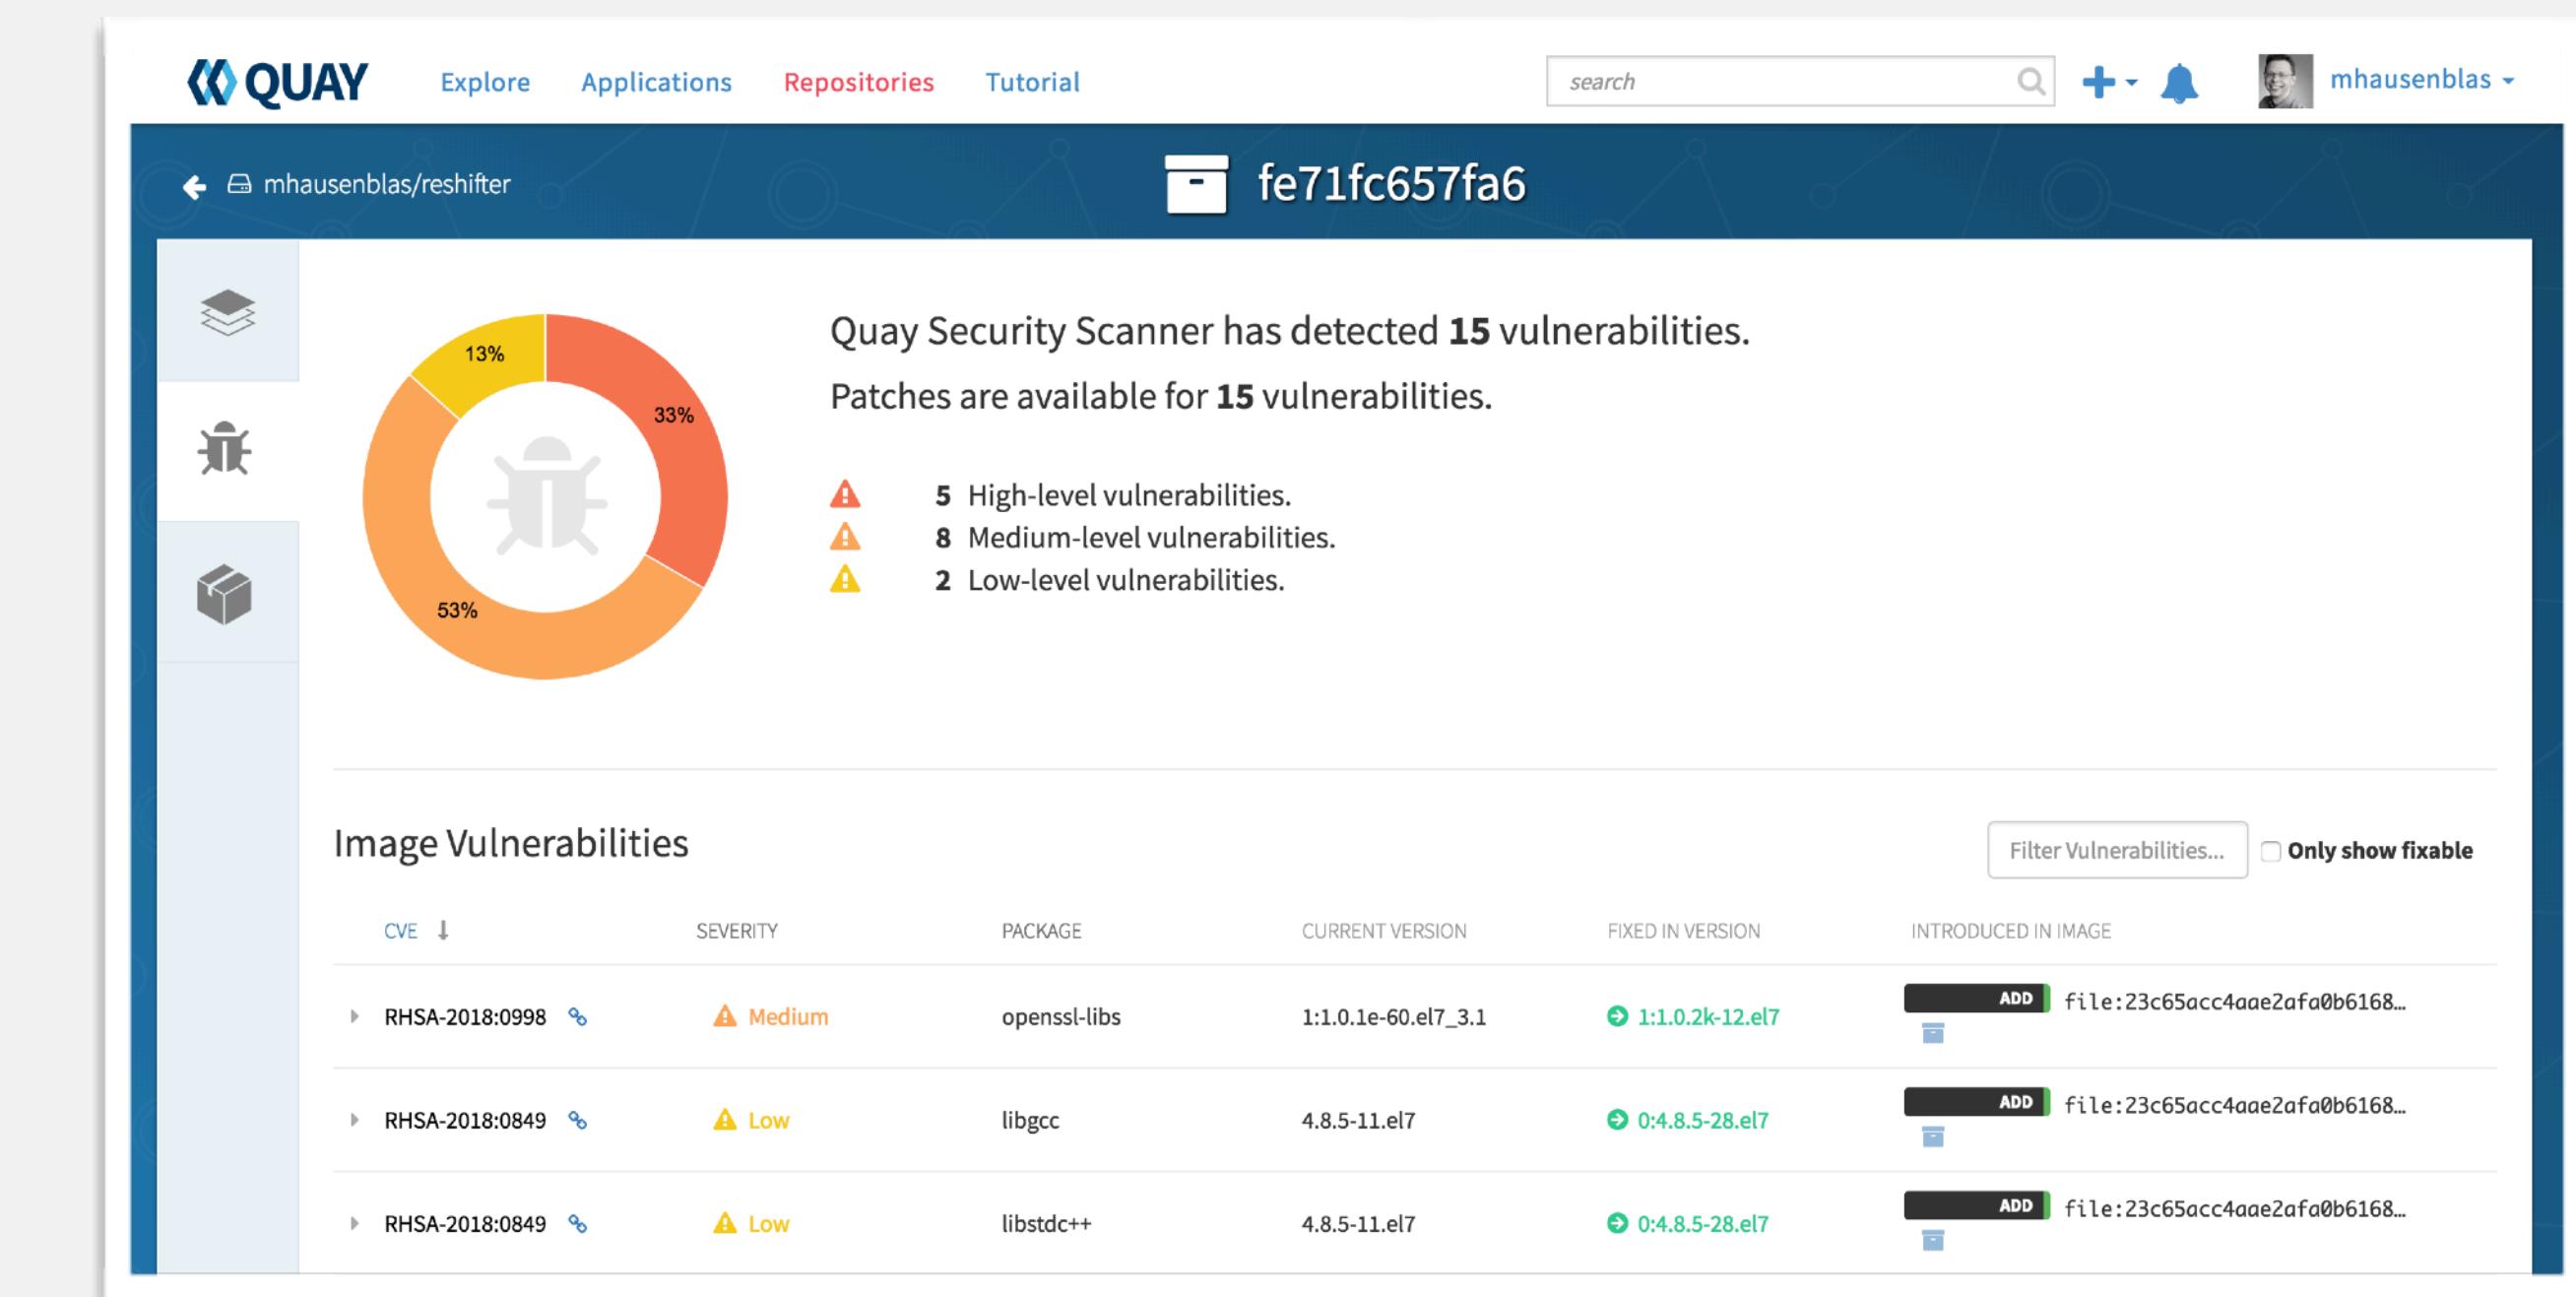
# What?

Build container images that ...

- have a small attack surface
- are checked for vulnerabilities
- are reproducible

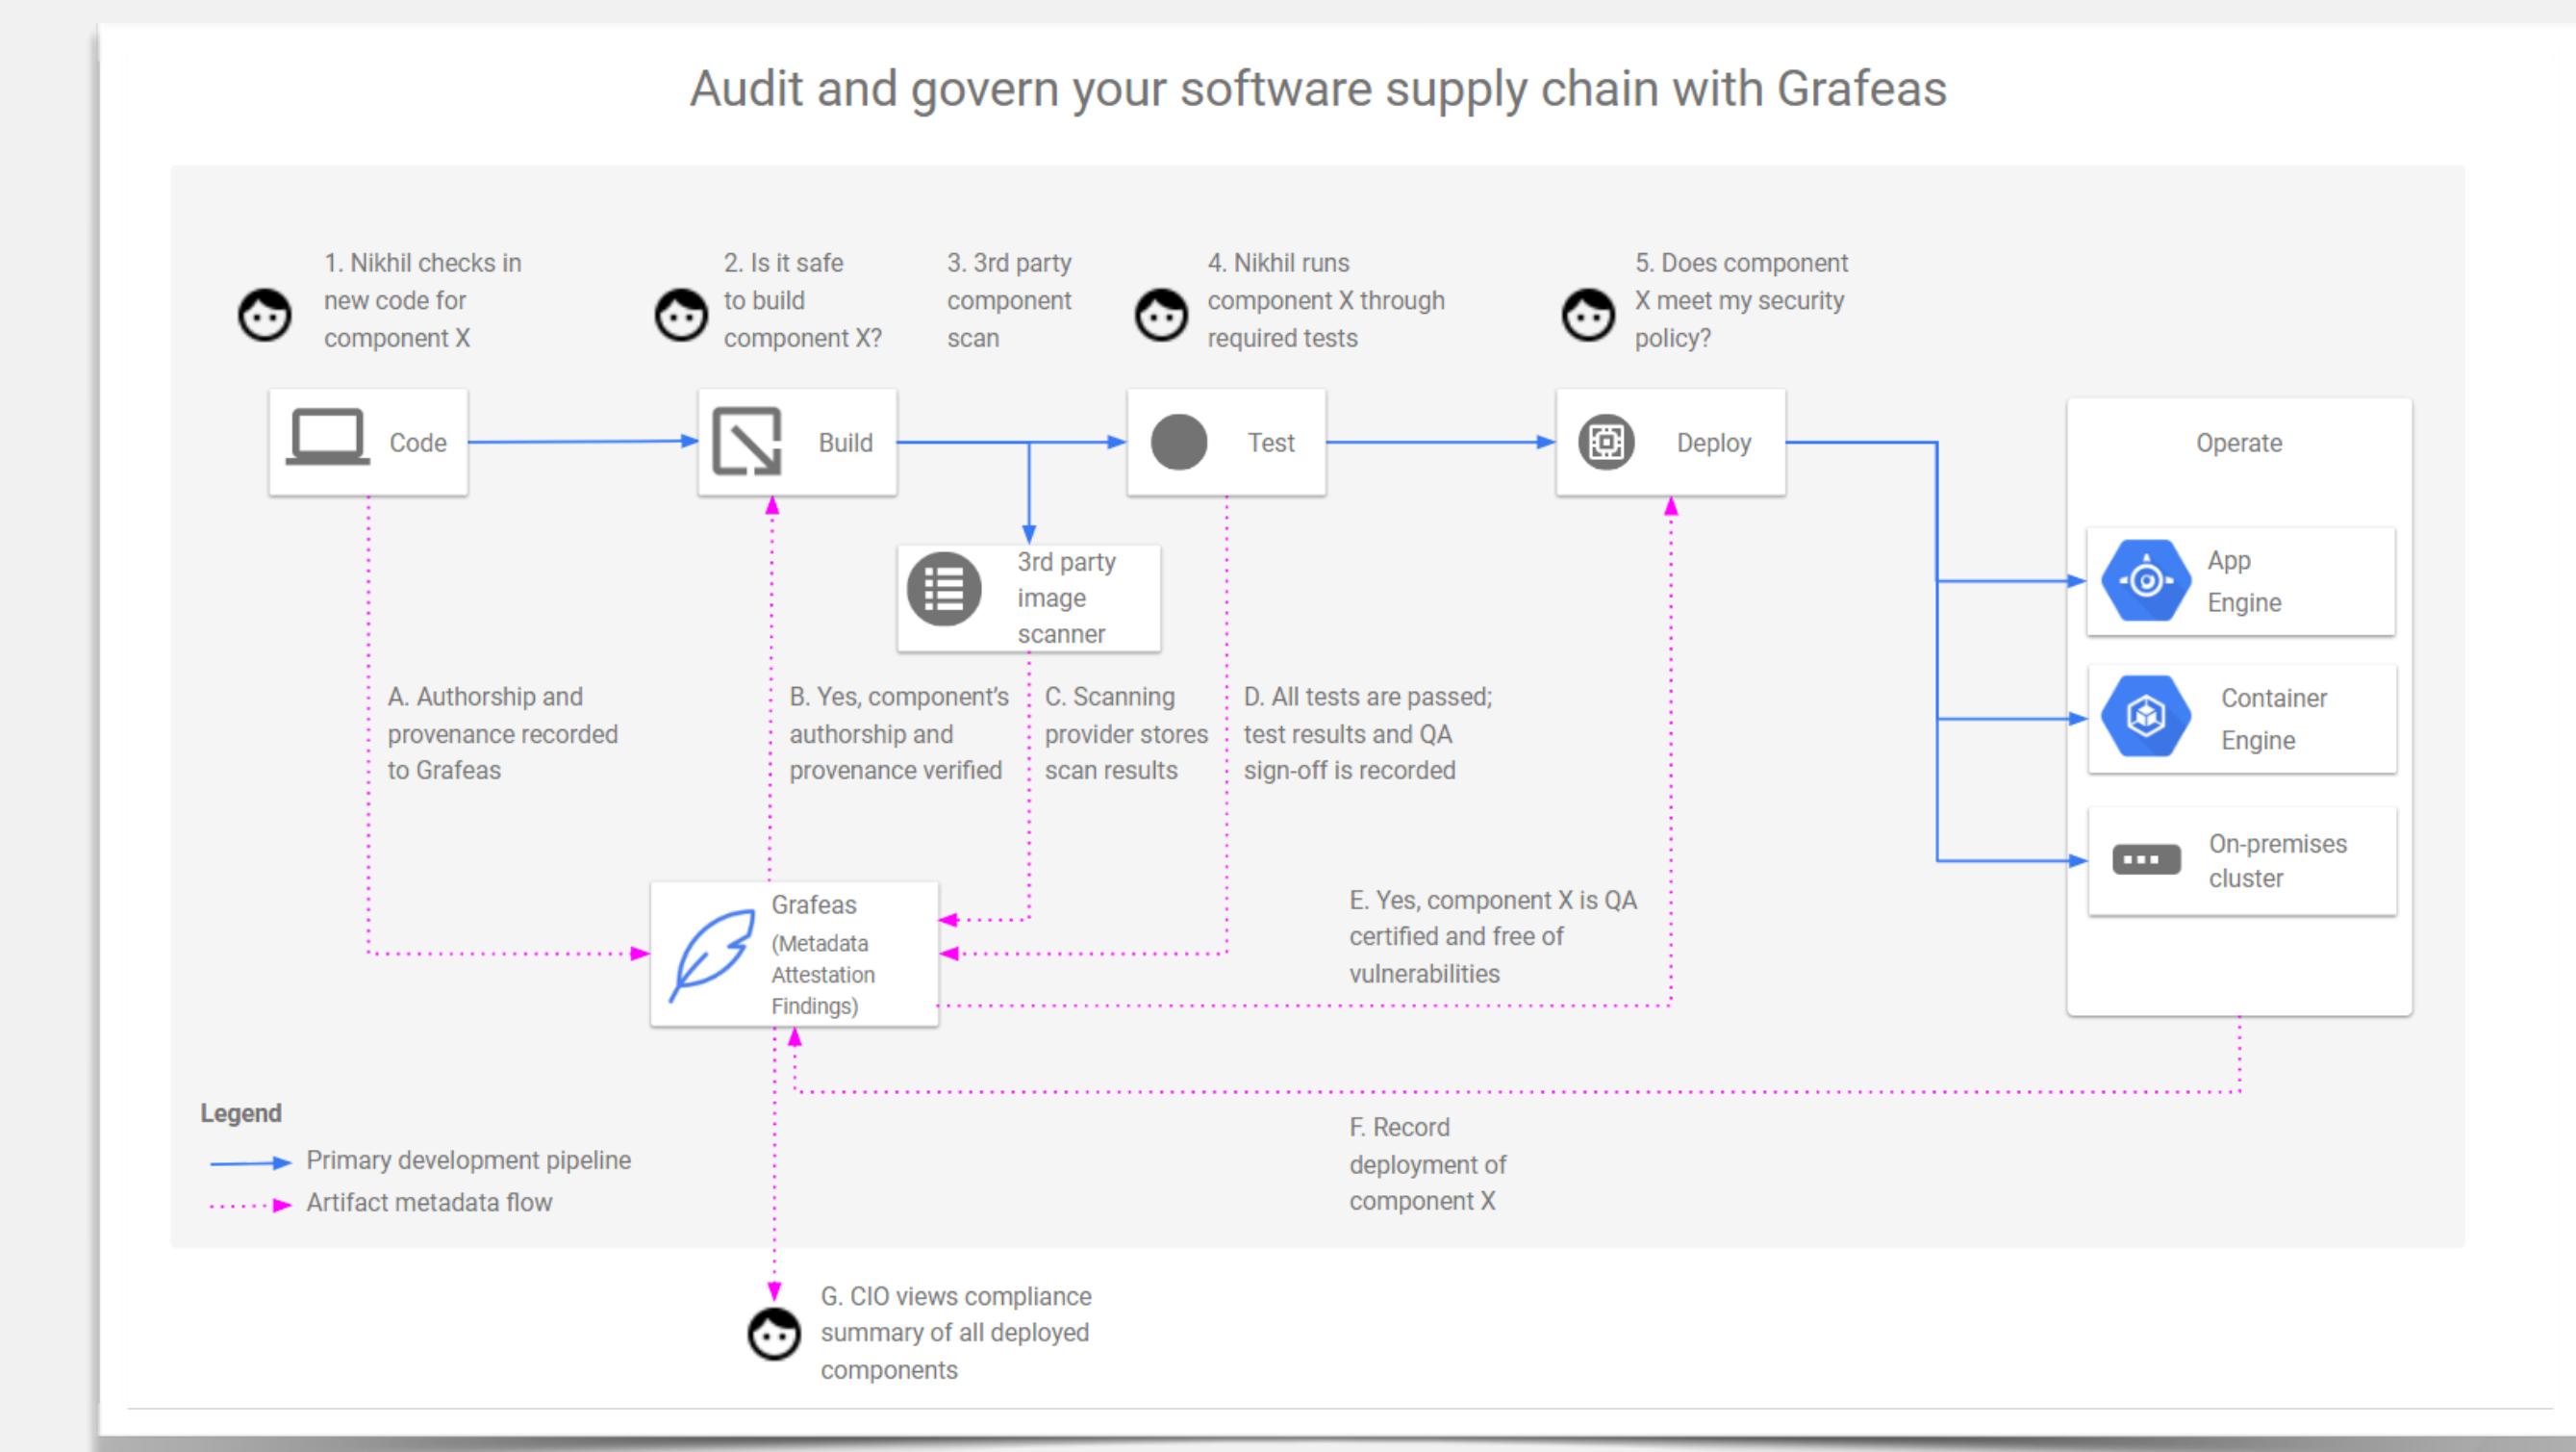
# Good practices

- use trusted base images
- define user in image
- perform automated CVE scans
- use private registries
- pin dependencies  
(reproducible builds)



# Grafeas

- supply chain management
- structured metadata API for annotating artefacts + enforcement (Kritis)
- backed by Google, JFrog, Red Hat, IBM, Black Duck, Twistlock, Aqua Security



Source: Introducing Grafeas: An open-source API to audit and govern your software supply chain

<https://grafeas.io>

# Tooling

- <https://docs.docker.com/docker-cloud/builds/image-scan/>
- <https://github.com/coreos/clair>
- <https://www.open-scap.org/tools/>
- <https://www.aquasec.com/use-cases/continuous-image-assurance/>
- <https://neuvector.com/container-compliance-auditing-solutions/>
- <https://github.com/theupdateframework/notary>
- <https://github.com/in-toto>

# Further reading

- Establishing Image Provenance and Security in Kubernetes
- Image Management & Mutability in Docker and Kubernetes
- Container security considerations in a Kubernetes deployment
- Building Container Images Securely on Kubernetes
- The OpenShift Build Process
- Introducing Grafeas: An open-source API to audit and govern your software supply chain

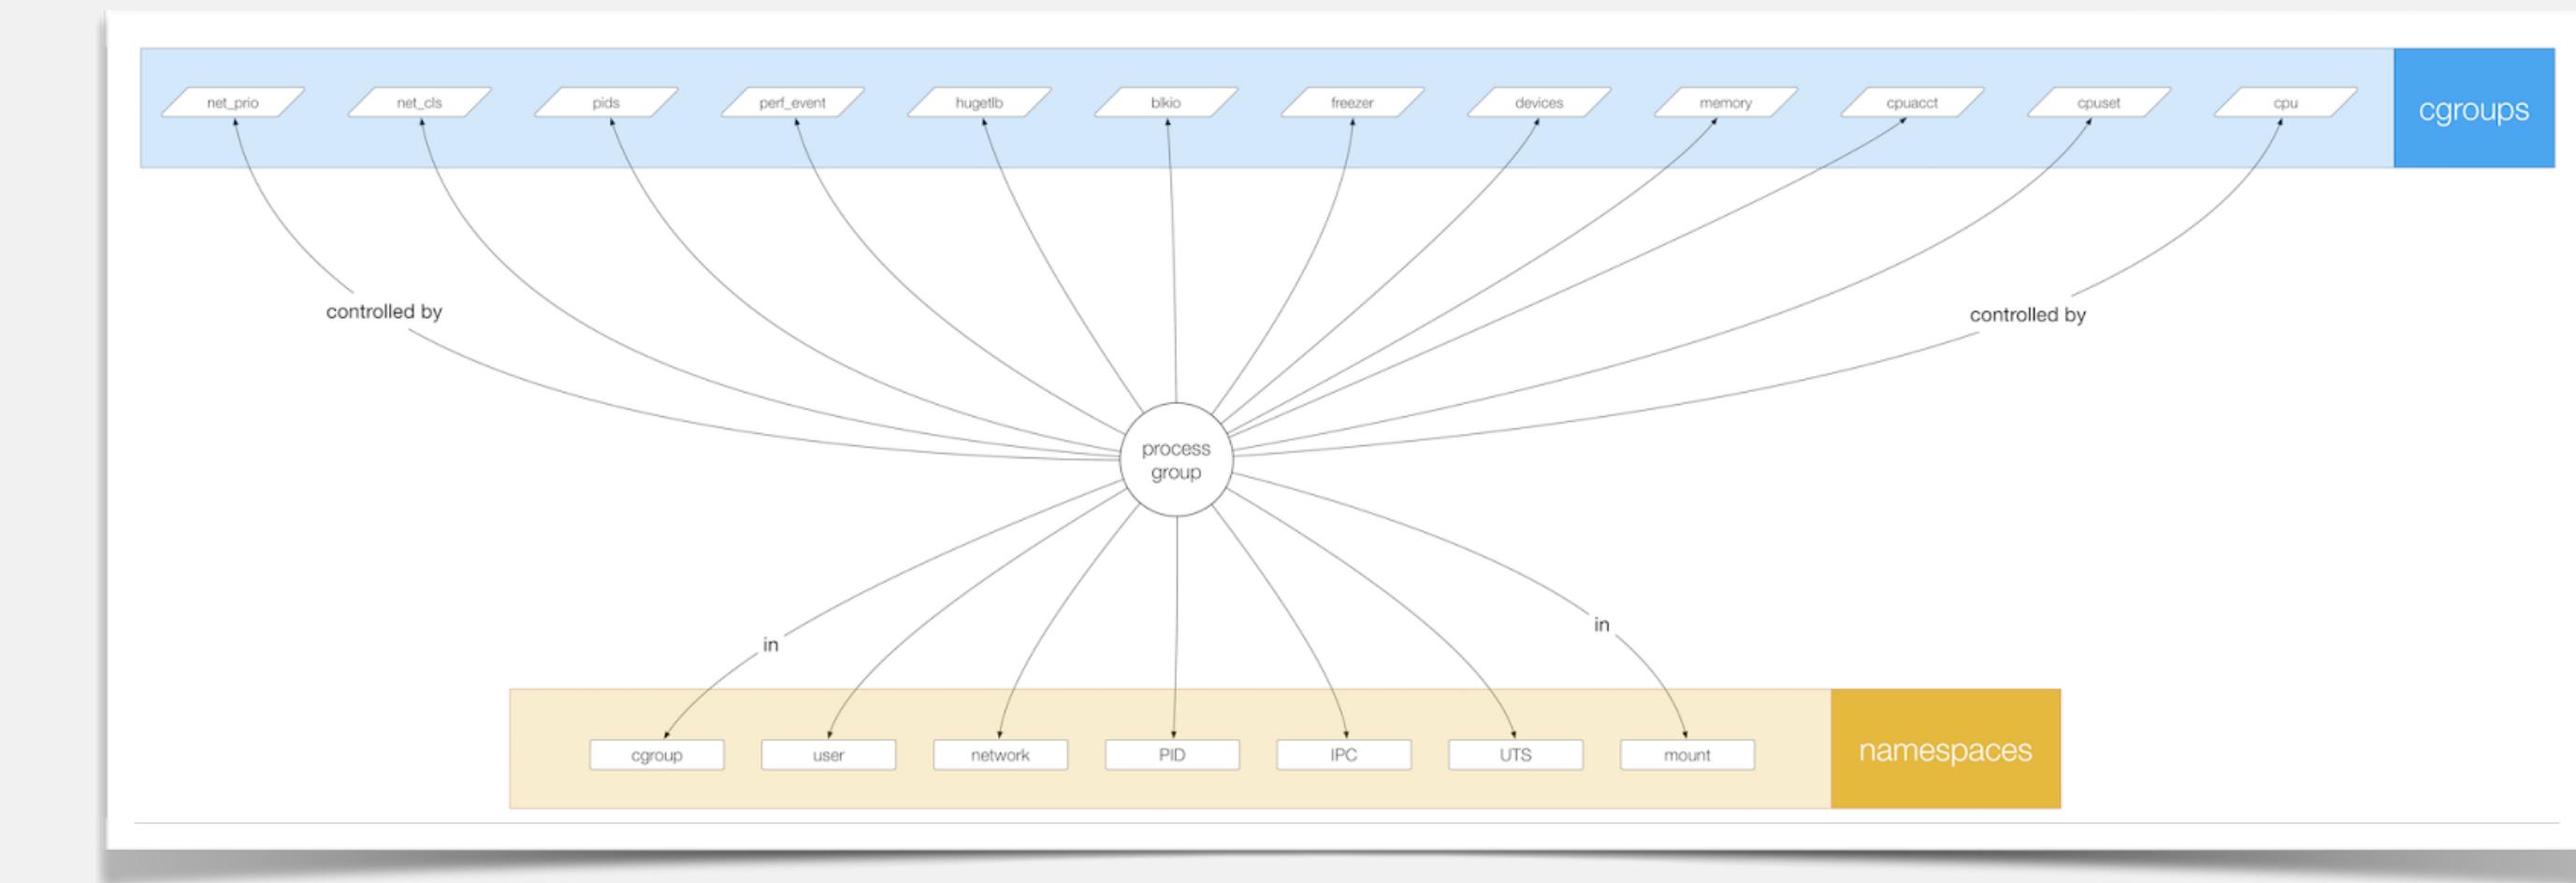
# Running containers

# What?

- Run containers ...
  - with known origin (image & registry)
  - use least privilege to carry out the task
  - do only minimal host mounts necessary

# Good practices

- verify the defaults
- don't run as root
- use security context & policies
- use security benchmarks



<http://containerz.info>

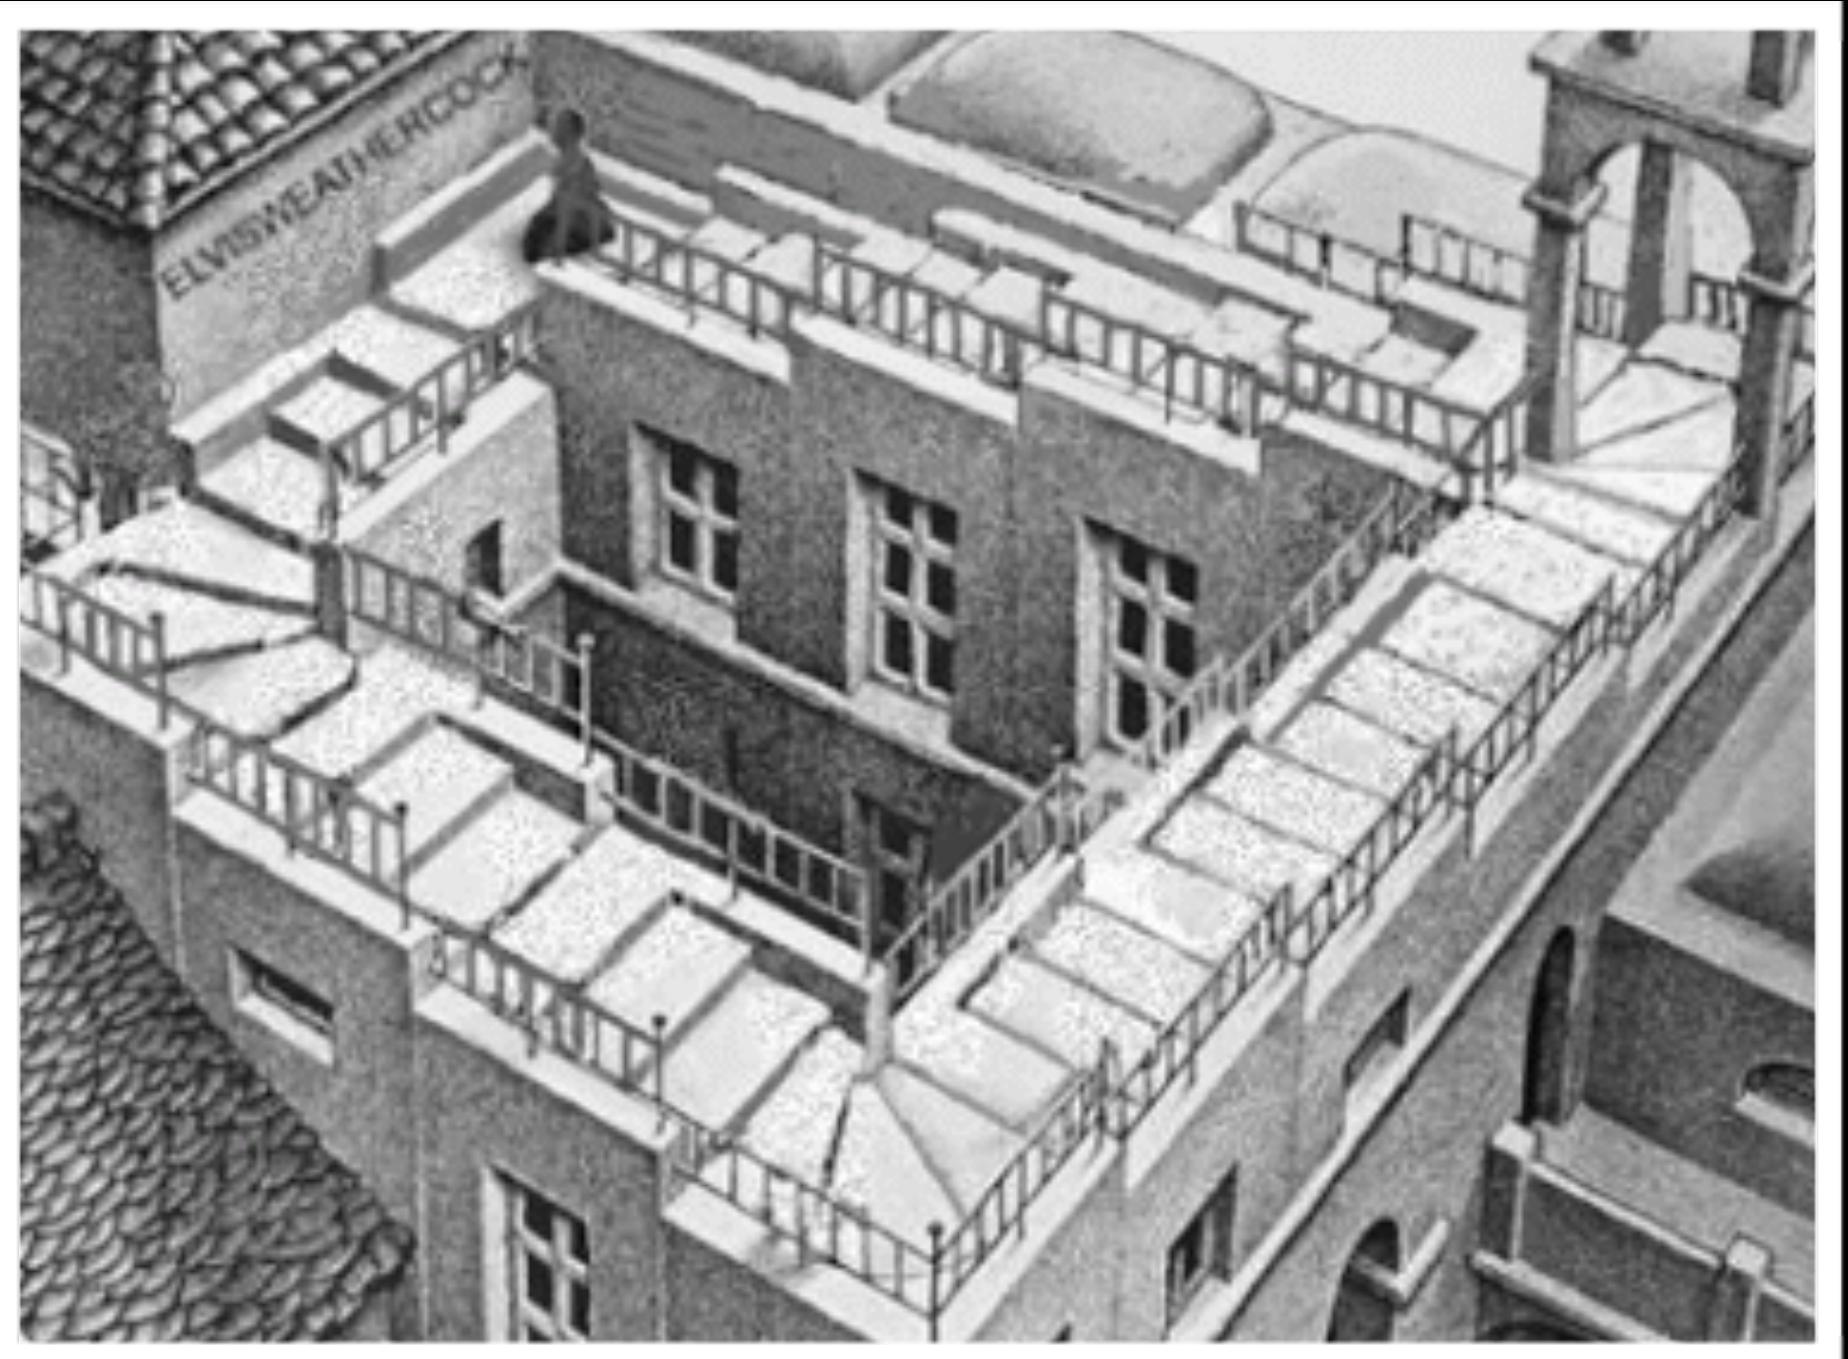
# Pod security context and policies

- security context
  - pod or container level
  - defines privilege and access control settings
  - seccomp, capabilities, SELinux, AppArmor
- security policies
  - cluster resource
  - via admission controller

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  securityContext:
    runAsUser: 0
    fsGroup: 2000
    volumes:
      - name: default-token-4q2zr
        emptyDir: {}
    hostNetwork: true
    hostPorts:
      - min: 0
        max: 65535
    hostIPC: true
    hostPID: true
    runAsUser:
      rule: 'RunAsAny'
    seLinux:
      rule: 'RunAsAny'
    supplementalGroups:
      rule: 'RunAsAny'
    fsGroup:
      rule: 'RunAsAny'

apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: privileged
  annotations:
    seccomp.security.alpha.kubernetes.io/allowedProfileNames: '*'
spec:
  privileged: true
  allowPrivilegeEscalation: true
  allowedCapabilities:
    - '*'
  volumes:
    - '*'
  hostNetwork: true
  hostPorts:
    - min: 0
      max: 65535
  hostIPC: true
  hostPID: true
  runAsUser:
    rule: 'RunAsAny'
  seLinux:
    rule: 'RunAsAny'
  supplementalGroups:
    rule: 'RunAsAny'
  fsGroup:
    rule: 'RunAsAny'
```

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/>



... demo time!

# Tooling

- <https://github.com/aquasecurity/kube-bench>
- <https://github.com/docker/docker-bench-security>
- <https://sysdig.comopensource/falco/>
- <https://kubesecc.io/>
- <https://www.twistlock.com/>

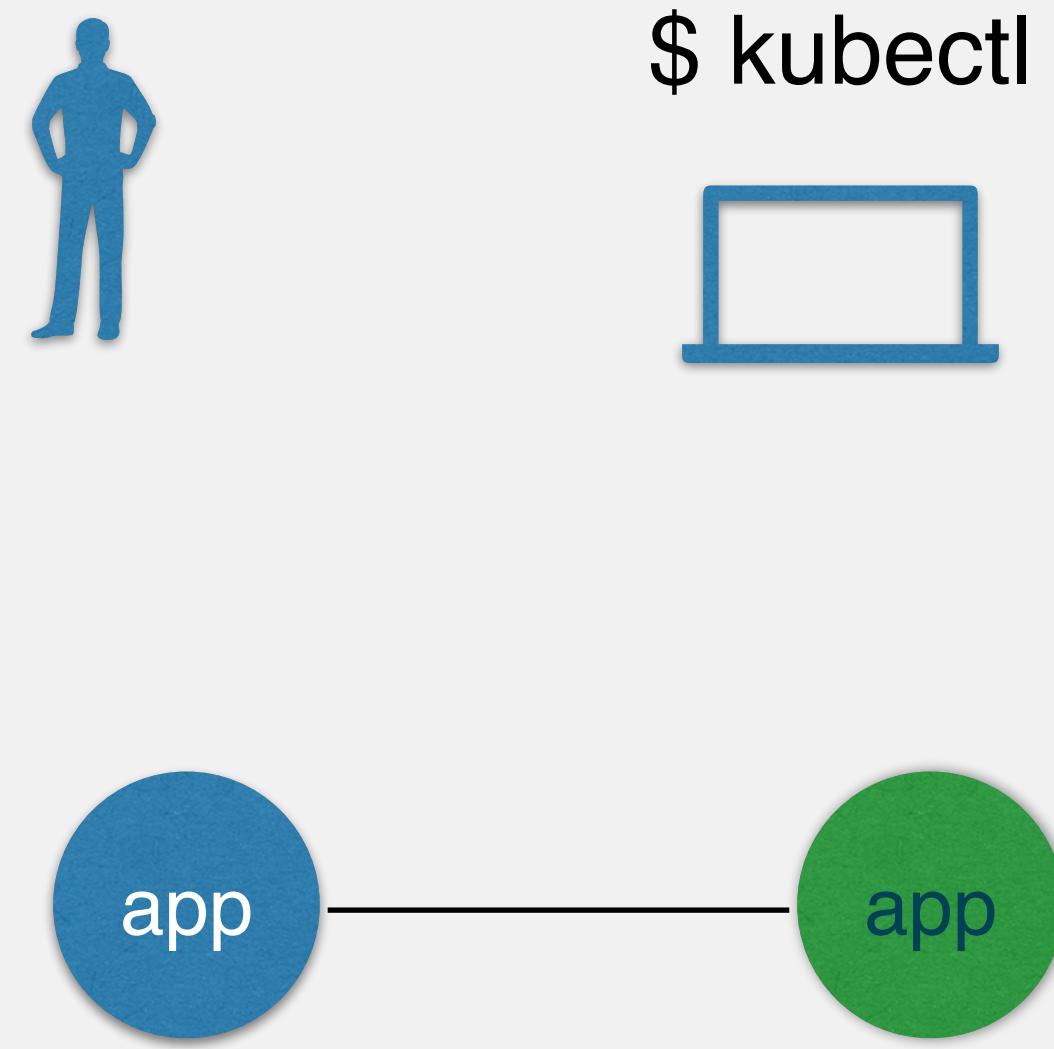
# Further reading

- Just say no to root (in containers)
- Exploring Container Mechanisms Through the Story of a Syscall ([slides](#) | [video](#))
- Improving your Kubernetes Workload Security
- Container Isolation at Scale (Introducing gVisor) ([slides](#) | [video](#))

# Authentication & Authorization

# Identity

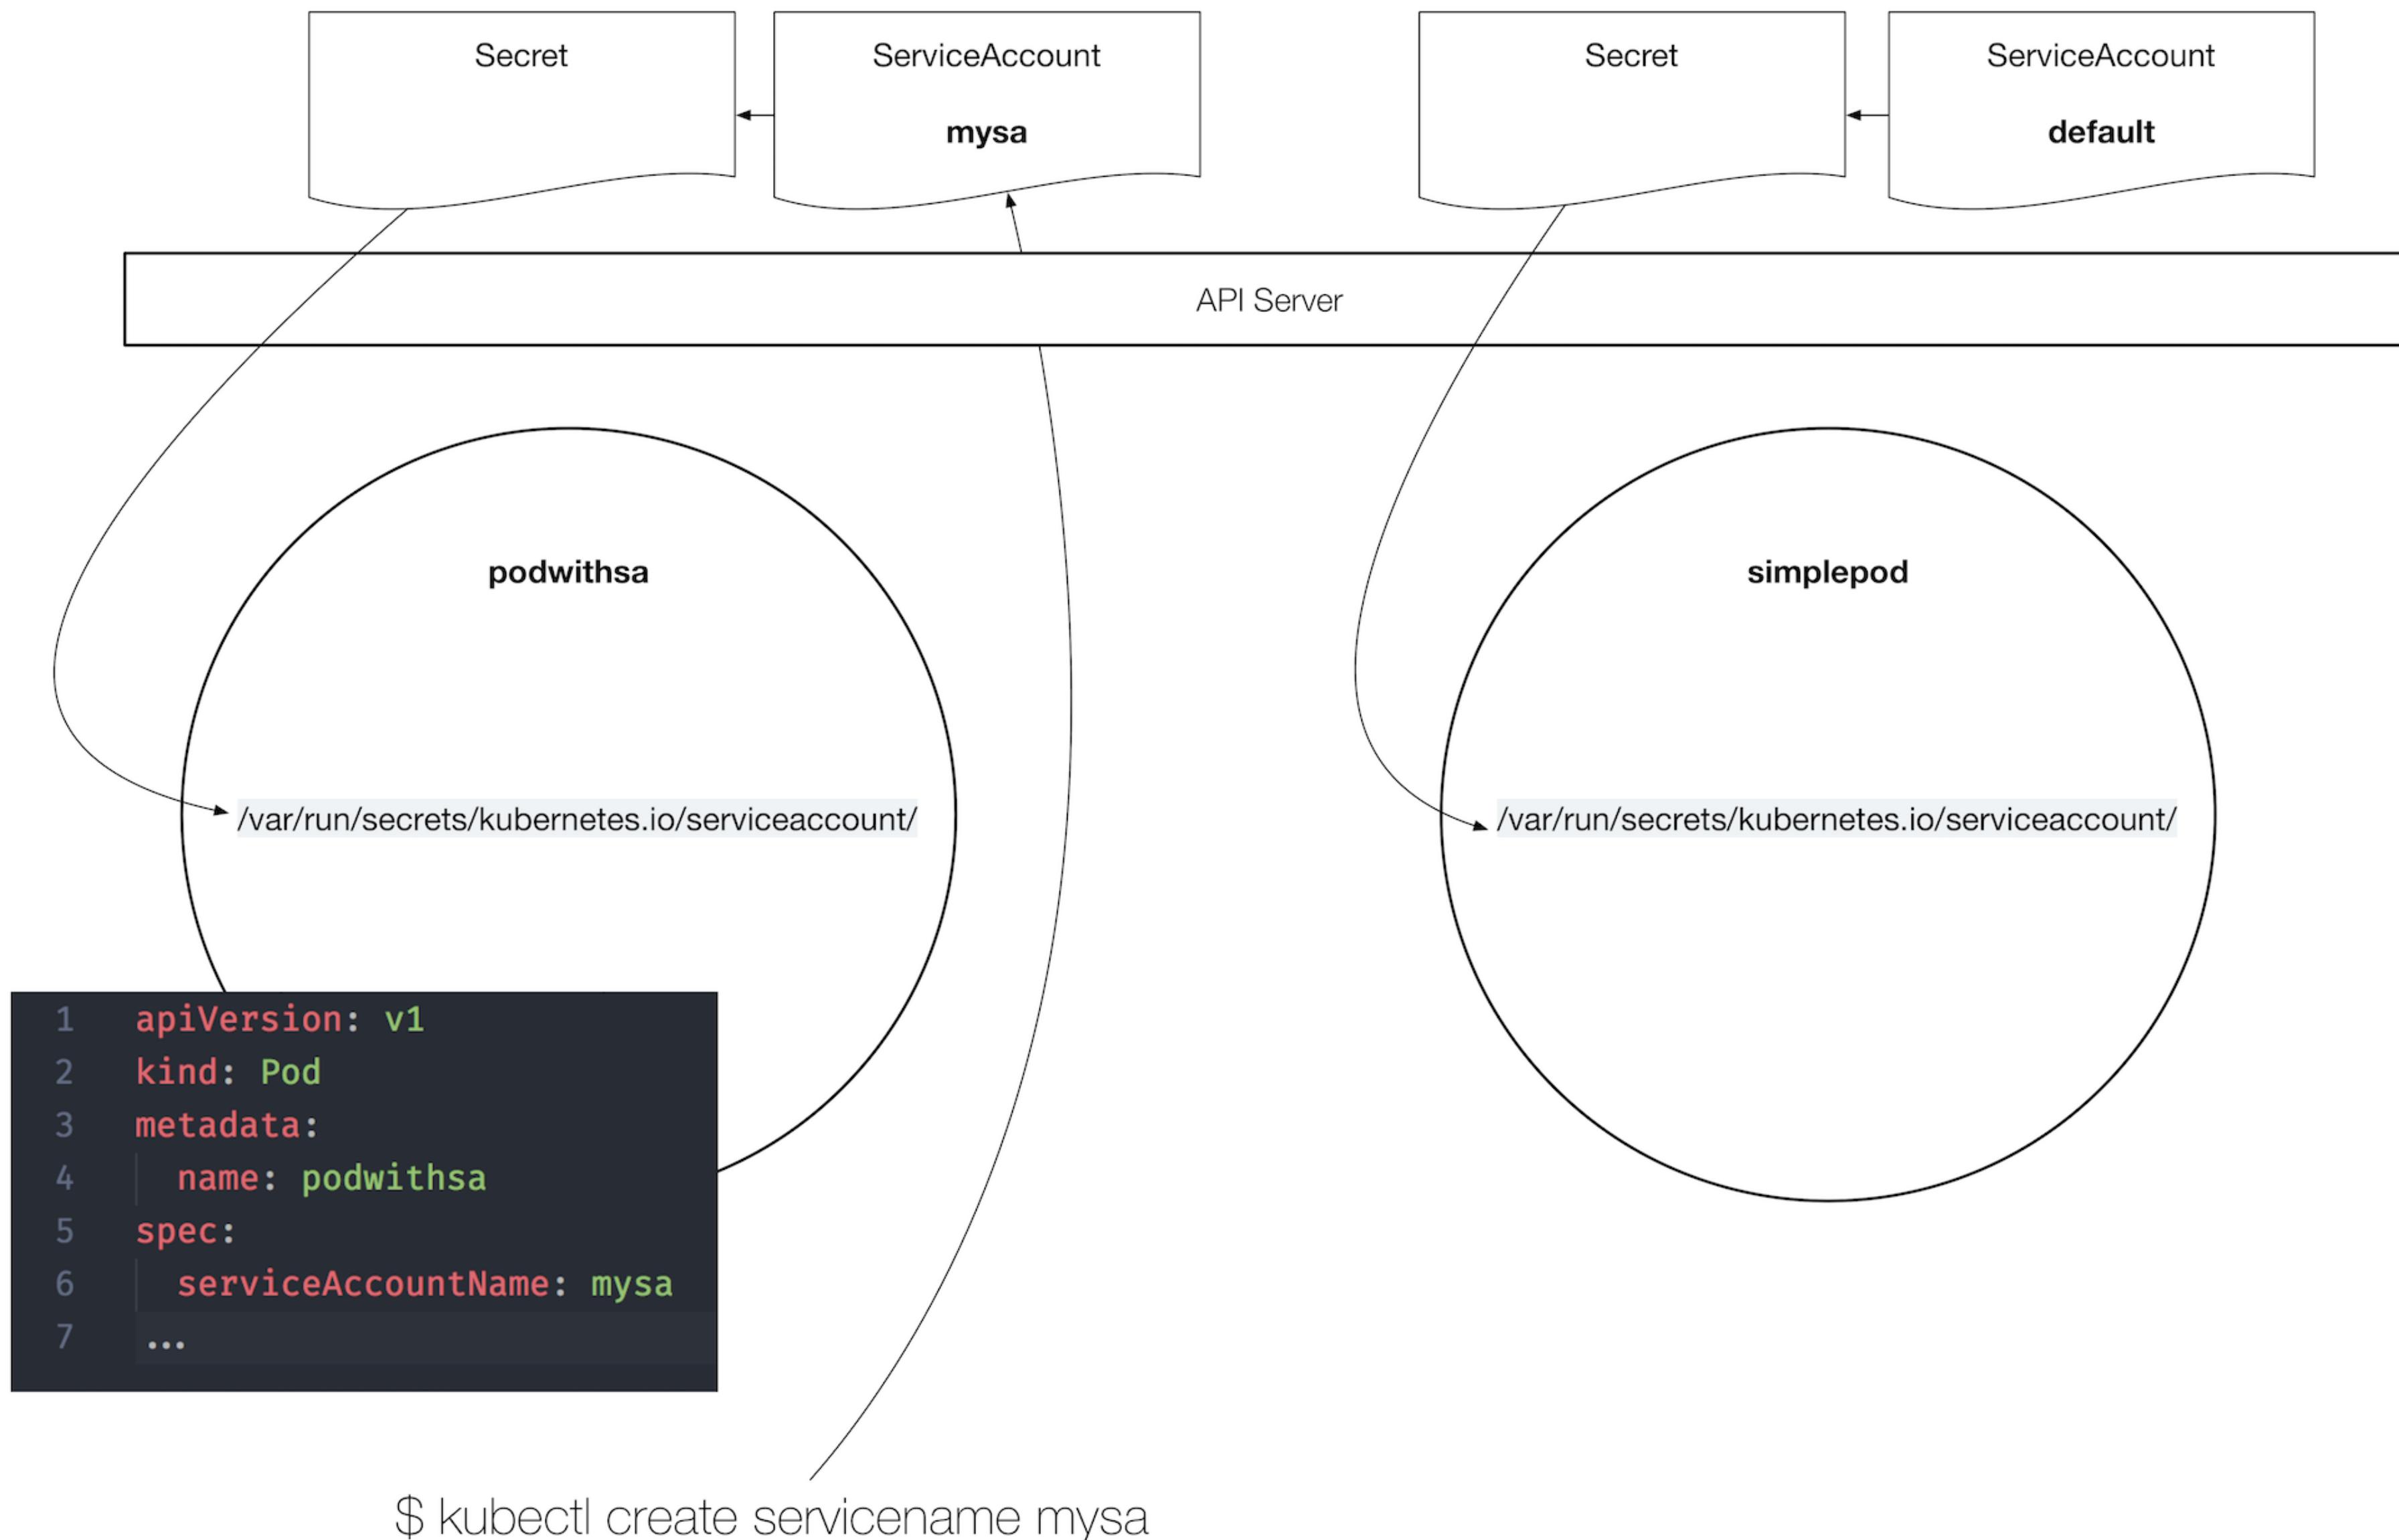
- human users
  - managed outside of Kubernetes
  - LDAP, SAML, Kerberos, etc.
- apps
  - running in containers in pods
  - first class resources via service accounts



# Service accounts

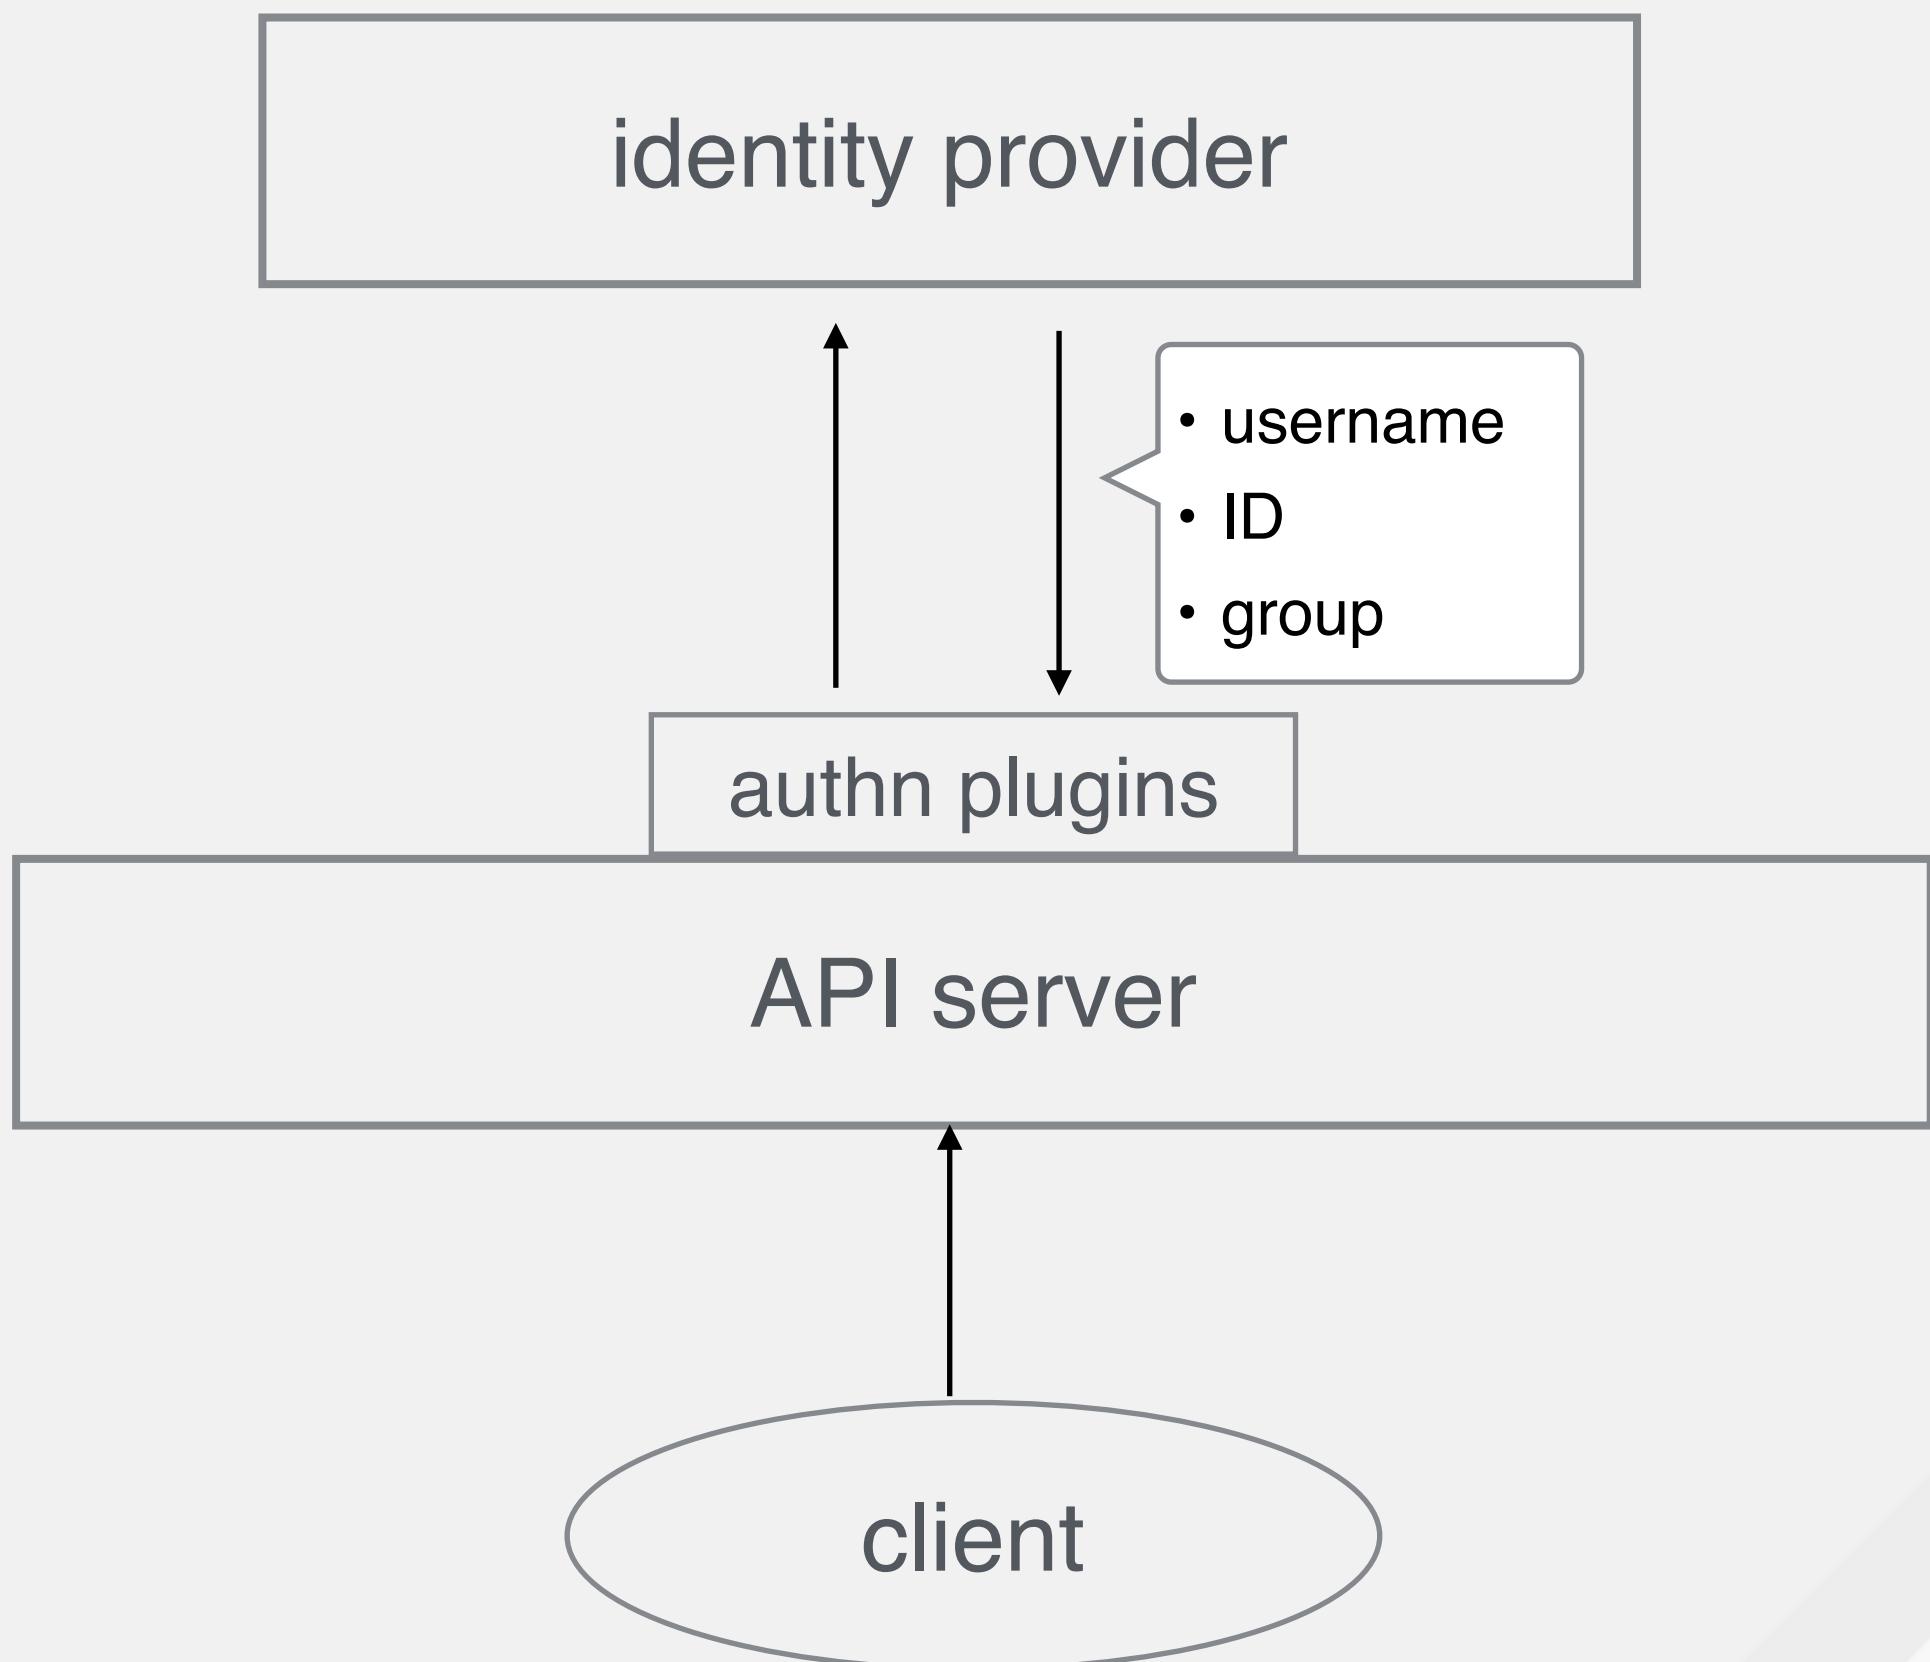
- provide identity for an app
- namespaced resources
- credentials via secret mounted into pod
- default service account per namespace

`system:serviceaccount:$NAMESPACE:$NAME`



# Authentication

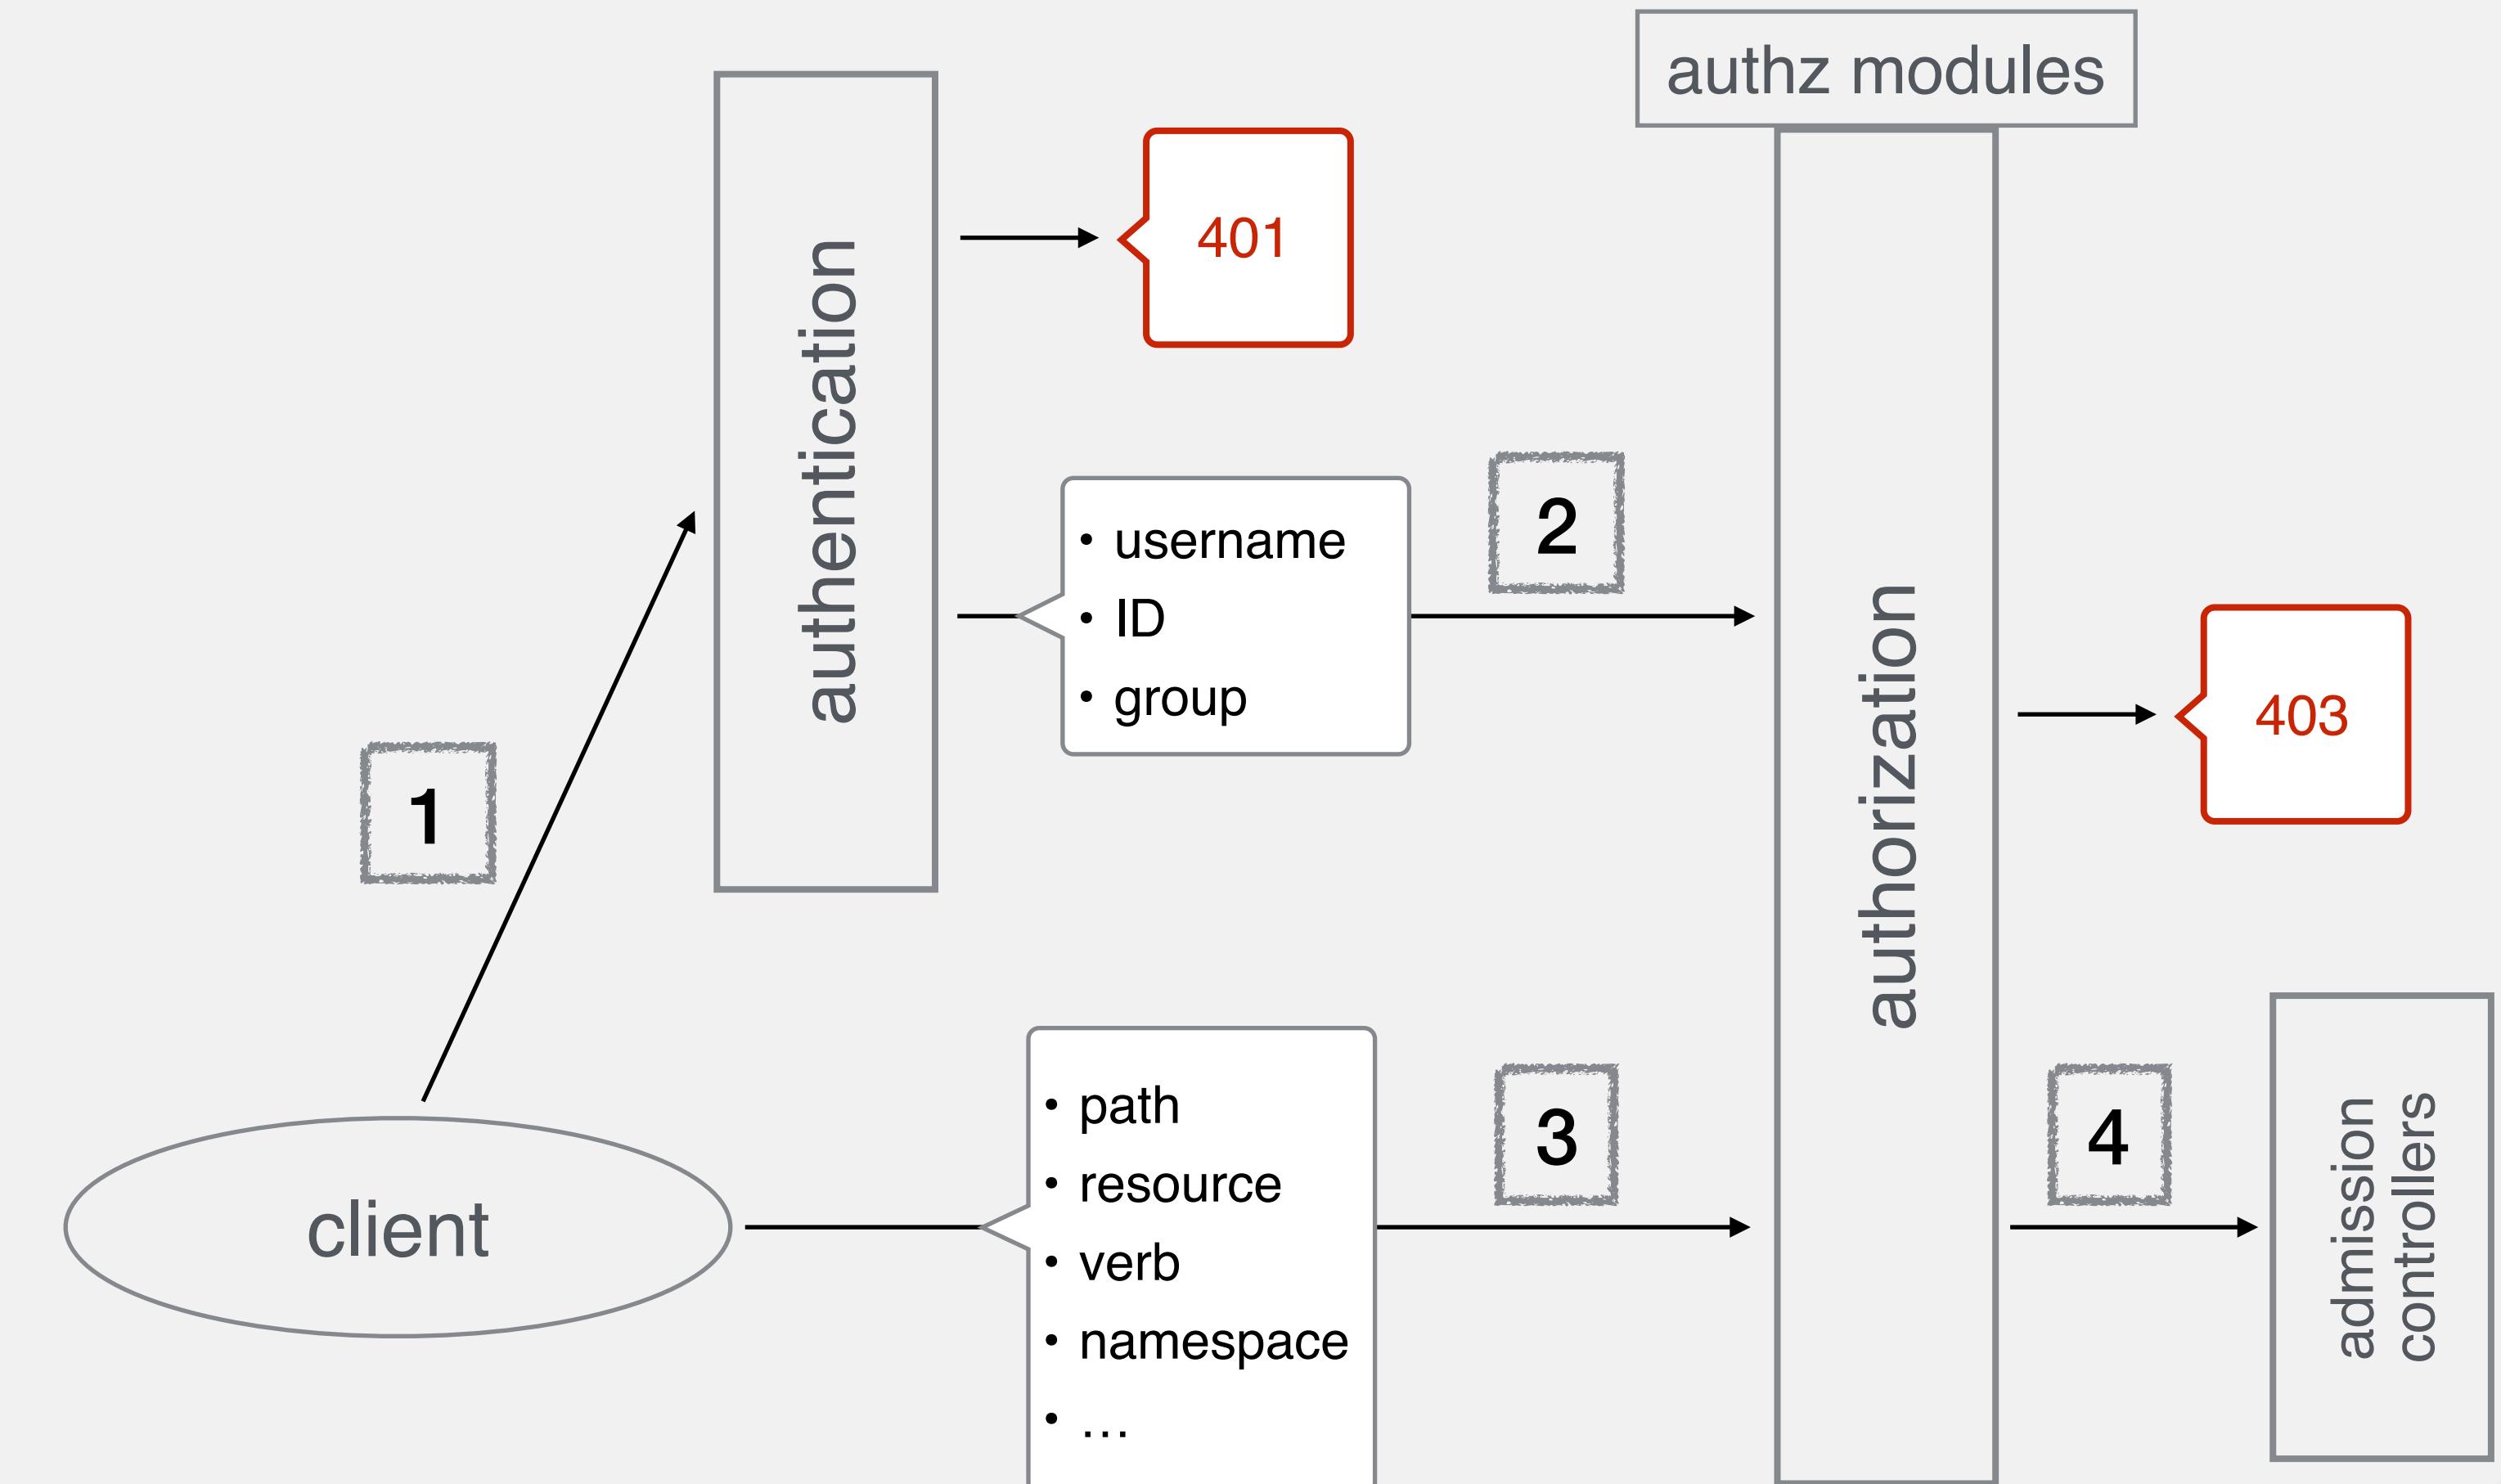
- static password/token file
- X509 client certs
- proxy+header
- OpenID Connect
- custom via Webhook



<https://kubernetes.io/docs/admin/authentication/>

# Authorization

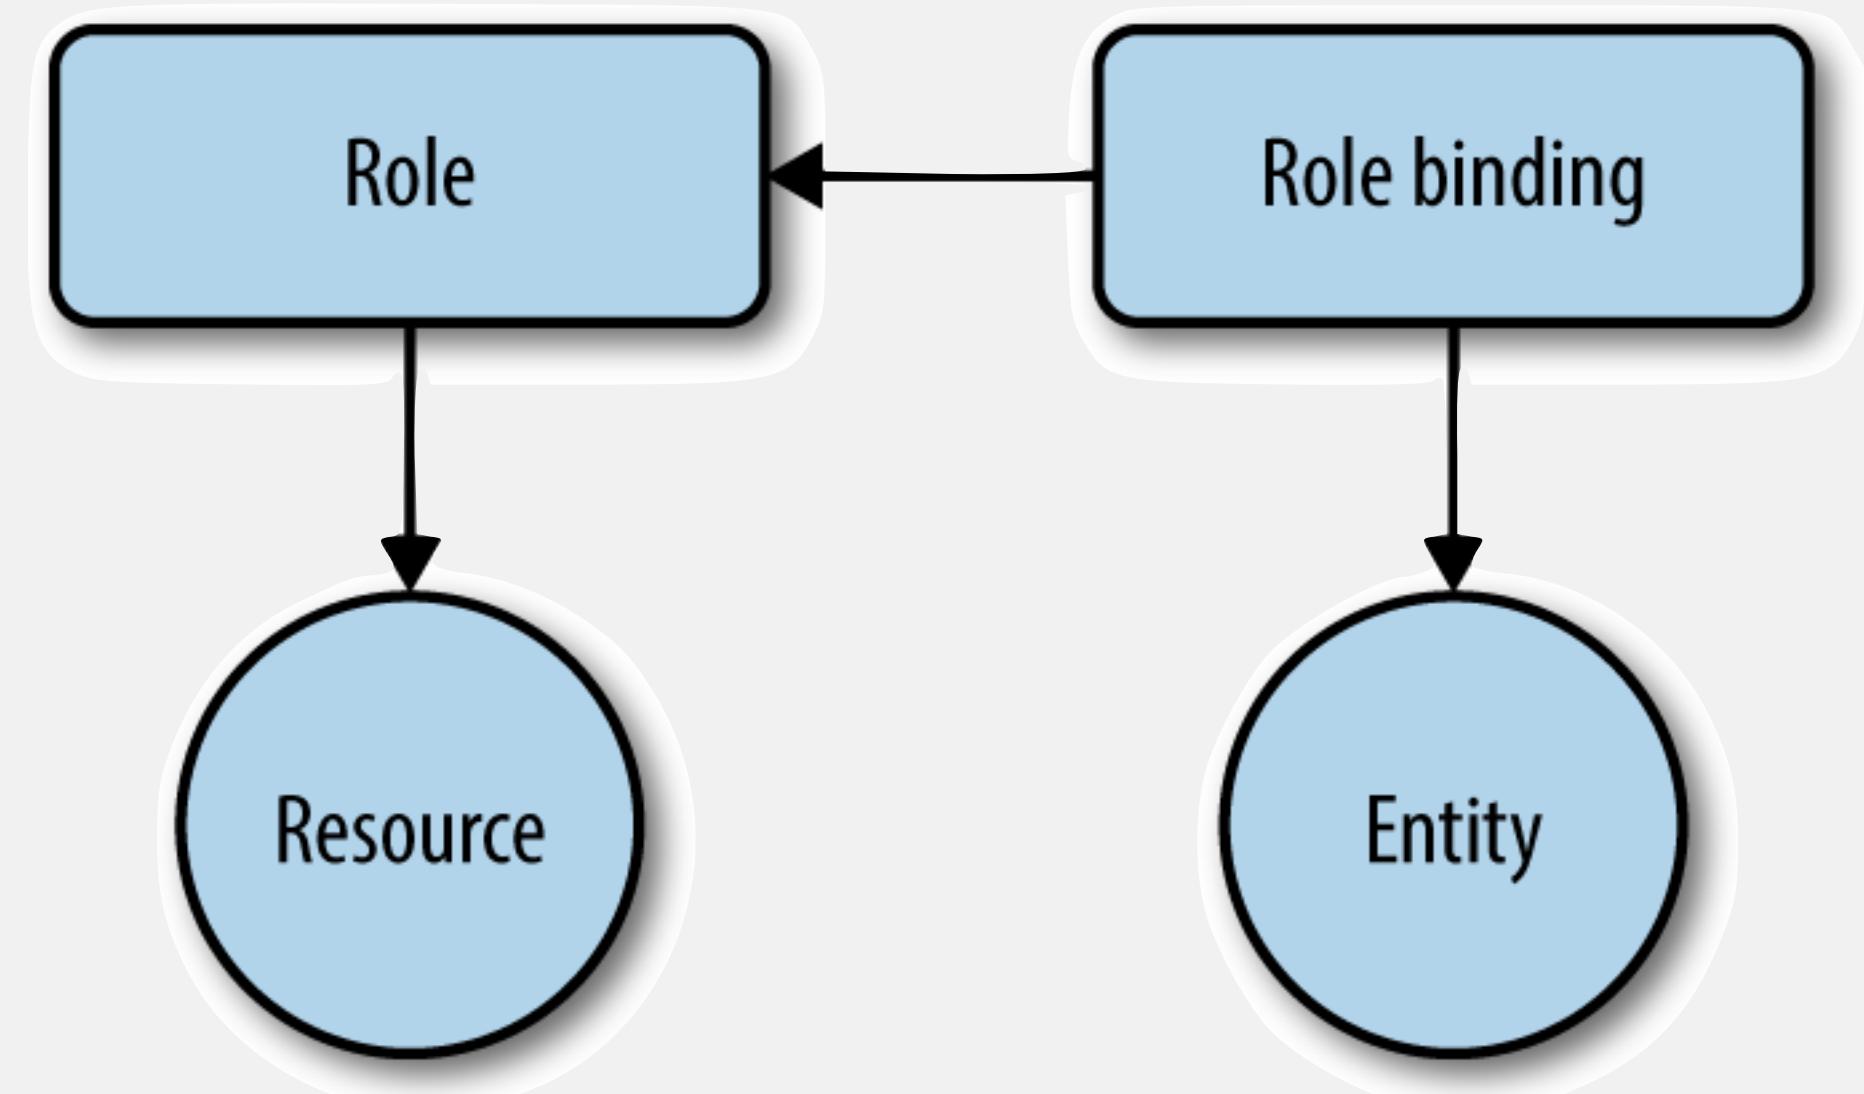
- Node (kubelet)
- ABAC (outdated)
- **RBAC**
- Webhook (external)



<https://kubernetes.io/docs/admin/authorization/>

# RBAC

- stable as of 1.8
- entities: service accounts, user, group
- scope: namespace or cluster
- roles and bindings
- privilege escalation prevention



<https://kubernetes.io/docs/reference/access-authn-authz/rbac/>

# Defaults

user-facing

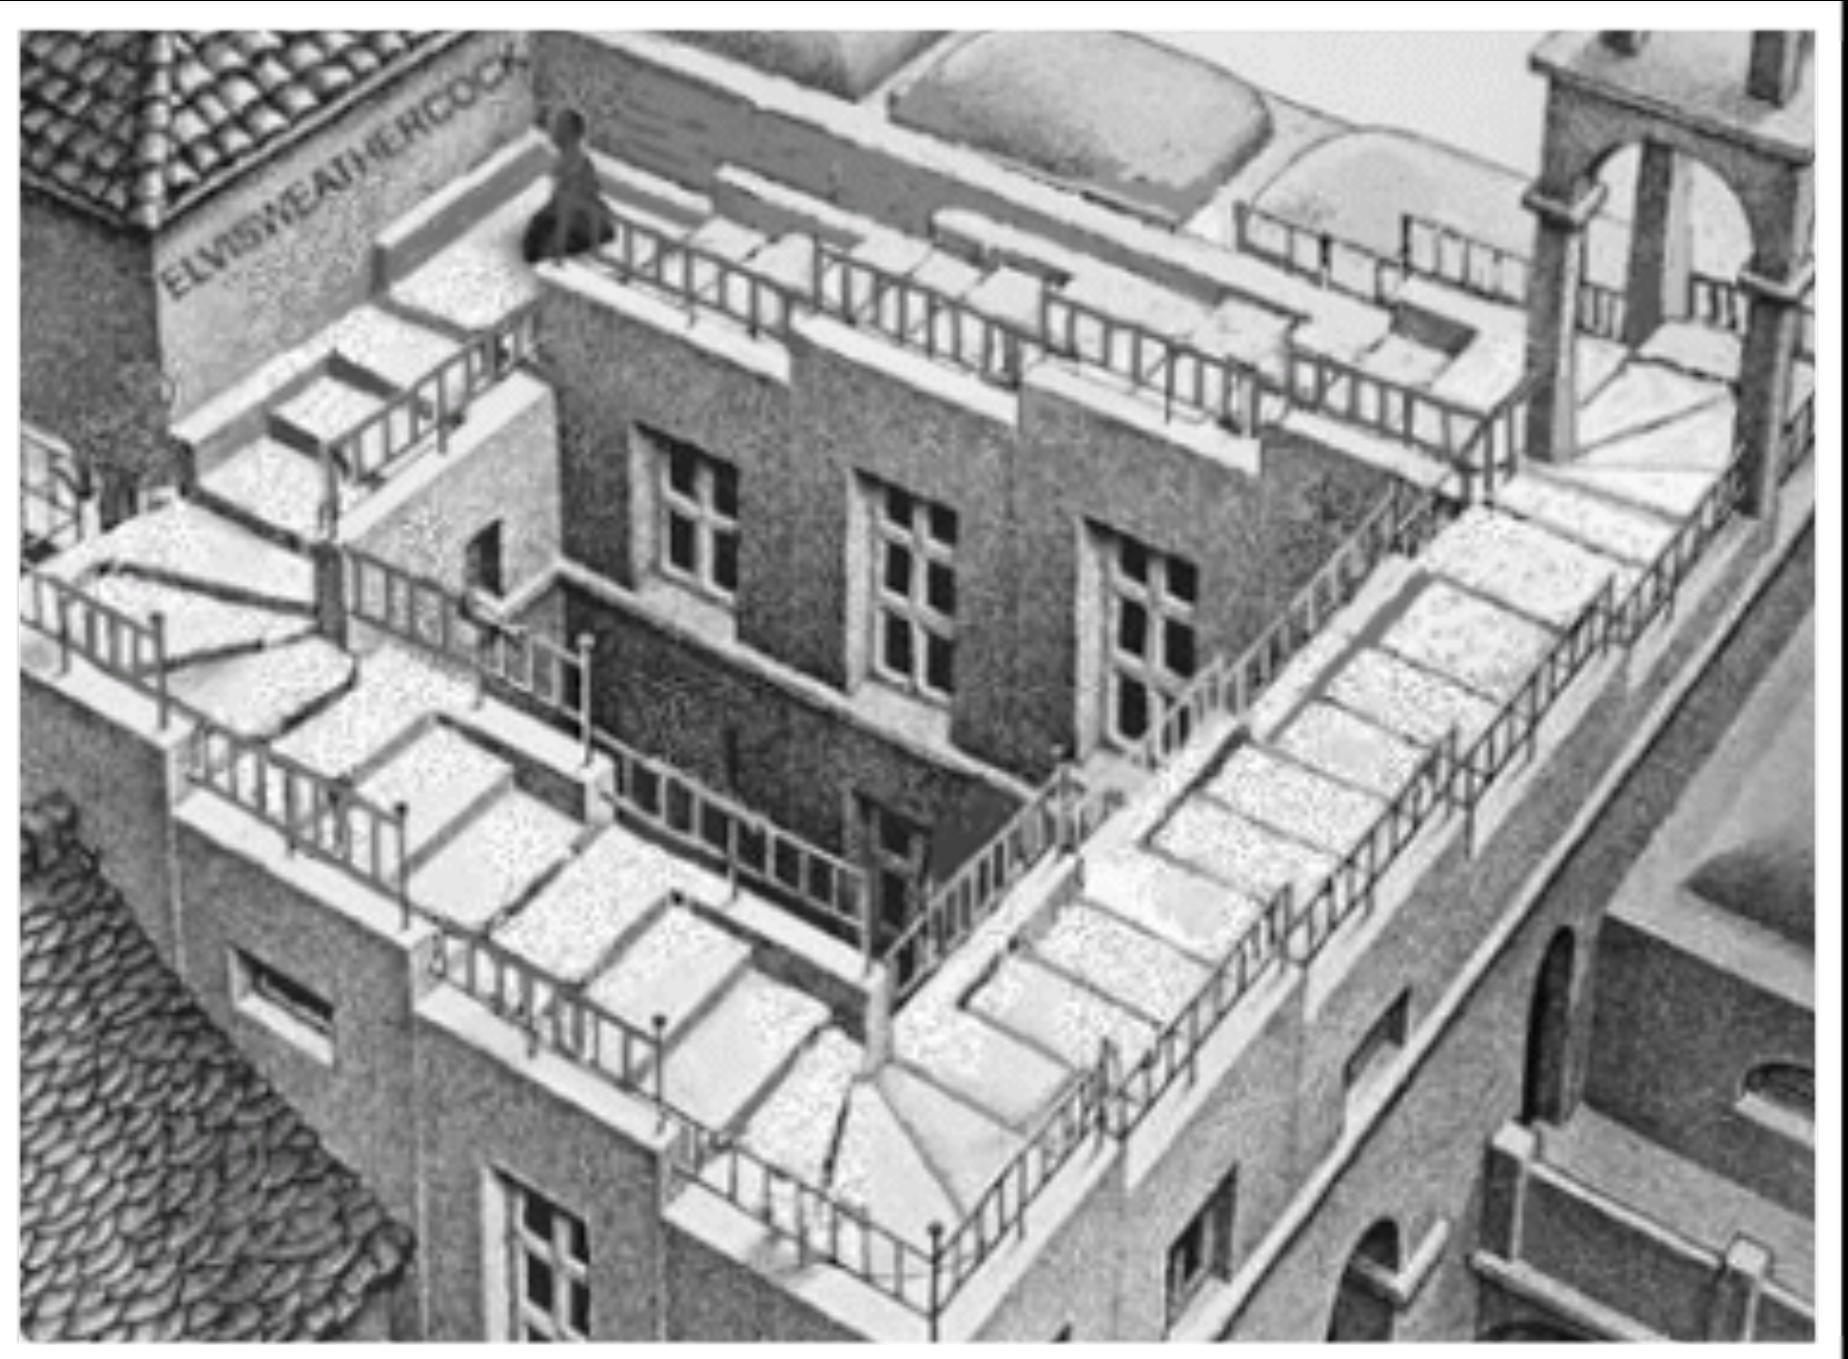
Default ClusterRole	Default ClusterRoleBinding
<b>cluster-admin</b>	<b>system:masters</b> group
<b>admin</b>	None
<b>edit</b>	None
<b>view</b>	None

core-components

Default ClusterRole	Default ClusterRoleBinding
<b>system:kube-scheduler</b>	<b>system:kube-scheduler</b> user
<b>system:kube-controller-manager</b>	<b>system:kube-controller-manager</b> user
<b>system:node</b>	None in 1.8+
<b>system:node-proxier</b>	<b>system:kube-proxy</b> user

other components

Default ClusterRole	Default ClusterRoleBinding
<b>system:auth-delegator</b>	None
<b>system:heapster</b>	None
<b>system:kube-aggregator</b>	None
<b>system:kube-dns</b>	<b>kube-dns</b> service account in the <b>kube-system</b> namespace
<b>system:kubelet-api-admin</b>	None
<b>system:node-bootstrapper</b>	None
<b>system:node-problem-detector</b>	None
<b>system:persistent-volume-provisioner</b>	None



... demo time!

# Tooling

- kubectl create, kubectl auth
- <https://github.com/coreos/dex>
- <https://github.com/heptio/authenticator>
- <https://github.com/liggitt/audit2rbac>

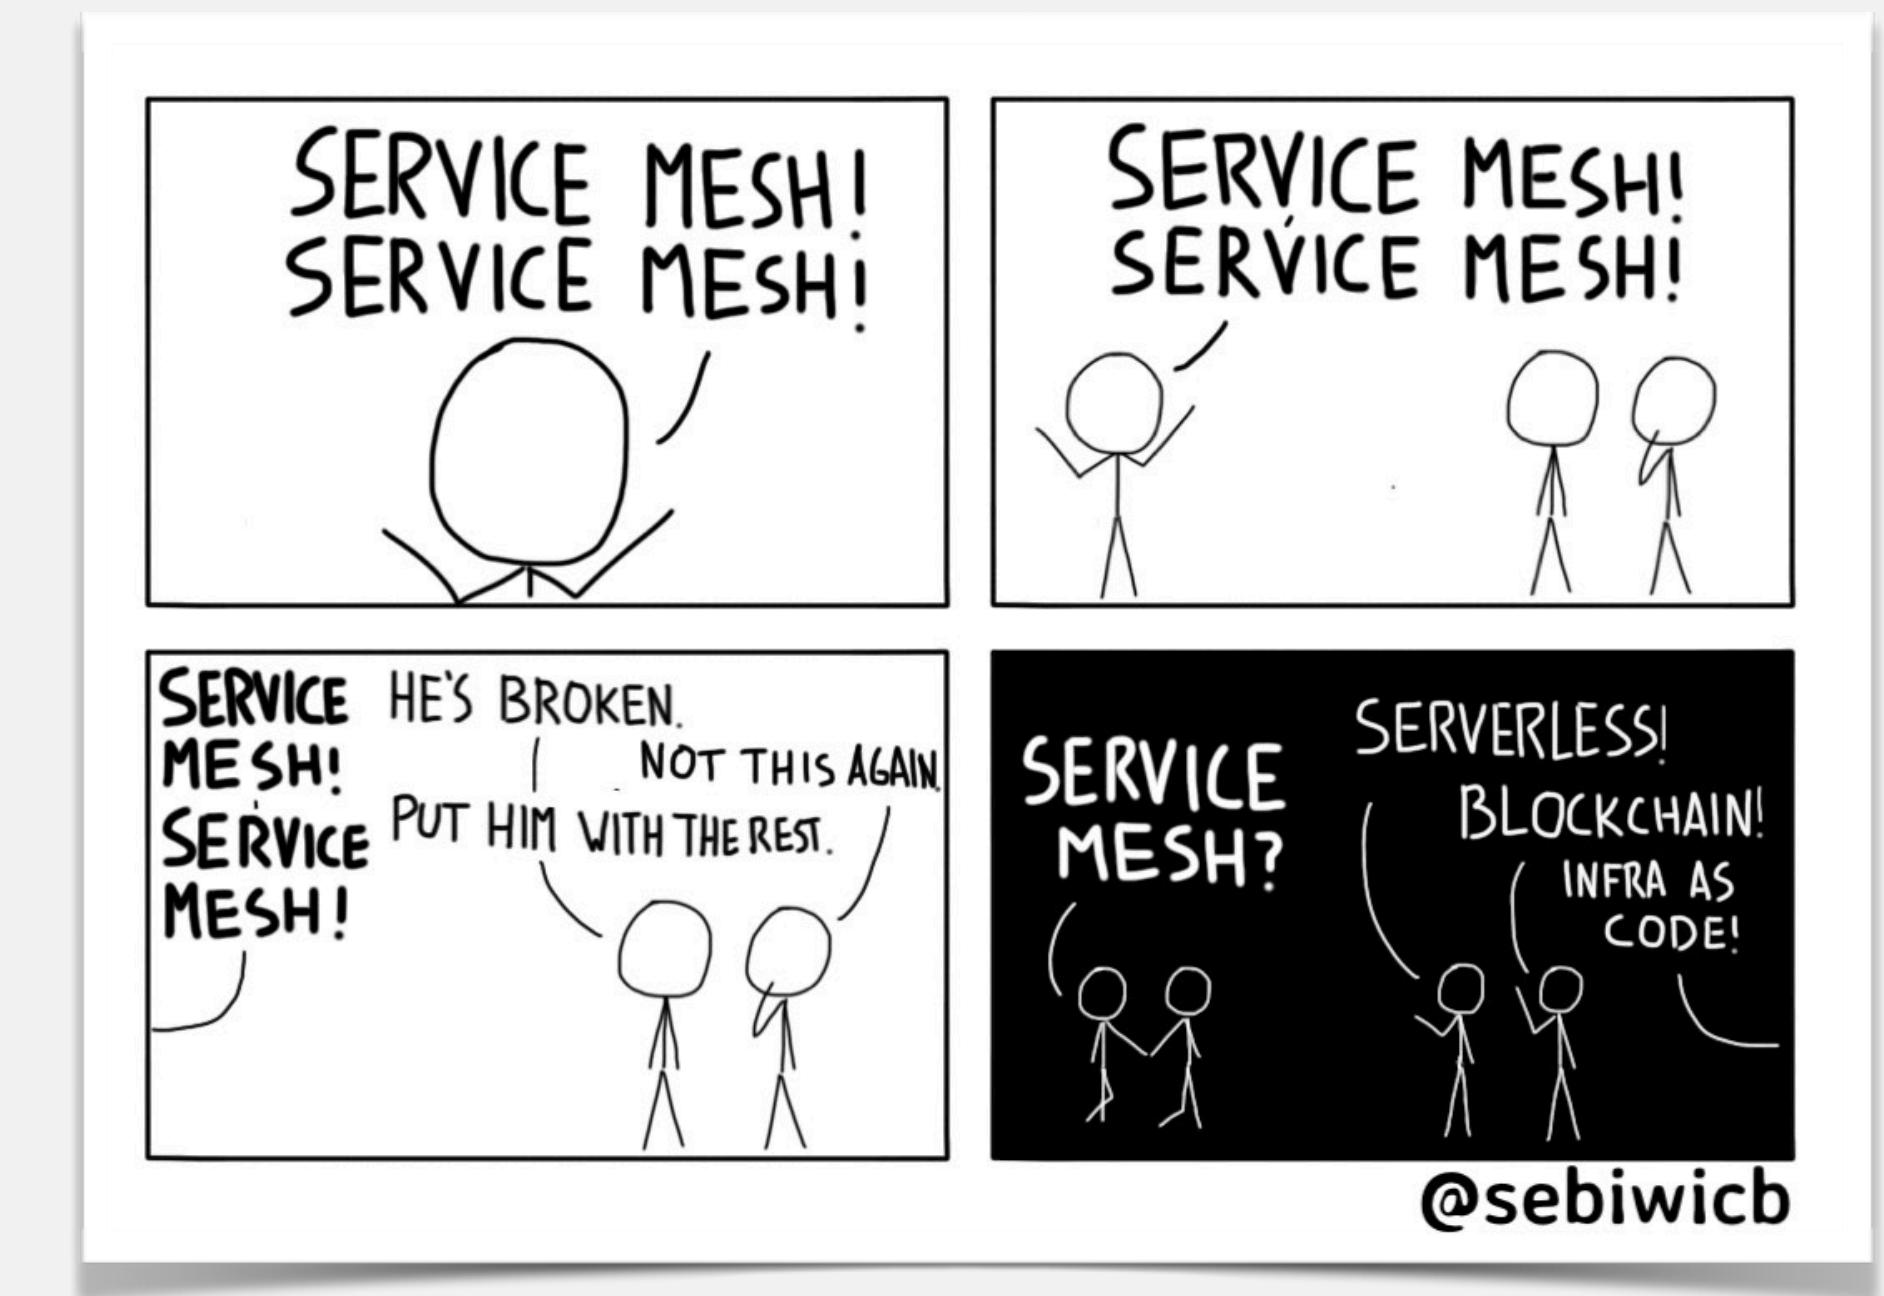
# Further reading

- Controlling Access to the Kubernetes API
- Kubernetes deep dive: API Server – part 1
- Certifik8s: All You Need to Know About Certificates in Kubernetes
- Kubernetes Auth and Access Control
- Effective RBAC
- Single Sign-On for Kubernetes: An Introduction
- Let's Encrypt, OAuth 2, and Kubernetes Ingress

# Communication

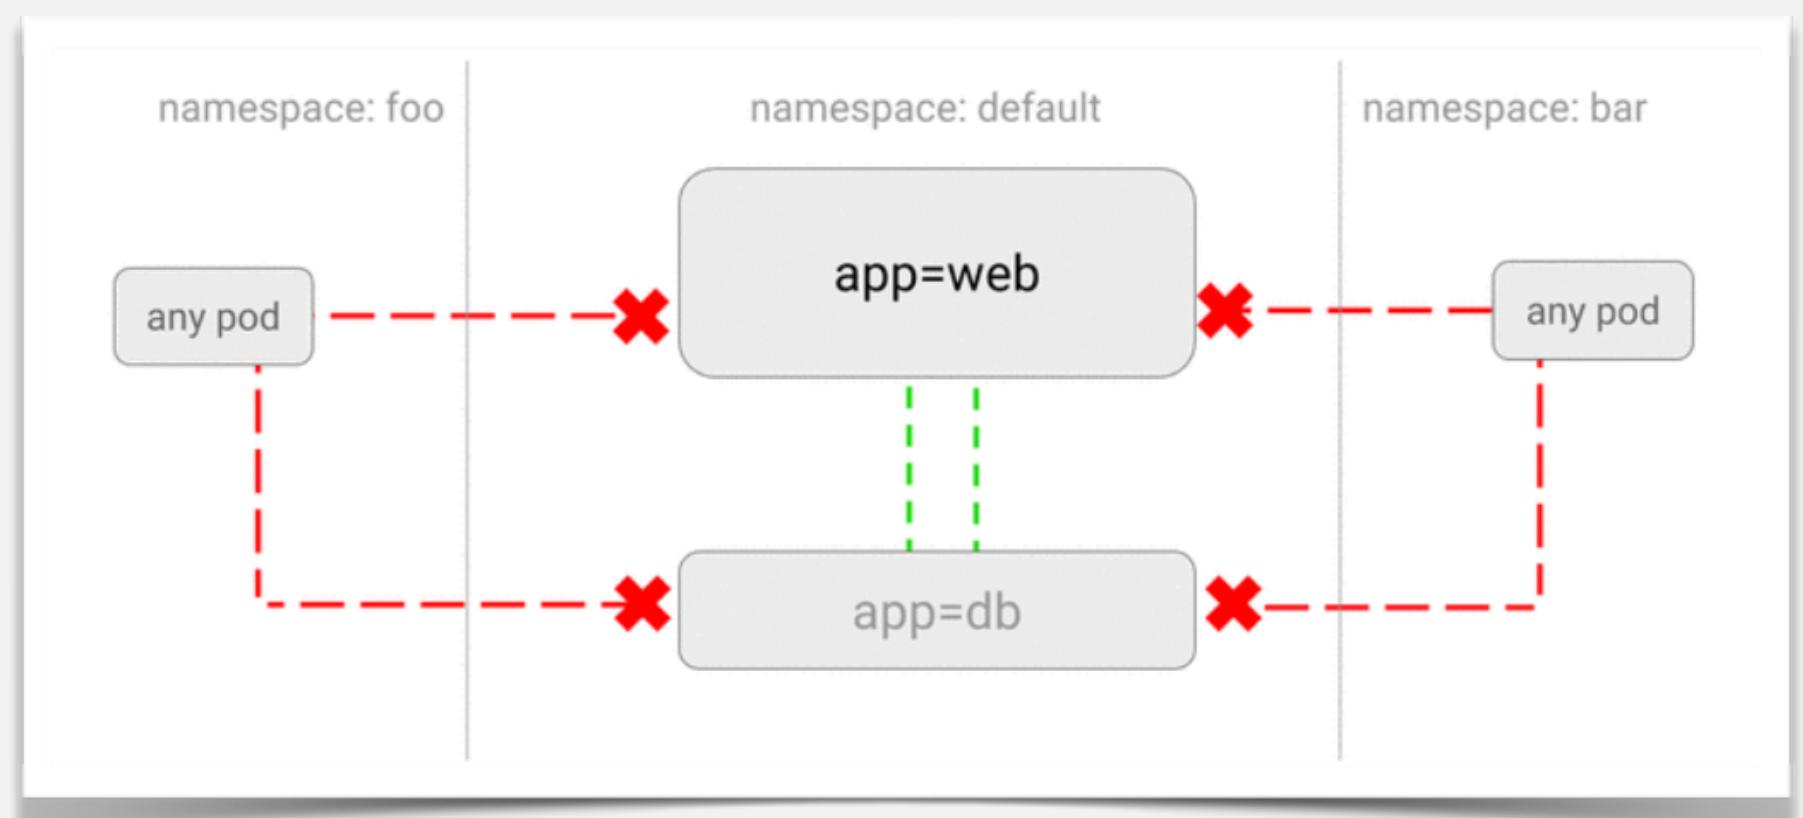
# What?

- encryption on the wire (TLS everywhere)
- network policies
- service meshes



# Network policies

- defines pod-to-pod communication
- enforced by network plugin

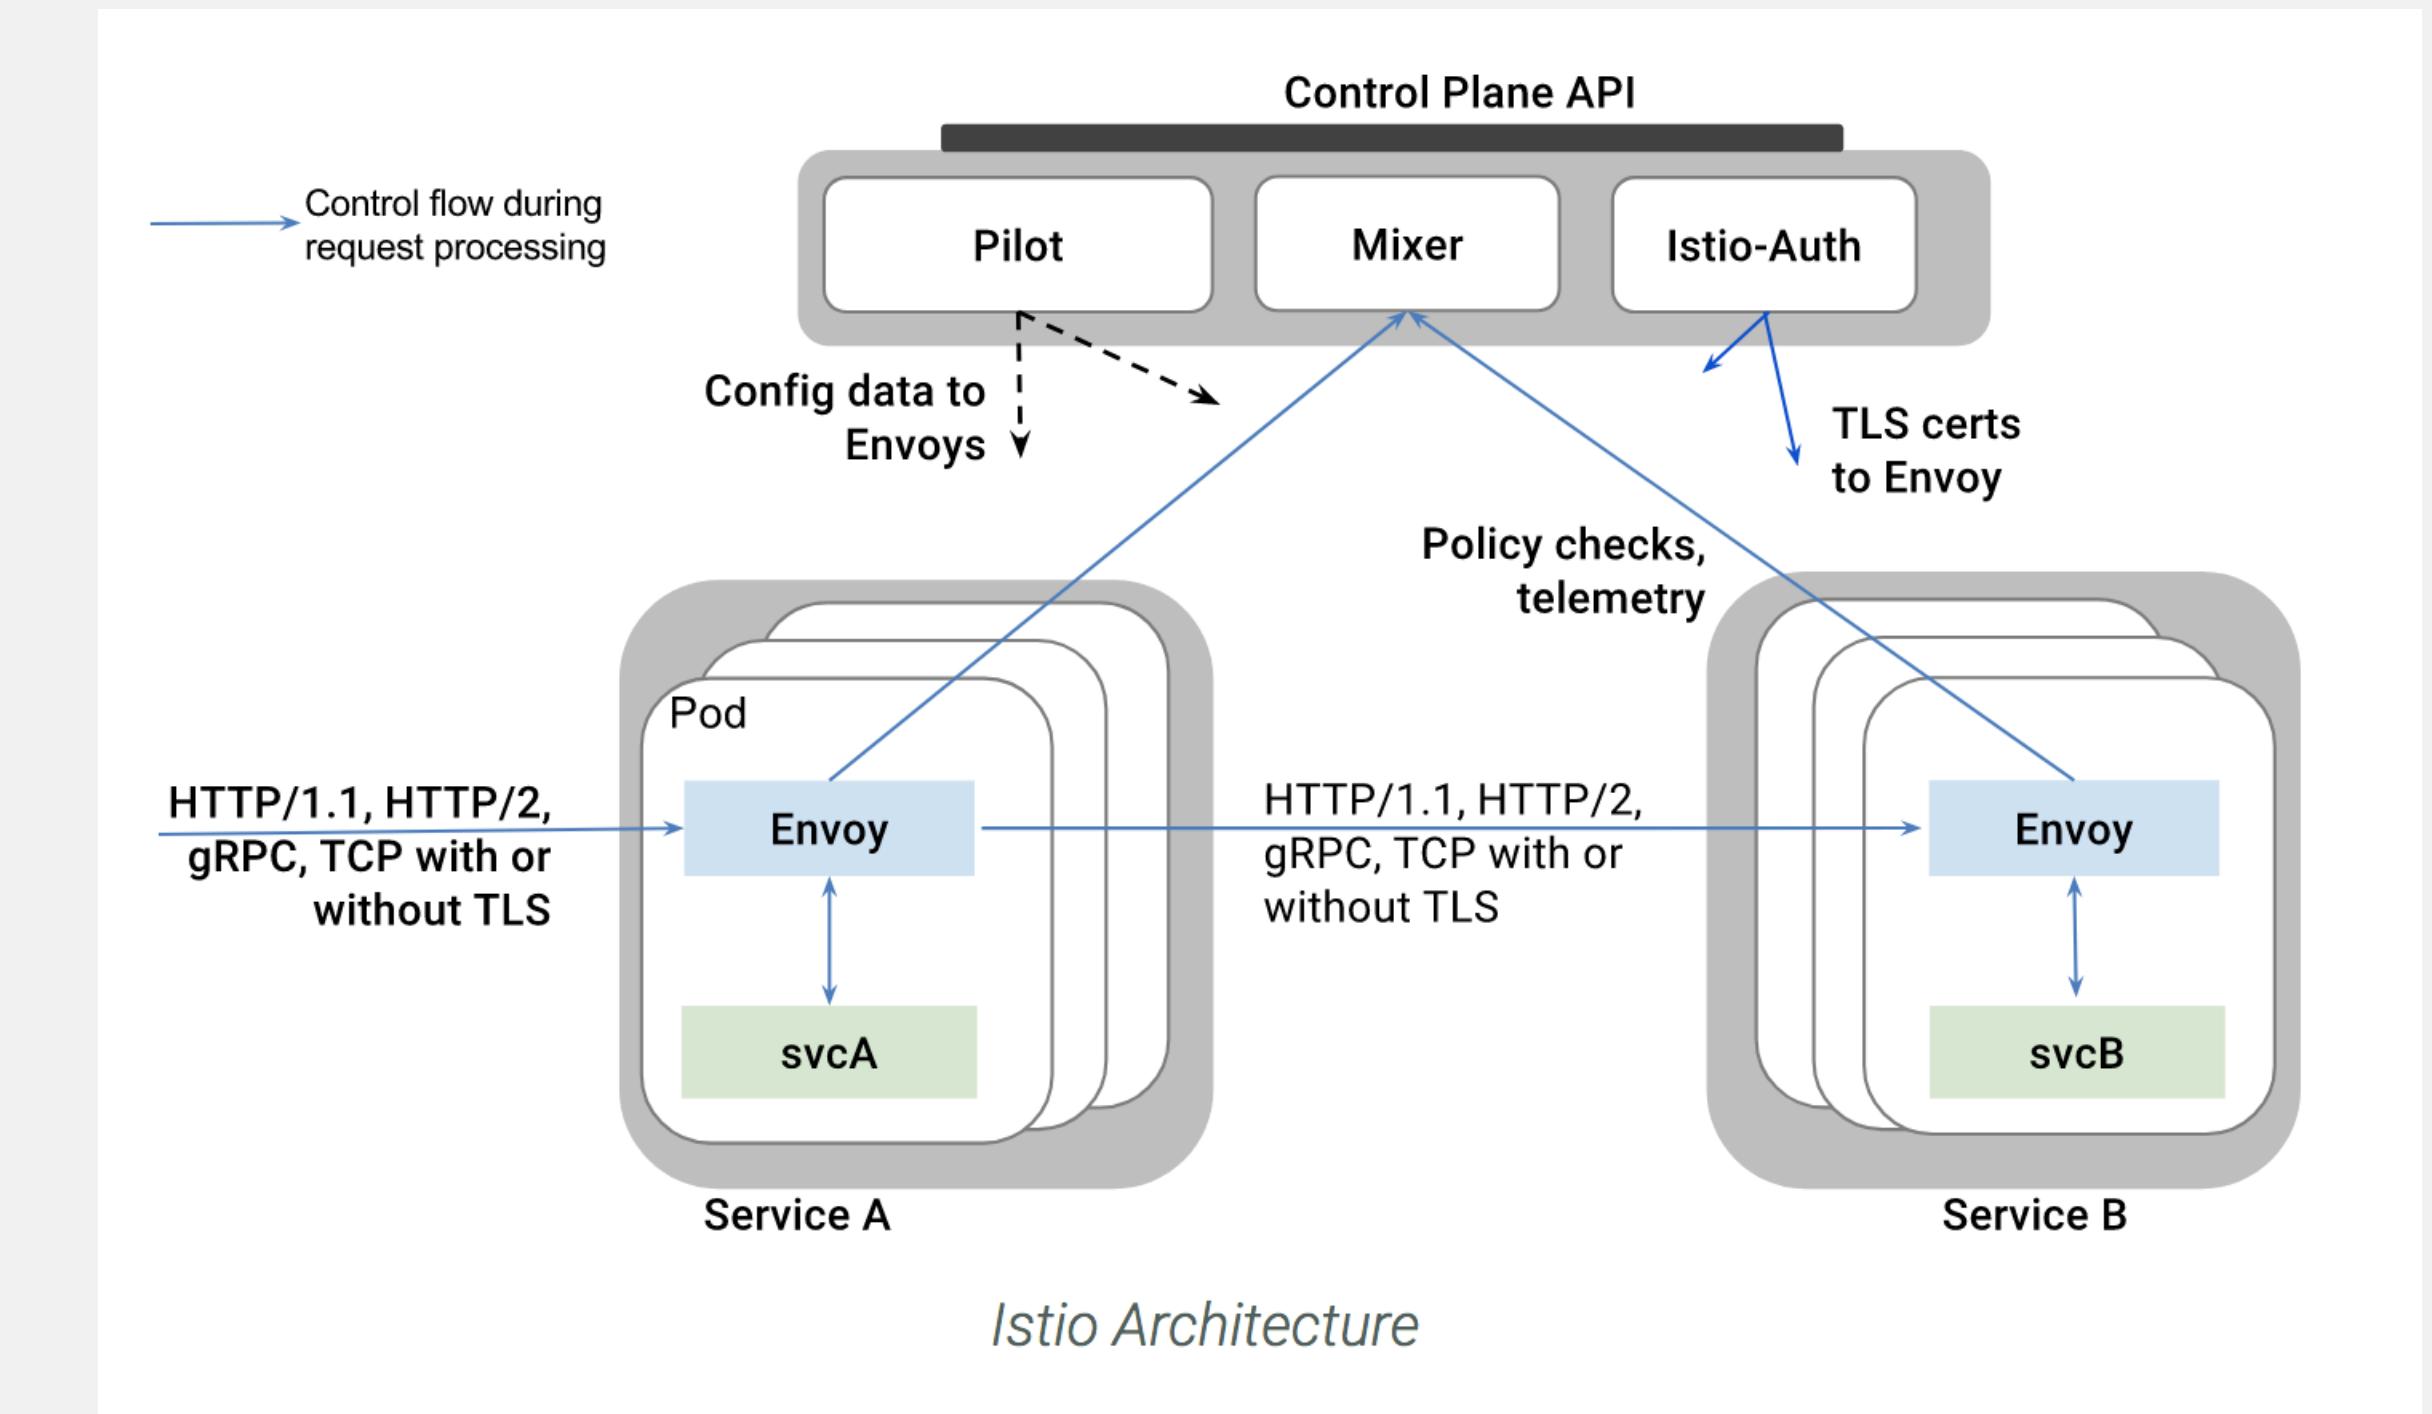


Source: Securing Kubernetes Cluster Networking by Ahmet Alp Balkan

<https://kubernetes.io/docs/concepts/services-networking/network-policies/>

# Service meshes

- traffic management
- policy enforcement
- monitoring & tracing
- no app code changes



[istio.io](https://istio.io)

# Tooling

- <https://github.com/aporeto-inc/trireme-kubernetes>
- <https://github.com/jetstack/cert-manager/>
- <https://spiffe.io/>
- <https://www.openpolicyagent.org/>
- <https://linkerd.io/>
- <https://conduit.io/>

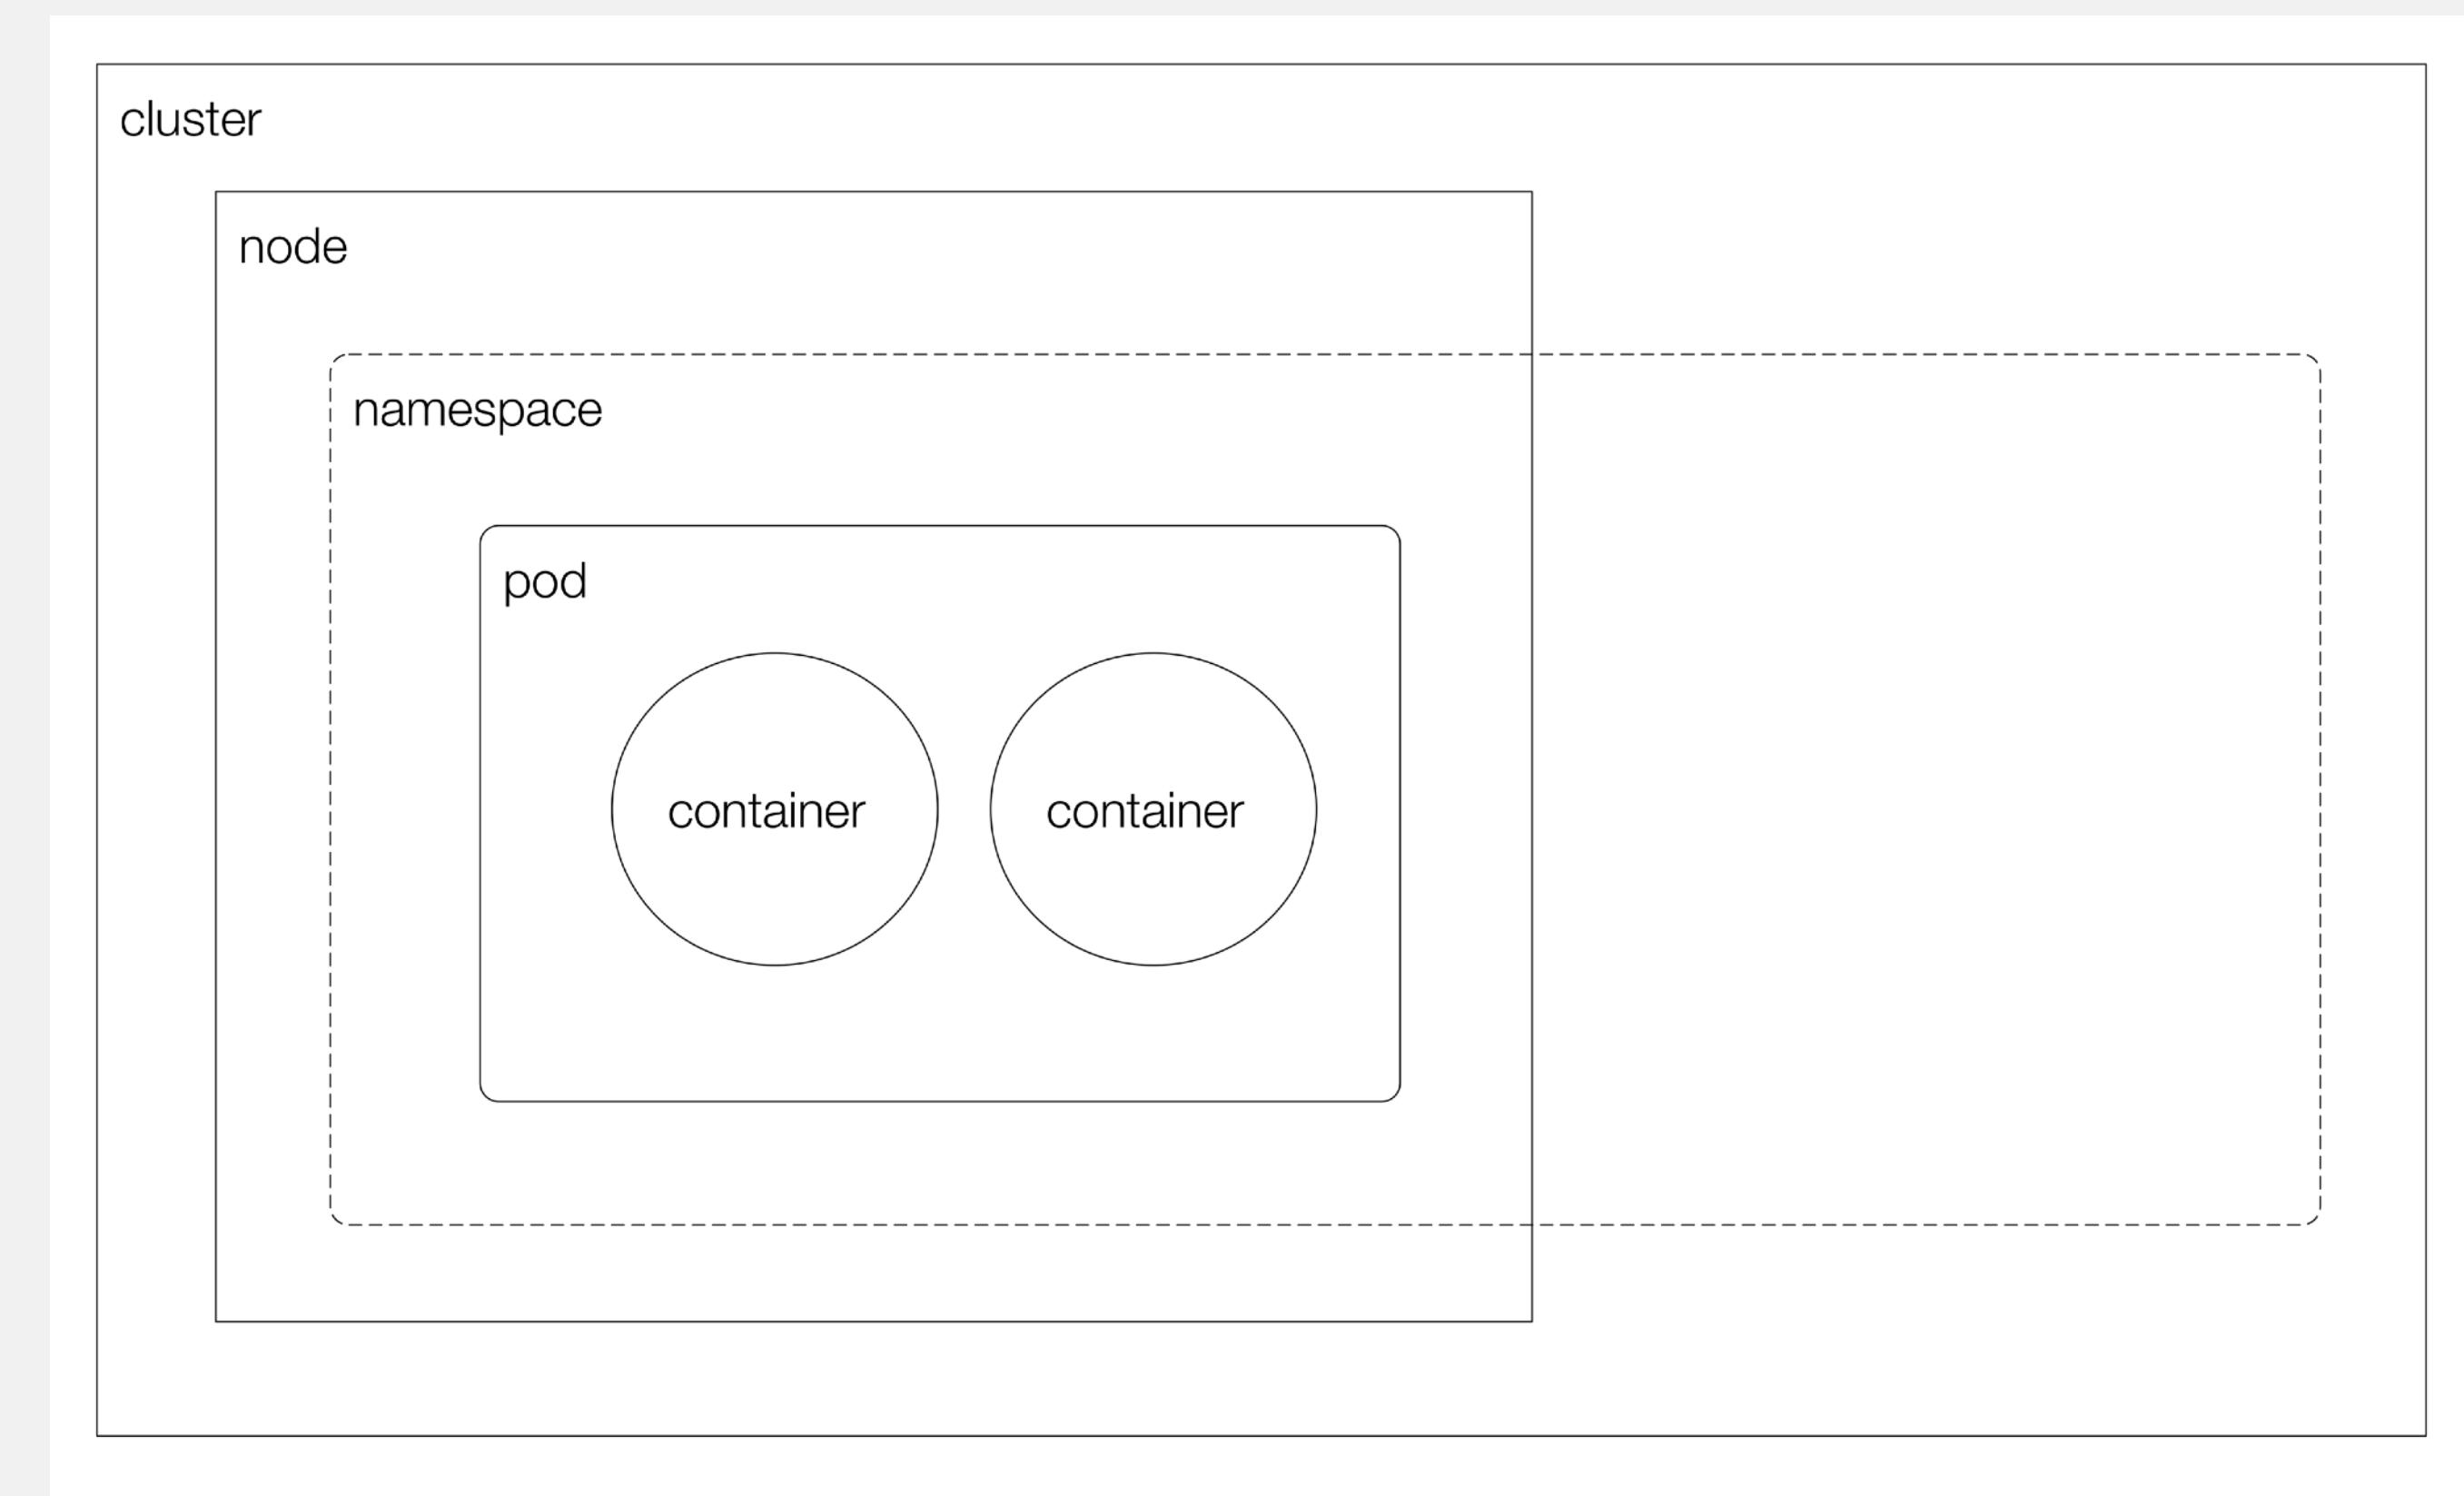
# Further reading

- How Kubernetes certificate authorities work
- Securing Kubernetes Cluster Networking
- Tutorials and Recipes for Kubernetes Network Policies feature
- Kubernetes Security Context and Kubernetes Network Policy
- Kubernetes Application Operator Basics

# Apps

# What?

- security boundaries
- segregation
- secrets



Based on: Exploring container security: Isolation at different layers of the Kubernetes stack

# Segregation

- namespace-level: visibility and access
- node-level:
  - separate sensitive workloads via affinity and taints
  - minimize blast radius (node authorizer)
- pod-level: limit communication via network policies or service mesh

# Secrets

- Namespaced objects to store sensitive information
- Access via volume or environment variable
- Data is stored in tmpfs volumes
- Per-secret size limit of 1MB
- Only base64 encoded, need to enable encryption at rest

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDF1MmU2N2Rm
```

<https://kubernetes.io/docs/concepts/configuration/secret/>

# Tooling

- <https://github.com/kelseyhightower/konfd>
- <https://github.com/hashicorp/vault-plugin-auth-kubernetes>
- <https://github.com/bitnami-labs/sealed-secrets>
- <https://github.com/shyiko/kubesec>
- <https://github.com/weaveworks/flux>

# Further reading

- Docs:
  - Configure a Security Context for a Pod or Container
  - Pod Security Policies
- Shipping in Pirate-Infested Waters: Practical Attack and Defense in Kubernetes
- Exploring container security: Isolation at different layers of the Kubernetes stack
- Security Best Practices for Kubernetes Deployment

# Control plane

# Good practices

- secure API server, etcd, dashboard
- secure kubelet
- limit access to cloud provider metadata
- limit access to metrics
- perform auditing

# Tooling

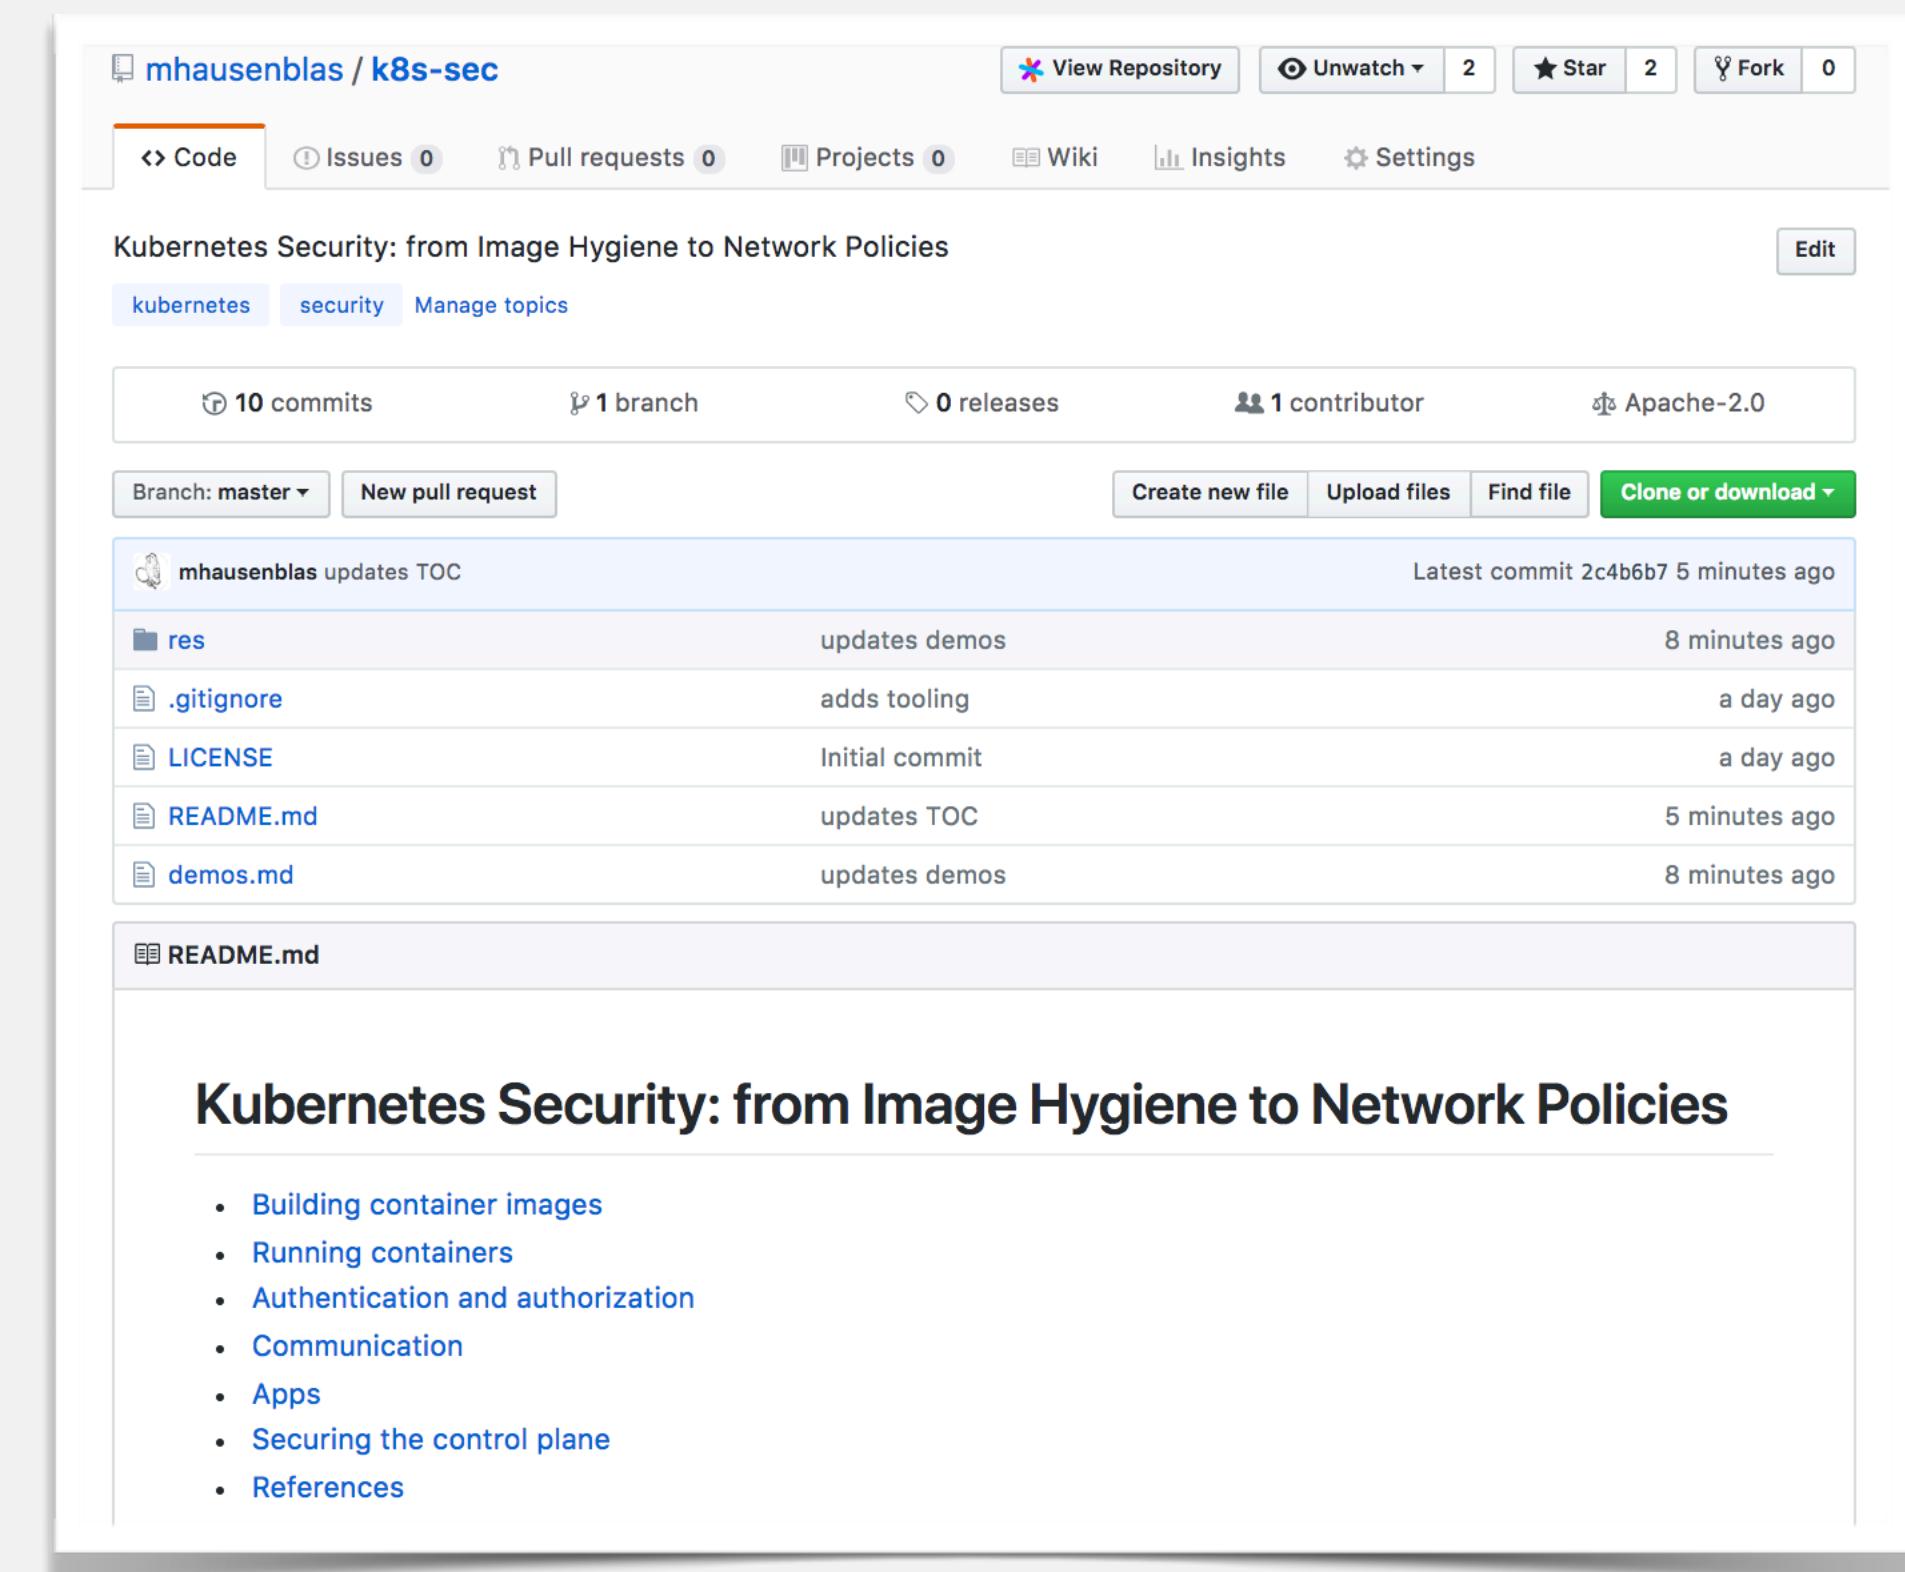
- <https://github.com/bgeesaman/kubeatf>
- <https://github.com/Shopify/kubeaudit>
- <https://k8guard.github.io/>
- <https://www.vaultproject.io/>

# Further reading

- Docs:
  - Securing a Cluster
  - Encrypting Secret Data at Rest
  - Auditing
- Securing Kubernetes components: kubelet, etcd and Docker registry
- K8s security best practices
- Kubernetes Security - Best Practice Guide
- Lessons from the Cryptojacking Attack at Tesla

# Resources

# Demos and references



<https://github.com/mhausenblas/k8s-sec>

# Articles

- NIST Special Publication 800-190: Application Container Security Guide  
<https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-190.pdf>
- Hacking and Hardening Kubernetes Clusters by Example, Brad Geesaman, KubeCon 2017  
<https://www.youtube.com/watch?v=vTgQLzeBfRU>
- Kubernetes Security Best Practices, Ian Lewis, FOSDEM 2018  
<https://www.youtube.com/watch?v= pzAwTC8KYV8>
- Continuous Kubernetes Security, Andrew Martin, microXchg 2018  
<https://www.youtube.com/watch?v=YtrA7eauSSg>
- What Does “Production Ready” Really Mean for a Kubernetes Cluster?  
<https://weave.works/blog/what-does-production-ready-really-mean-for-a-kubernetes-cluster>



redhat<sup>®</sup>

# learn.openshift.com

A screenshot of the learn.openshift.com website. At the top right, there are links for "Report an Issue" and "SIGN UP TO OPENSHIFT ONLINE FOR FREE". The main header says "OPENSHIFT" and "RED HAT OPENSHIFT Interactive Learning Portal". Below this, a text block explains that the portal provides pre-configured OpenShift instances for experimentation. There are three course cards at the bottom: "Foundations of OpenShift" (START COURSE), "Building Applications On OpenShift" (START COURSE), and "OpenShift Playgrounds" (START COURSE).



[plus.google.com/+RedHat](https://plus.google.com/+RedHat)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[twitter.com/RedHatNews](https://twitter.com/RedHatNews)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)