

Securing Systems at Cloud Scale

Ian Massingham, Technical Evangelist

 @IanMmmm



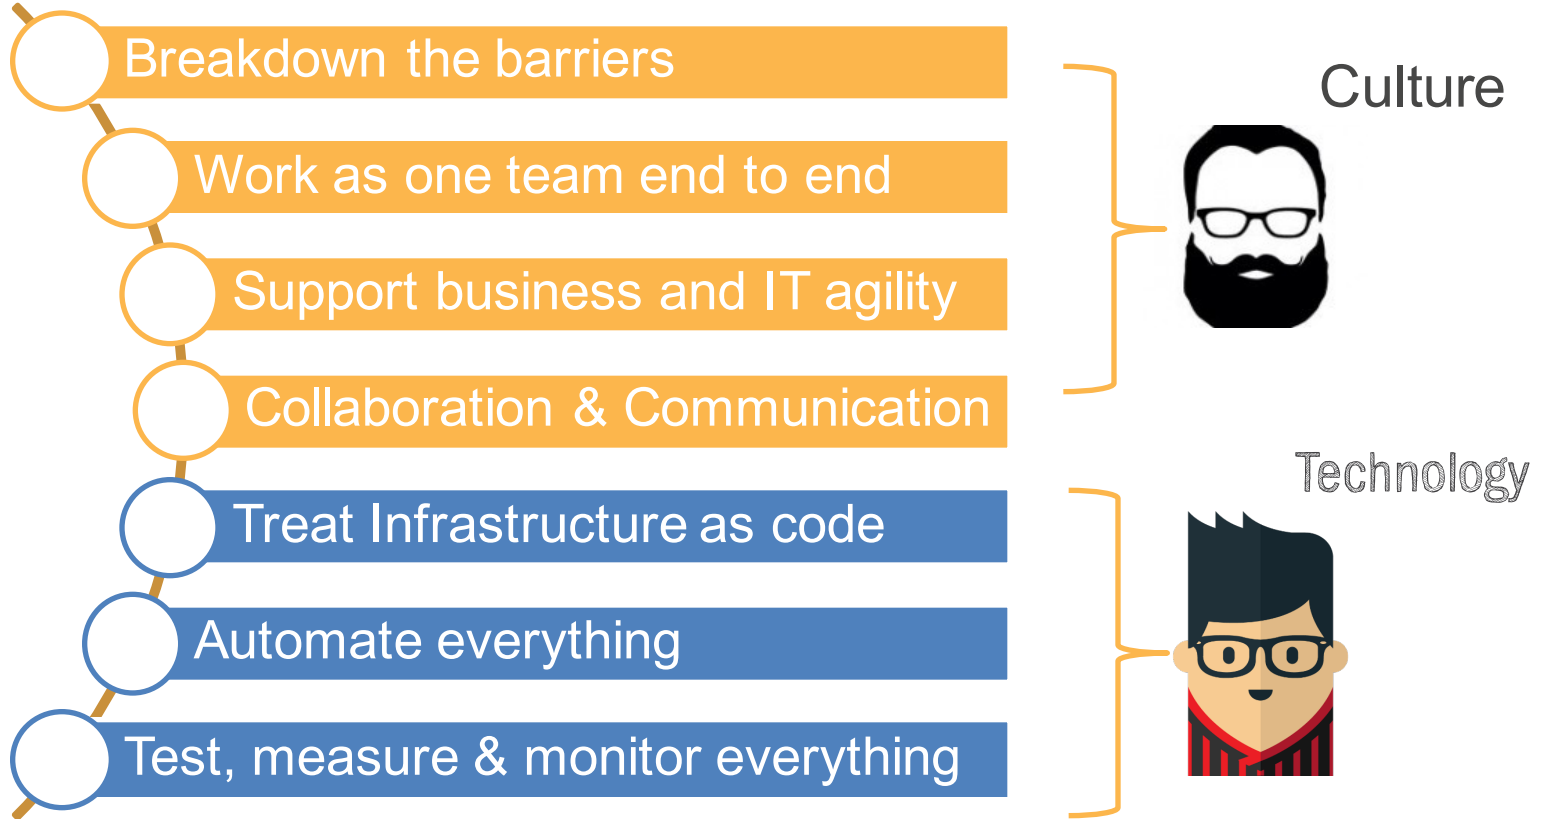
DevSecOps

Integrating Security with DevOps

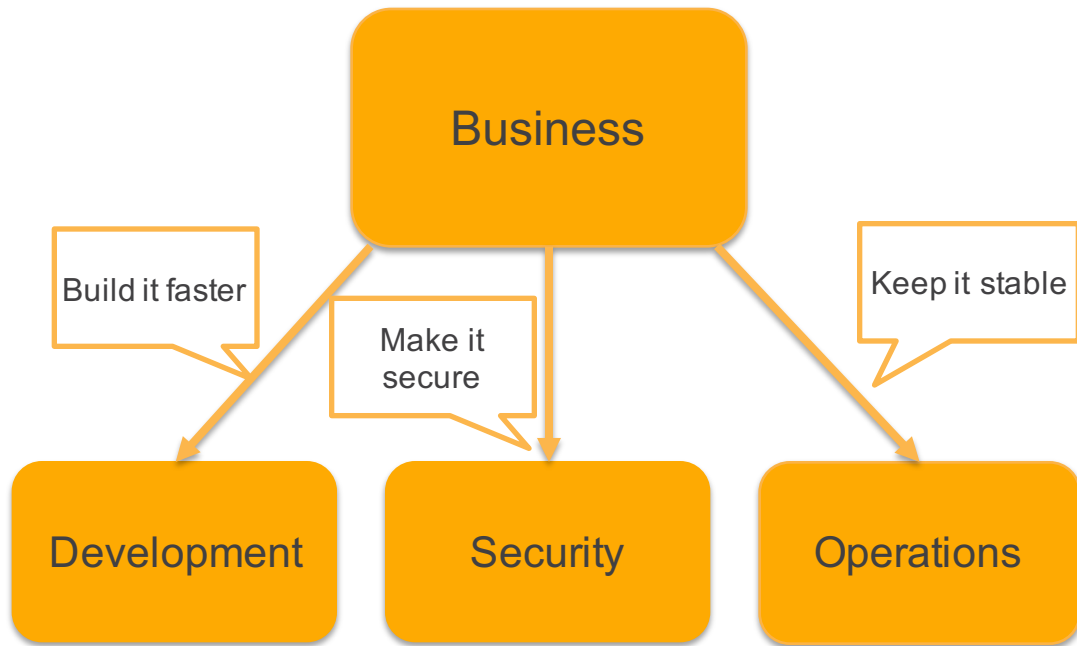
What is DevOps?

Cultural
Philosophy + Practices + Tools

What is DevOps

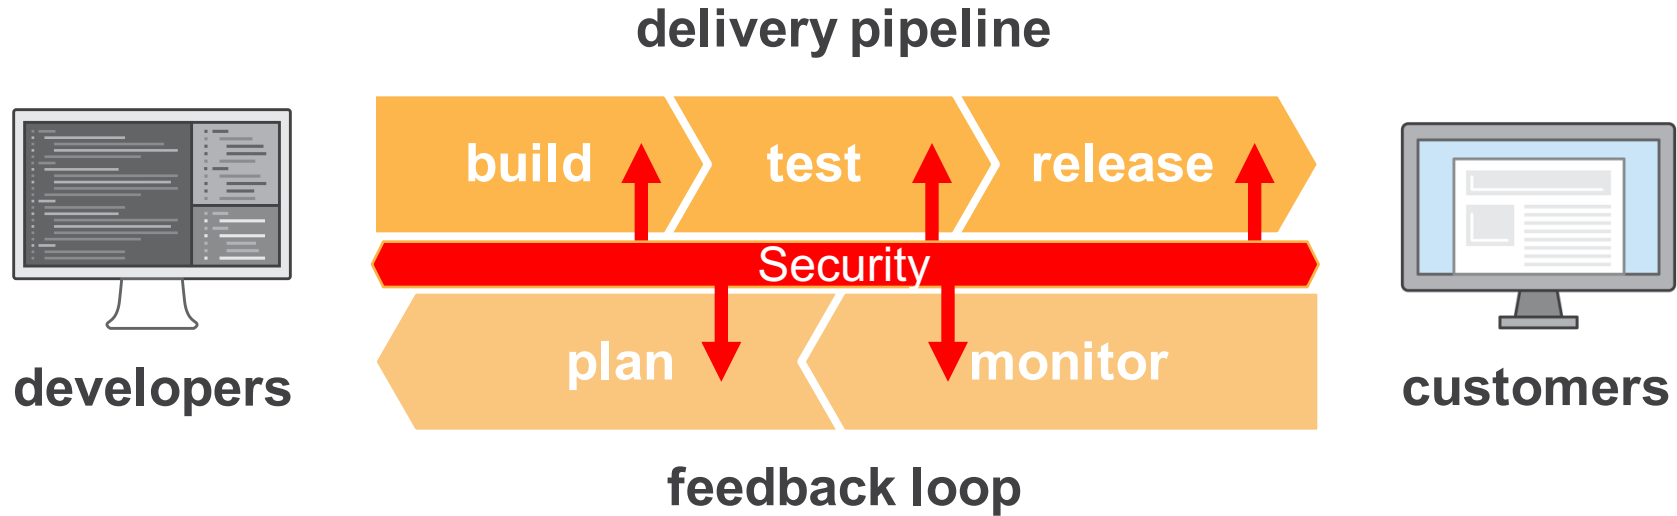


Competing Forces



What is DevSecOps

Software development lifecycle

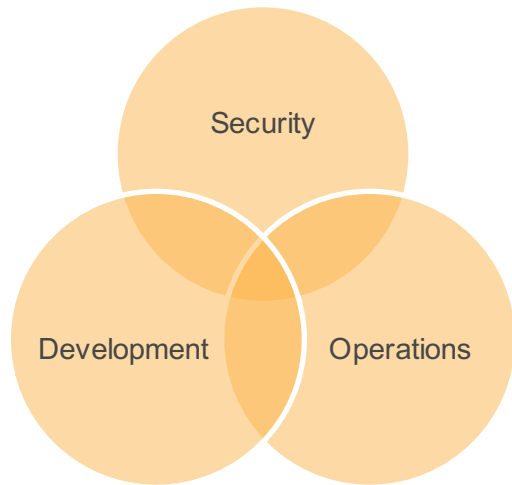


DevOps = Efficiencies that speed up this lifecycle

DevSecOps = Validate building blocks without slowing lifecycle

Who is DevSecOps

- DevSecOps is
 - Team/Community, not a person
 - Automated and autonomous security
 - **Security at scale**
- DevSecOps role
 - **Not** there to audit code
 - Implement the control segments to validate and audit code and artifacts as part of the CI/CD process





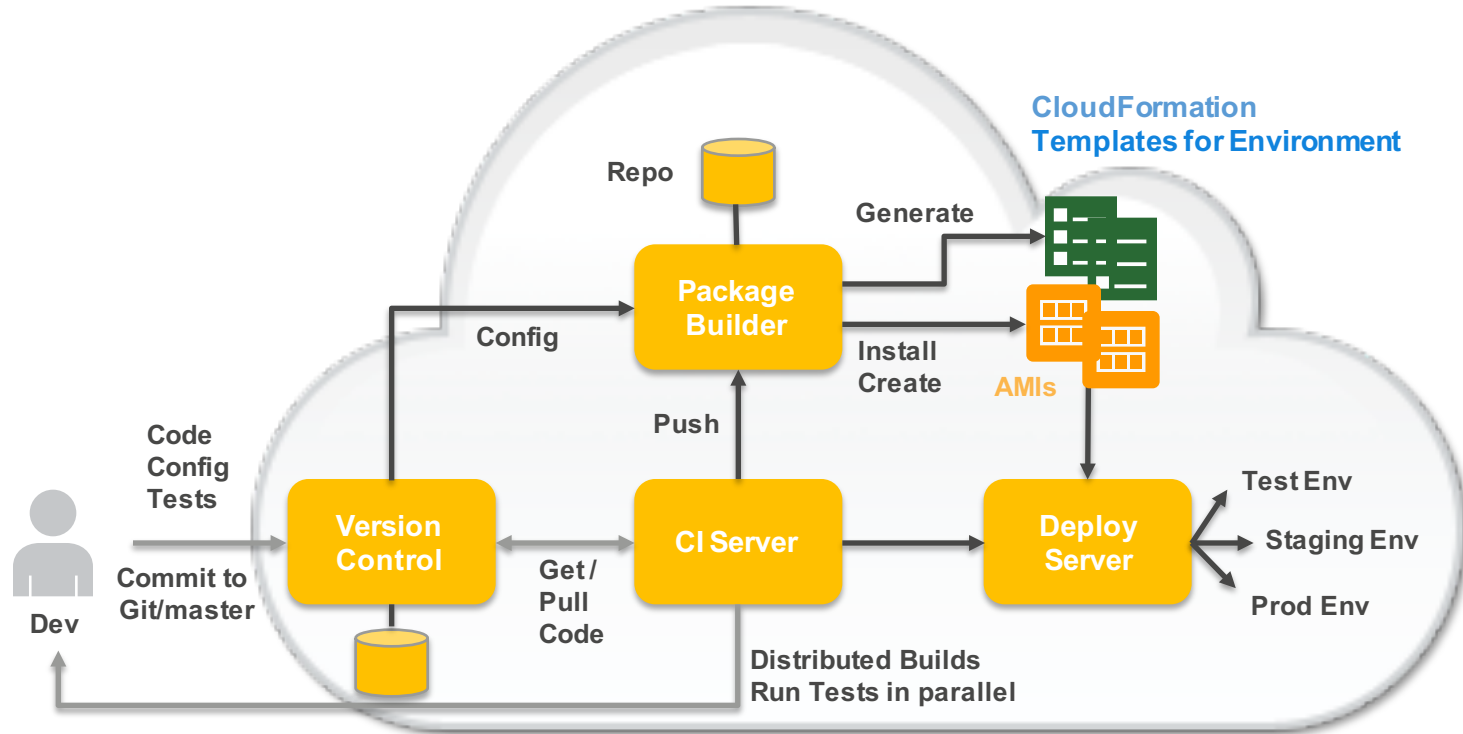
Continuous Integration

Techniques and tools to implement the **continuous process of applying quality control**; in general, small pieces of effort, applied frequently, to improve the quality of software, and to reduce the time taken to deliver it.

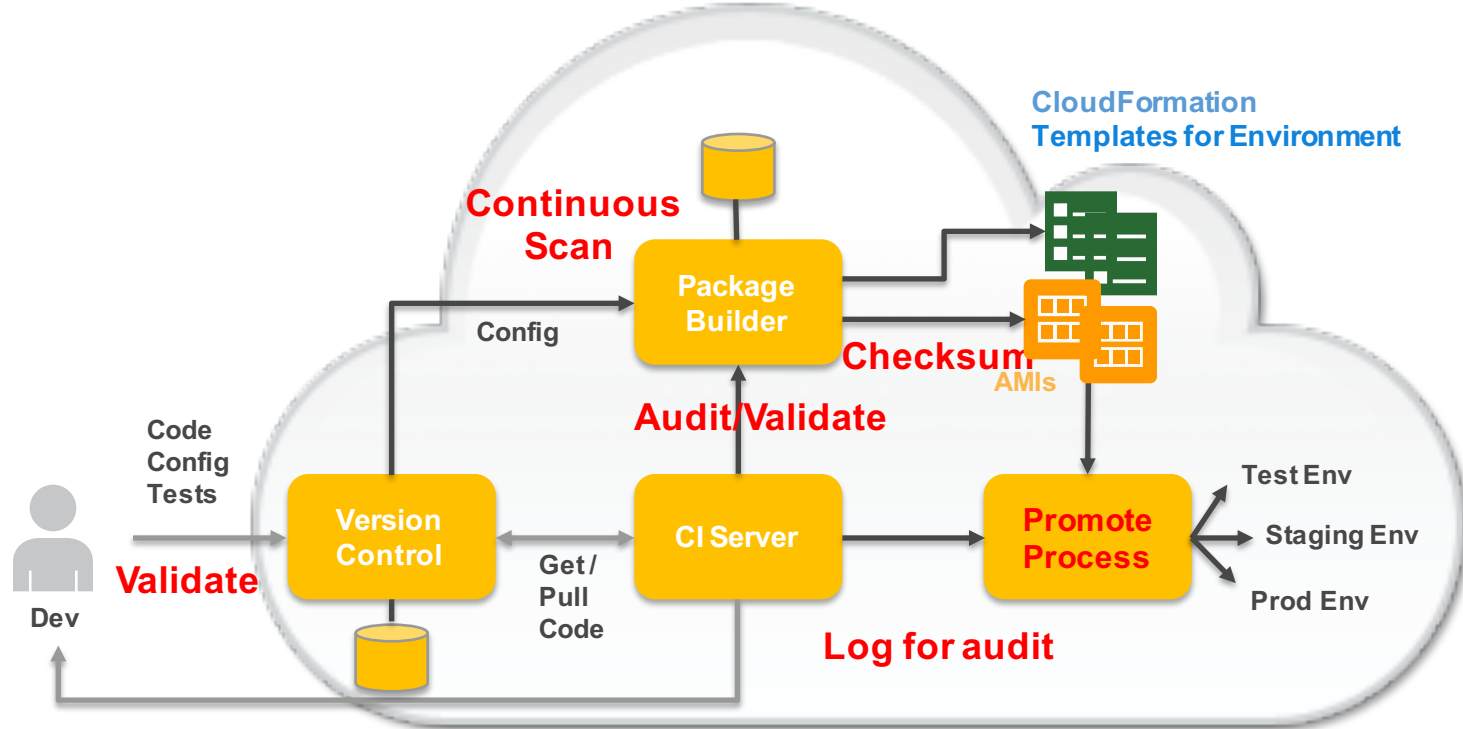
Continuous Delivery

Techniques and tools to **improve the process of software delivery**, resulting in the ability to rapidly, reliably, and repeatedly push out enhancements and bug fixes to customers at low risk and with minimal manual overhead.

CI/CD for DevOps

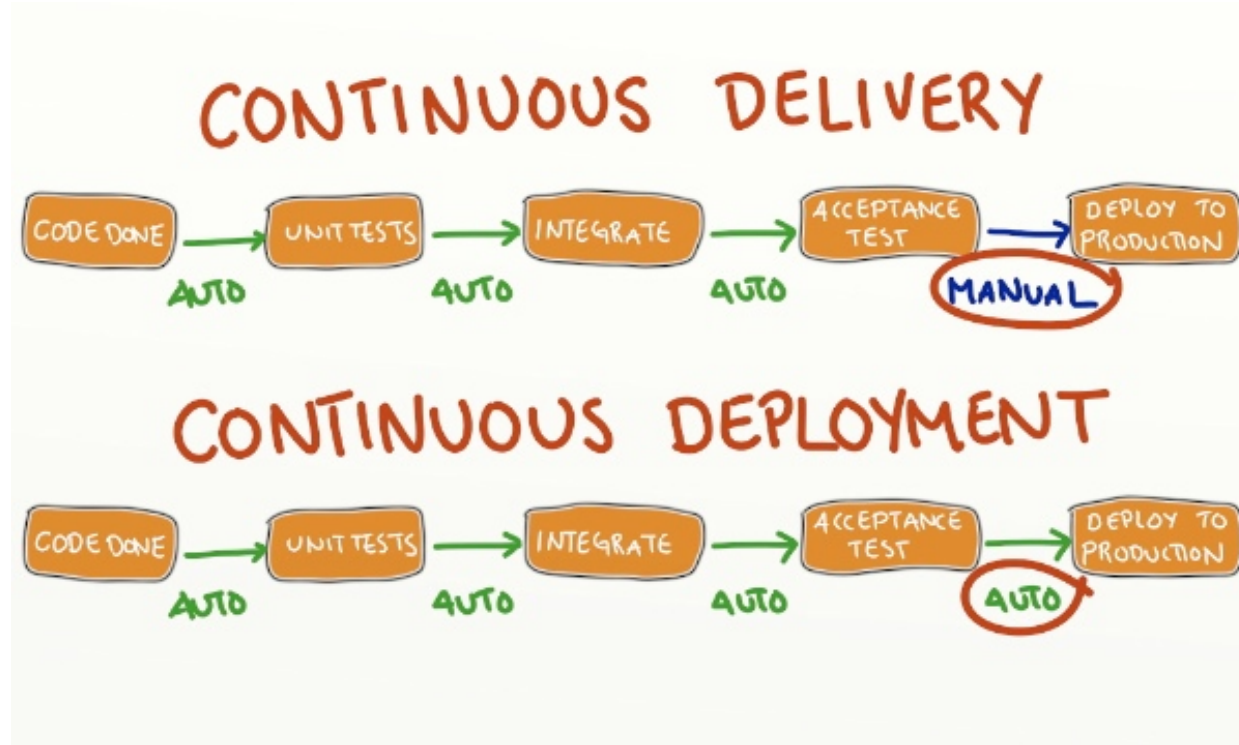


CI/CD for DevSecOps



Send Build Report to Security
Stop everything if audit/validation failed

Promotion Process in Continuous Deployment





What Does DevSecOps CI/CD Give Us?

- **Confidence** that our code is **validated** against corporate security policies.
- Avoid infrastructure/application failure in a later deployment due to different security configuration
- Match DevOps **pace of innovation**
- Audit and alert
- **Security at scale!**



Governance First

Infrastructure is code



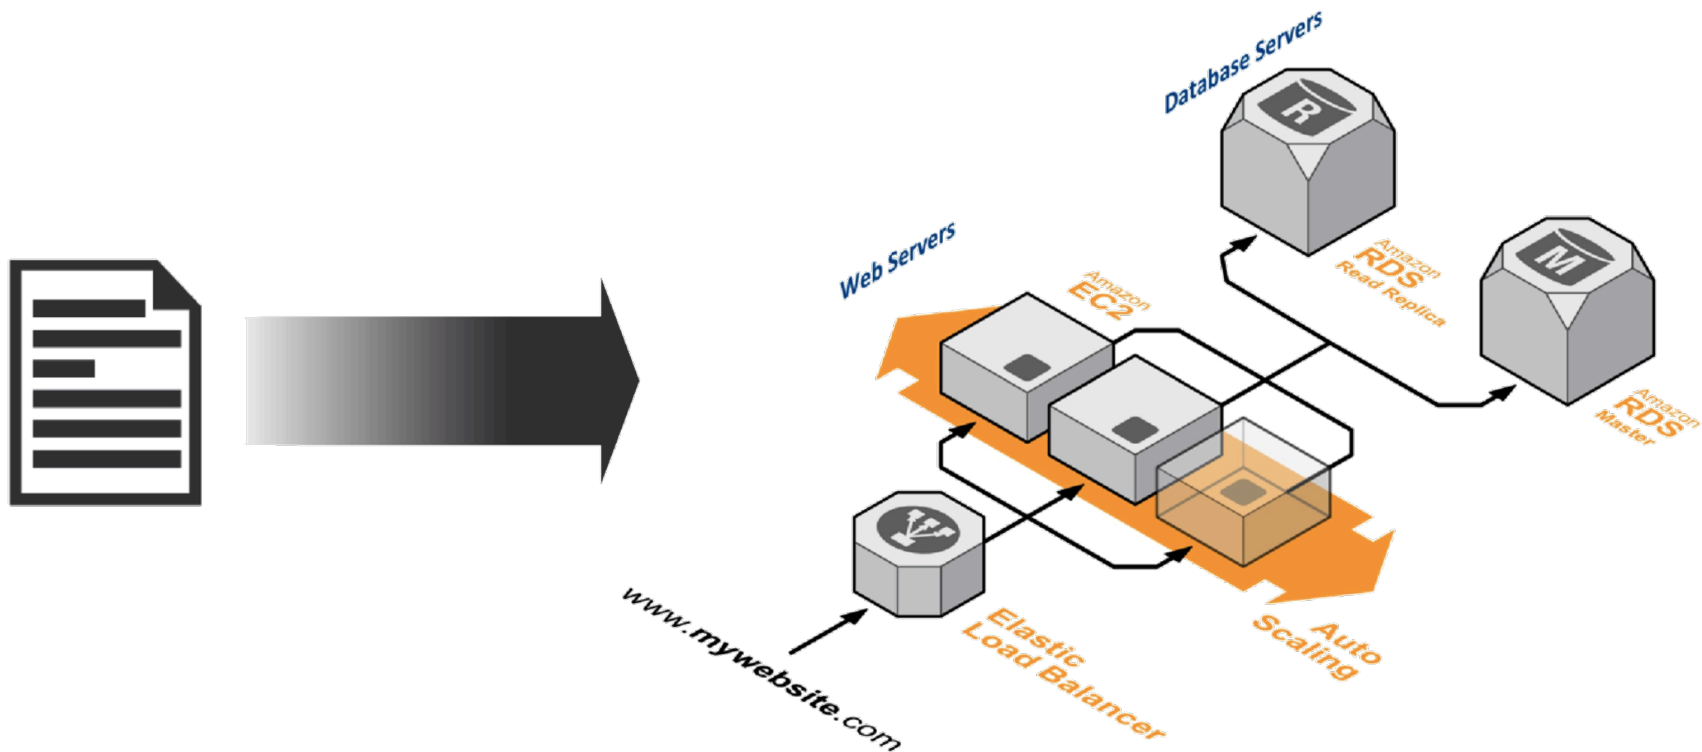
AWS CloudFormation Primer

Allows you to define a “template” which is composed of different “resources” and then provision that template into repeatable, live, “stacks”.

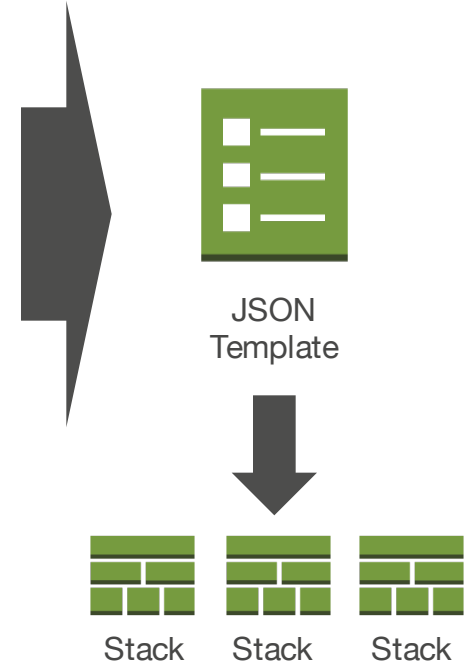
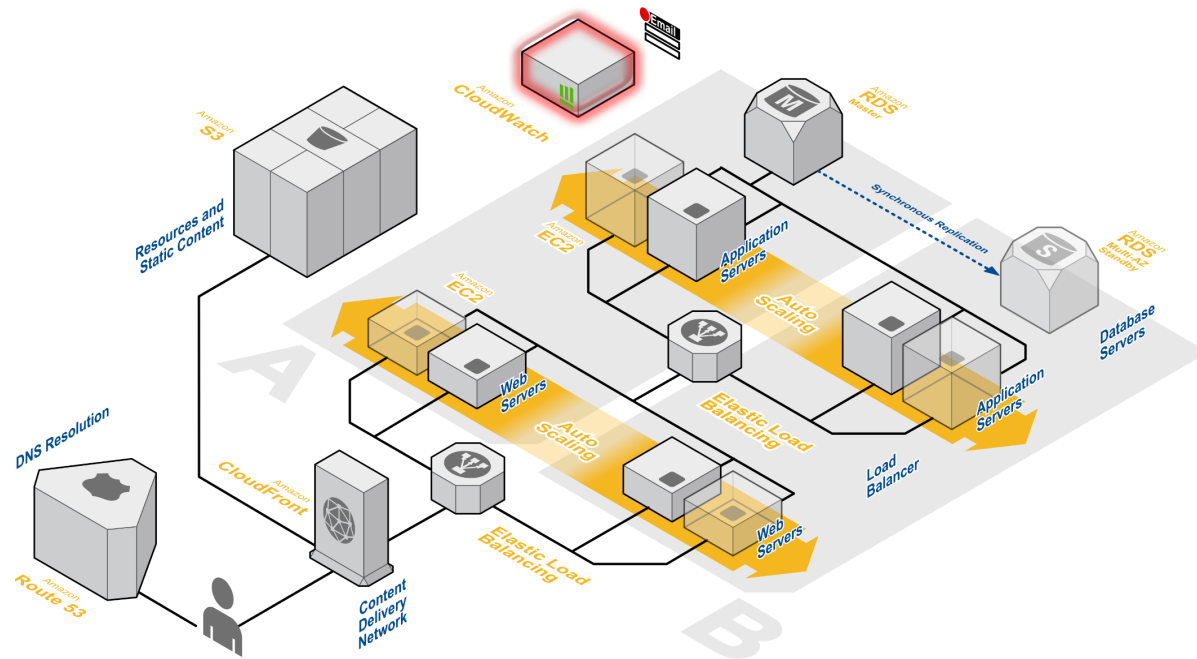
CloudFormation (CFn) providing a **single service interface** and abstracts the underlying API semantics and implementation mechanics.

CFn templates can hook into external configuration management frameworks such as Jenkins/Chef/Puppet/etc.

AWS CloudFormation Primer



AWS CloudFormation Stacks





CFn for Security Practitioners

“Infrastructure **is** code”

- We already know how to manage “code” lifecycle
- Let CFn perform state changes and govern who calls CFn

Split ownership templates

- Merge code from Sec, Dev and Ops
- Baseline and build/state management

Provides inventory and documentation

- Resources are the same, just scaled

Integrate with Service Catalog

- Provides UI
- Use constraints (example: tags)

```

graph LR
    A[code] --> B[11/11/11]
    B --> C((Execute))
    C --> D[11/11/11]
  
```



Create Skeleton

Define Resources



Split Ownership Configurations

Who knows your solution best?

- Dev, Infra, Sec...?
- Delegate ownership

Use Yaml and split file into chunks or functions

- Separate file sources with access control – Use IAM/VPC-E/etc.
- Push files -> Validate -> Merge files -> Validate -> Deploy -> Validate

Jenkins for deployment

- Promotion flows
 - Move from manual to Automation based on validation quality
- Excellent for merging jobs of split configurations

Merging

From single file or multiple files

- Maintain access control using policies
- Use different source stores if needed

Based on function/state

Reusable patterns

Maintain order, especially of validation

- Security validation last to execute
- Security should always win

```
2 baseline: &baseline
3   AWSTemplateFormatVersion: "2010-09-09"
4   Description: "Service Catalog dependencies."
5   Parameters: {}
6   Mappings: {}
7   Resources: &Resources
8     ManagementSecurityGroup: &ManagementSecurityGroup
9       Type: "AWS::EC2::SecurityGroup"
10      Properties:
11        GroupDescription: "Security group for management traffic"
12        VpcId: "<VPC-ID>"
13      Outputs: {}
14
15 SecDevOps-Production:
16   <<: *baseline
17   Resources:
18     <<: *Resources
19     ManagementGroupIngressRuleA:
20       Type: "AWS::EC2::SecurityGroupIngress"
21       Properties:
22         GroupId:
23           Ref: "ManagementSecurityGroup"
24         IpProtocol: "tcp"
25         FromPort: 22
26         ToPort: 22
27         CidrIp: "10.10.10.0/24"
28     ManagementGroupIngressRuleB:
29       Type: "AWS::EC2::SecurityGroupIngress"
30       Properties:
31         GroupId:
32           Ref: "ManagementSecurityGroup"
33         IpProtocol: "tcp"
34         FromPort: 22
35         ToPort: 22
36         CidrIp: "10.10.20.0/24"
37
38 SecDevOps-Test:
39   <<: *baseline
40   Resources:
41     <<: *Resources
42     ManagementGroupIngressRule:
43       Type: "AWS::EC2::SecurityGroupIngress"
44       Properties:
45         GroupId:
46           Ref: "ManagementSecurityGroup"
47         IpProtocol: "tcp"
48         FromPort: 22
49         ToPort: 22
50         CidrIp: "10.10.20.0/24"
```

Validation

Keep track of what section you are validating

- Stage vs Prod
- Merged vs separated

Validate often and log/alert

- Validate part and end result
- Run-time validation

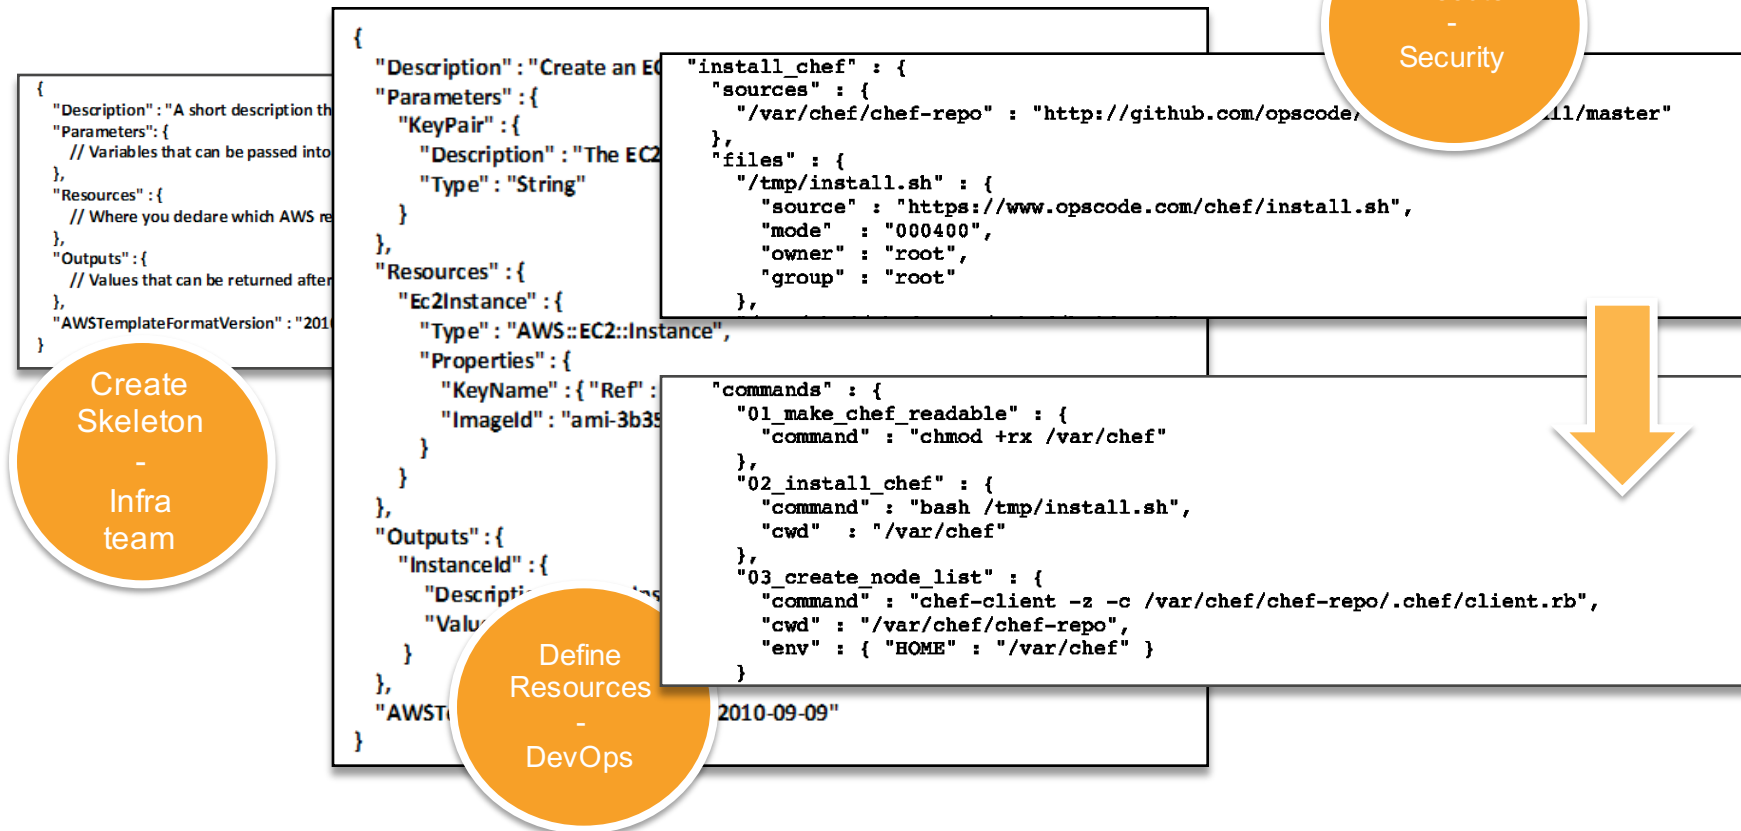
Use external agents

- AWS Simple Work Flow
- AWS Lambda
- Etc.

```
f = File.open("template.json", "rb")
g = f.read
g.gsub("","","\n").each_line {|line|
  if (line.downcase.include? "cidrip") && !line.include?("0.0.0.0/0")
    cfnCorpCIDRArray.push line[/((([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-4])[0-9]{1,3})\.([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-4])[0-9]{1,3})\.([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-4])[0-9]{1,3})/
    p errorCIDRSource
    fail = true
  elsif (line.downcase.include? "cidrip") && line.include?("0.0.0.0/0")
    cfnCIDRArray.push line[/((([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-4])[0-9]{1,3})\.([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-4])[0-9]{1,3})\.([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-4])[0-9]{1,3})/
    p errorCIDRSource
    fail = true
  end
}

# Check Corp CIDR blocks
strFail = ""
cfnCorpCIDRArray.each {|y|
  found = false
  corpCIDRArray.each {|x|
    cidr = NetAddr::CIDR.create(x)
    if ((cidr.contains?(y)) || (y == x))
      found = true
    end
  }
  if !found
    p errorCorpCIDRSource + " (#{"y"})"
    fail = true
  end
}
```

Structured deployment using Split ownership





Applied Governance

Building processes to segment and encapsulate application packages



Applied Governance deconstructed

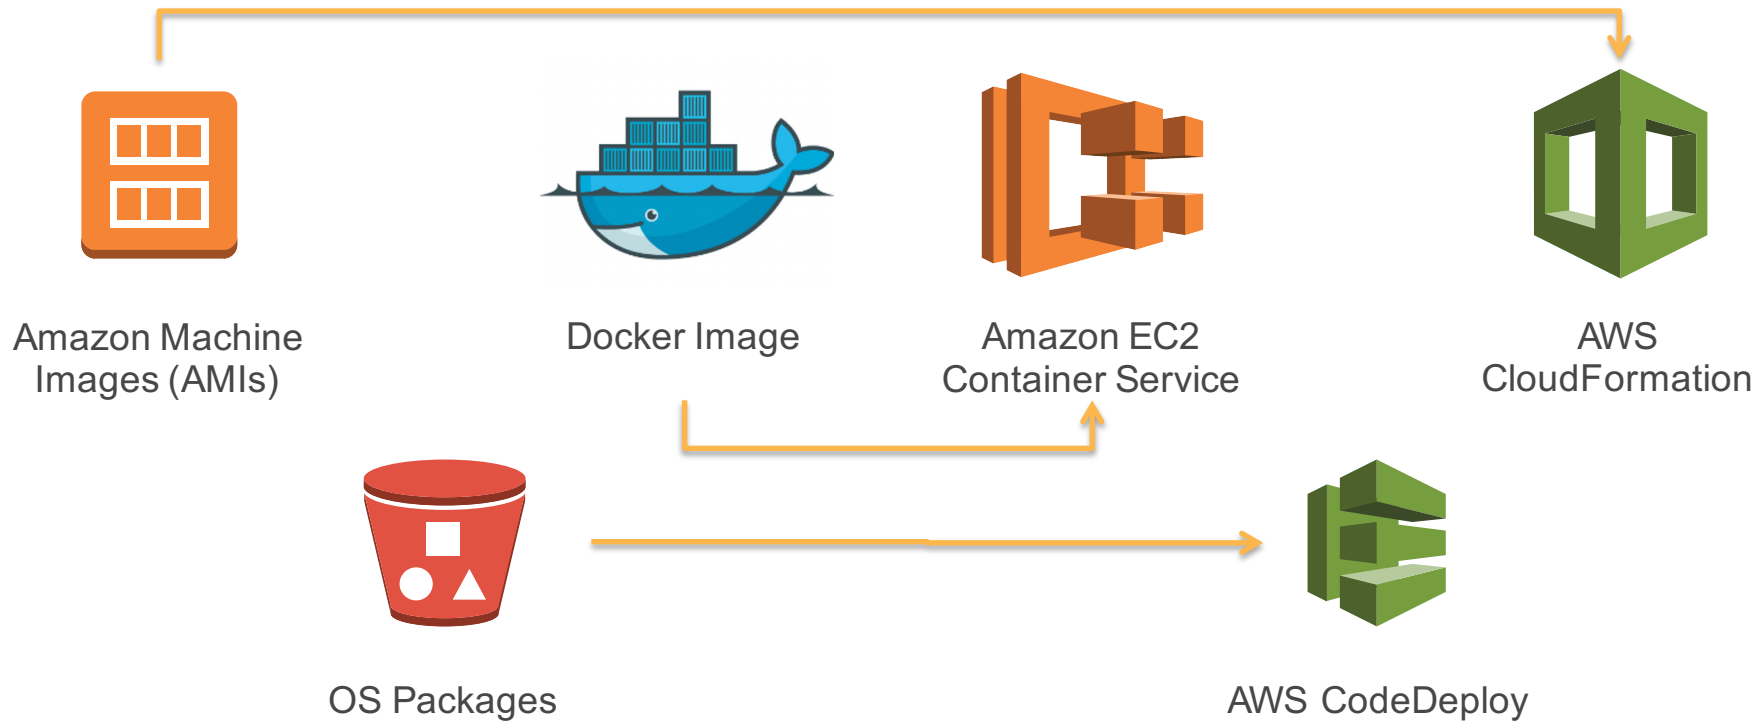
Deployment mechanisms for software artifacts

Configuration building blocks

Managing segmentation and security dependencies

...and a real scenario to wrap this all together.

Deployment Mechanisms for Software Artifacts



Deployment Mechanisms for Software Artifacts

Software Artifacts



Amazon Machine
Images (AMIs)

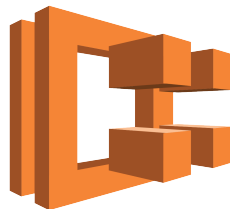


Docker Images



OS Packages

Deployment Services



Amazon EC2
Container Service



AWS
CloudFormation

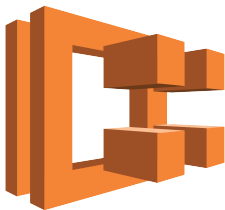


AWS CodeDeploy

Configuration building blocks



CloudFormation
Template



```
"name": "wordpress",  
"links": [  
  "mysql"  
],  
"image": "wordpress",  
"essential": true,  
"portMappings": [  
  {  
    "containerPort": 80,  
    "hostPort": 80
```

Task Definition



```
version: 0.0  
os: operating-system-  
files:  
  source-destination-  
permissions:  
  permissions-specific  
hooks:  
  deployment-lifecycle
```

Application
Specification File
(AppSpec file)



...and more.



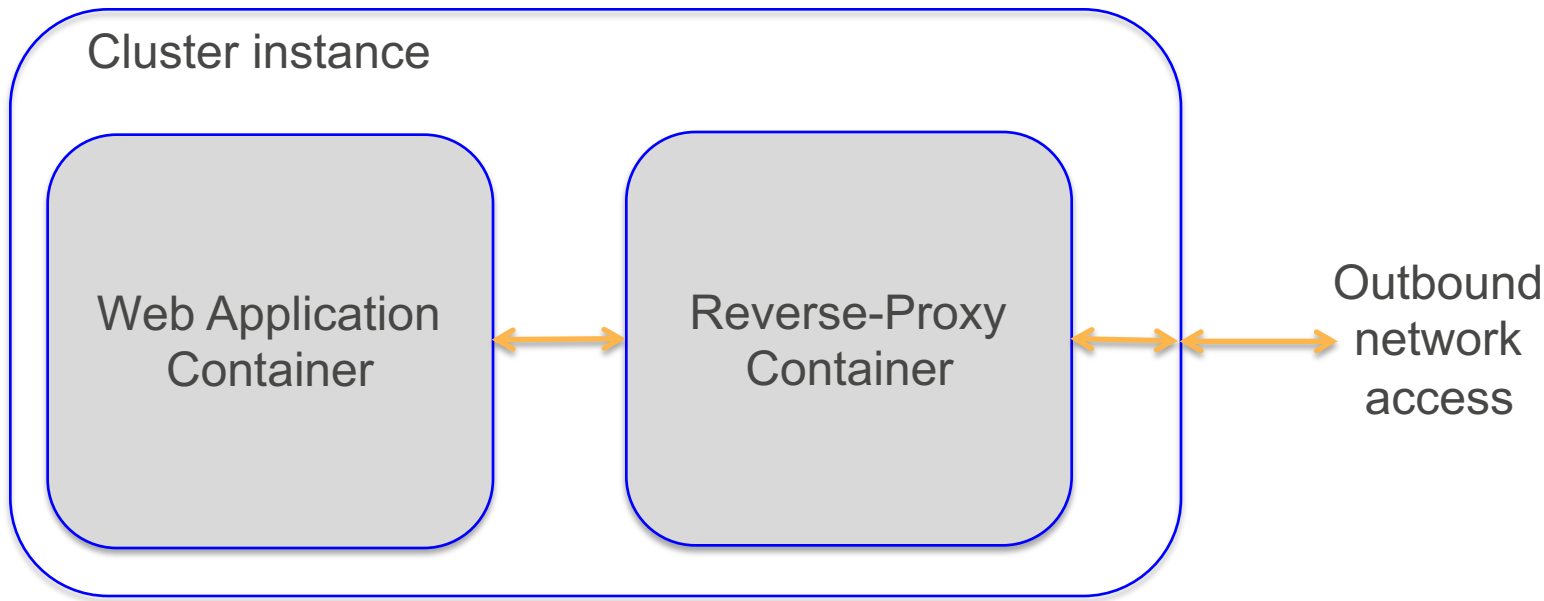
Managing segmentation and security dependencies

Why use these configuration building blocks to enforce security?

- Remove accidental conflicts
- Make security processes continuous and automatic
- Encapsulate software artifacts and implement controls one level up
- Control changes to these mechanisms via IAM

Example - Amazon EC2 Container Service

Rate limiting proxy





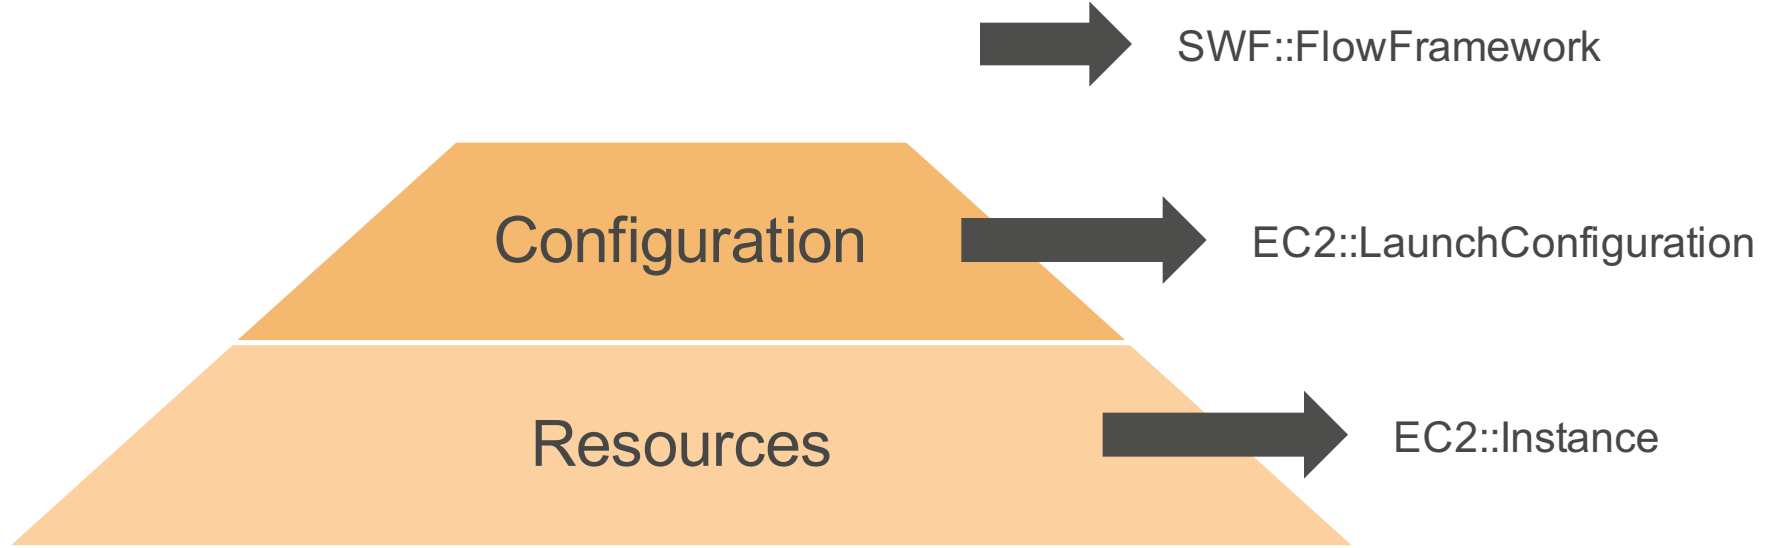
Scaling Trust Decisions

Autonomous security

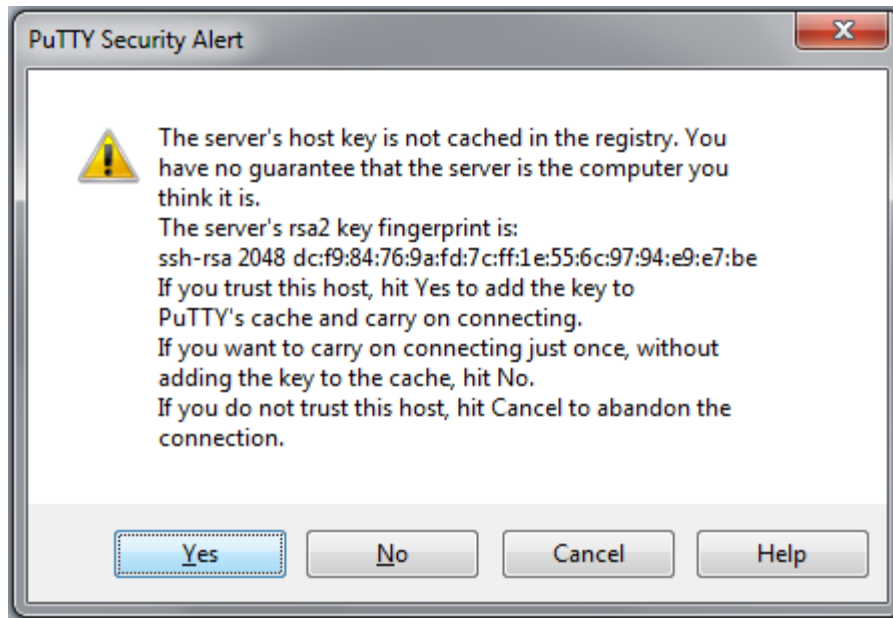
Does this flow feel familiar

1. Deploy Instance (*Chef got it so no worries...*)
2. Follow checklist #1 for patches (*Oops...New patches arrived...time to update Chef*)
3. Follow latest security list (Sec team don't have access to Chef yet...)
4. Put in production
5. SecOps called...they have a new list (*oh dear...*)
6. Pull from production
7. Rinse and repeat

Design time vs run time decisions



Why do we need these trust decisions?





Thought experiment: what is the root of trust for your Amazon EC2 instances?

How can we translate that into our W's?

- **Who** [is that instance]?
- **What** [is it's configuration]?
- **Where** [is the instance provisioned]?
- **When** [was it launched]?
- **Why** [is this instance alive]?

Scaling trust decisions - Allows

Allows Split Ownership

Allows intelligent decisions

- Inspect – Decide – Apply – **Validate** – Action - Log

Forced

- **Comply or be destroyed**

Security inventory for audits at scale

- Gather information from multiple sources
 - Host based
 - EC2 service
 - Created artifacts

Unique per resource action and information

- Keys/Certificates



Scaling trust decisions - Do

Use combination of design time and run time decisions

Tie to authentication on-demand for secure access of automated processes

Combine with Continuous scanning

- Code artifacts
- Resources (AWS Config)

Don't mistake run time decisions with post deployment

- Run-time happen at creation but before being made available

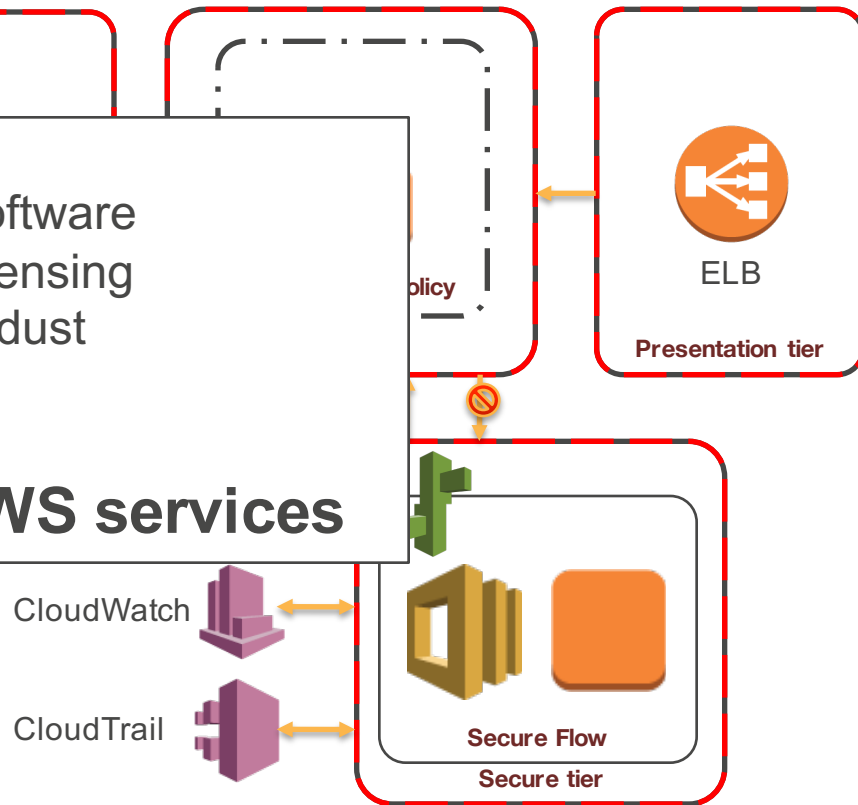
Push pattern – web server certificate

1. AutoScaling policy triggers
 - Hook Pending:Wait
2. Flow
 1. Create certificate *U
 2. Start web service
 - Certificate in me
 3. Delete certificate fro
 4. Install AWS CloudV
 5. Create SSH_RSAK
 6. Replace SSH_RSA
3. Validation
 1. Check CW Logs API
 2. Run WGET
4. Logging
 1. Log information/state in DynamoDB (IAM!)
5. Continue/terminate instance/hook
 1. Instance InService/terminated

No need for:

- Third party software
- Expensive licensing
- Magical fairy dust

All using AWS services



Run-time security flow – best practices

Avoid pull patterns

- Requires an available source/path from target system
- Own/compromise the target and you own the access process

Push ensures only outgoing traffic

- Instances have no path back to source
- Instances have no knowledge of the process/flow

Ephemeral/on-demand scripting on target

- Store in memory
- Create scripts on target on-demand at execution time
- Temporary credentials for first run

Keep source store on-instance (worker node) or S3 using dedicated VPC with VPC-Endpoint for S3 policy



Tooling - Triggering

AutoScaling Group

- Use Lifecycle hooks
- Automatic repair of non approved instances
- SNS for integration with Workers/Lambda
- Use least privileges on AutoScaling Groupd instances
 - Use IAM Instance Roles
- Unique credentials per instance



Tooling - Worker

Simple WorkFlow

- Flow engine for Ruby/Java
- Allow synchronous or asynchronous flows
- Deciders/Workers
 - EC2 or Lambda
- Can be integrated into Elastic Beanstalk
- Use internal or external source for validation tasks/tests
 - Checksum files before use
 - Use Lambda/Elastic Beanstalk workers



Real-life trust applications

- Agent deployment/validation
- Integration with Configuration Management (Chef/Puppet)
- Unique per-instance configuration
- Cross-instance network authentication
- Joining an Active Directory

Ask yourself what makes this instance trustworthy for any of the above and how can I scale that trust?

Thank You



**Please let us have your
feedback on this event!**

**Don't forget to leave your badge and collect
your \$50 AWS Credit Token**